



# MLOps Fundamentals

## Student Guide

SG-840034-002

**skillsoft**  
**global knowledge**<sup>TM</sup>



# Course Information

## Copyright

---

Copyright © 2025 by Global Knowledge Training LLC.

The following publication, *MLOps Fundamentals Student Guide*, was developed by Global Knowledge Training LLC. All rights reserved. No part of this publication may be reproduced or distributed in any form or by any means without the prior written permission of the copyright holder.

Products and company names are the trademarks, registered trademarks, and service marks of their respective owners.

## *Welcome!*

---

Thank you for selecting Global Knowledge as your training provider. Global Knowledge is the worldwide leader in IT training and learning services. We empower organizations, teams, and individuals with the skills and best practices necessary to leverage the technologies critical for sustained success and to create competitive advantage.

Learn more at [www.globalknowledge.com](http://www.globalknowledge.com).

## *Certification and career paths*

---

For more information on certification and additional courses that can help you achieve your career goals, please visit our Web site at [www.globalknowledge.com](http://www.globalknowledge.com).

---

# Table of Contents

---

Introduction .....	1
Module 1: MLOps and Data.....	3
Topic 1: Introduction to MLOps .....	4
Topic 2: Machine Learning Lifecycle Overview.....	8
Topic 3: MLOps Components and Tools.....	12
Topic 4: Setting Up an ML Project .....	18
Topic 5: Data Management Fundamentals .....	25
Topic 6: Feature Stores.....	35
Topic 7: Model Development.....	45
Topic 8: Implementing a Basic ML Pipeline .....	54
Module 2: Development, Experimentation and Evaluation .....	71
Topic 1: Model Development Strategies .....	72
Topic 2: ML Model Interpretability and Explainability.....	82
Topic 3: Implementing Algorithms .....	92
Topic 4: Experiment Tracking and Model Evaluation .....	99
Topic 5: Setting Up MLflow for Experiment Tracking .....	107
Topic 6: Evaluating Models.....	114
Topic 7: Hyperparameter Tuning Techniques .....	118
Topic 8: Automated Hyperparameter Tuning.....	125
Module 3: Deployment, Monitoring, and Pipelines.....	135
Topic 1: Model Serving and Development Strategies.....	136
Topic 2: Legal and Compliance Issues in MLOps .....	146
Topic 3: Containerizing ML Models with Docker .....	153
Topic 4: Deploying Models to Cloud Platforms .....	159
Topic 5: Federated Training and Edge Deployments.....	165
Topic 6: CI/CD for ML .....	174

## MLOps Fundamentals Student Guide

Topic 7: Setting Up CI/CD for Pipeline for ML.....	184
Topic 8: Monitoring and Maintaining ML Systems .....	190
Topic 9: Implementing Monitoring Tools .....	197
Appendix: Knowledge Check Answer Key.....	205
Module 1: MLOps and Data.....	205
Module 2: Development, Experimentation and Evaluation .....	212
Module 3: Deployment, Monitoring, and Pipelines.....	217

## Course Agenda

### Module 1: MLOps and Data

- Introduction to MLOps
- Machine Learning Lifecycle Overview
- MLOps Components and Tools
- Setting Up an ML Project
- Data Management Fundamentals
- Feature Stores
- Model Development
- Implementing a Basic ML Pipeline

## Course Agenda

### Module 2: Development, Experimentation, and Evaluation

- Model Development Strategies
- ML Model Interpretability and Explainability
- Implementing Algorithms
- Experiment Tracking and Model Evaluation
- Setting up MLflow for Experiment Tracking
- Evaluating Models
- Hyperparameter Tuning Techniques
- Automated Hyperparameter Tuning

## Course Agenda

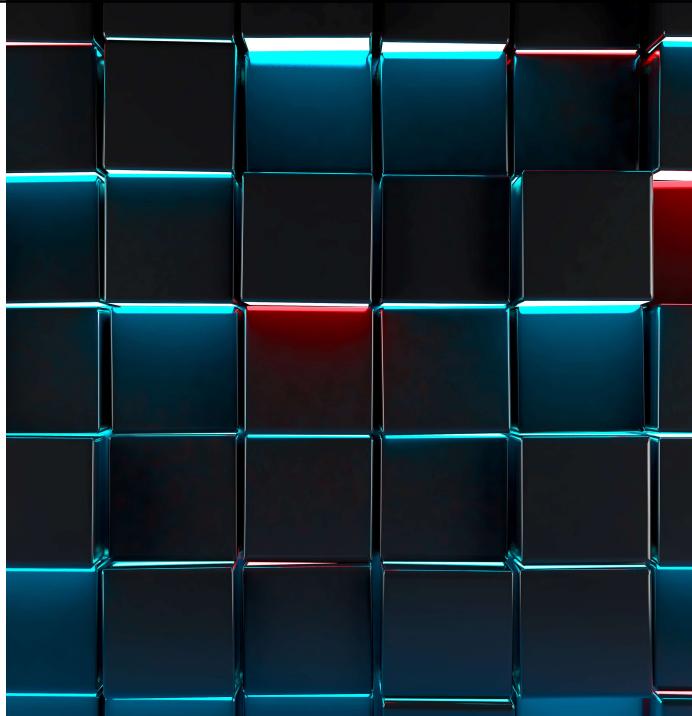
### Module 3: Deployment, Monitoring, and Pipelines

- Model Serving and Deployment Strategies
- Legal and Compliance issues in MLOps
- Containerizing ML Models with Docker
- Deploying Models to Cloud Platforms
- Federated Training and Edge Deployments
- CI/CD for ML
- Setting up CI/CD Pipelines for ML
- Monitoring and Maintaining ML Systems
- Implementing Monitoring Tools

## Course Objectives

After completing this course, you will be able to

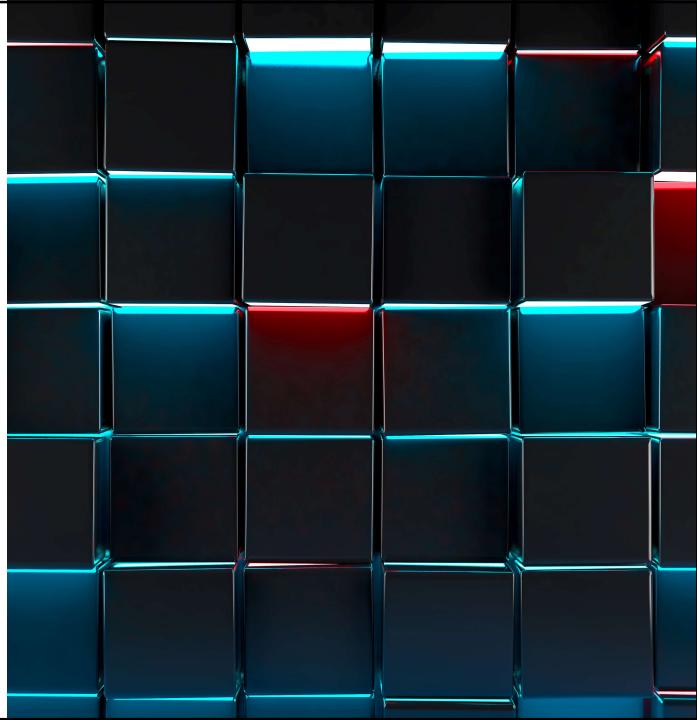
- Implement the MLOps lifecycle for effective model management and deployment
- Automate ML workflows using CI/CD and experiment tracking tools
- Deploy and serve models with Docker and cloud platforms
- Optimize model performance with advanced hyperparameter tuning



## Course Objectives

After completing this course, you will be able to

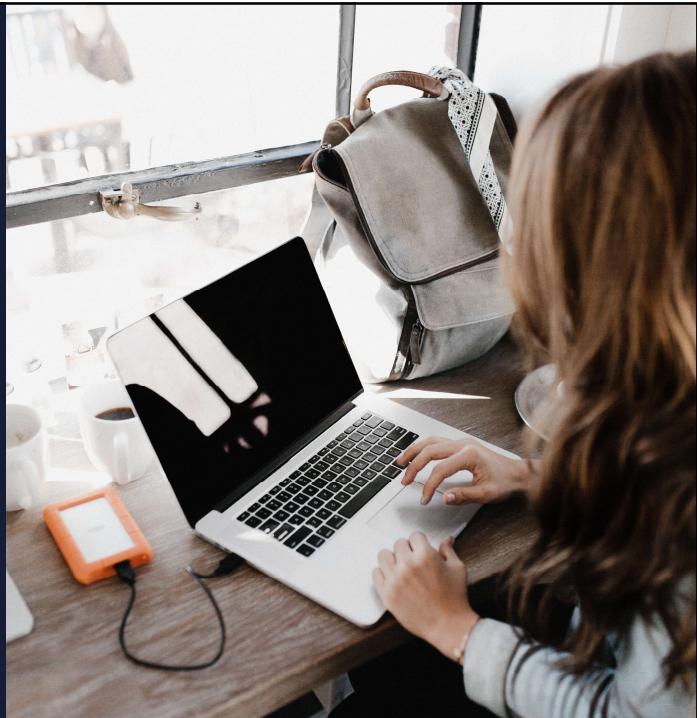
- Monitor and maintain ML models in production with real-time alerts
- Ensure model interpretability and compliance with regulatory standards



## Module 1: MLOps and Data

After today's module, you will be able to

- Understand the fundamentals of MLOps and its importance in ML projects
- Manage data lifecycles, versioning, and governance for effective data management
- Perform exploratory data analysis, preprocessing techniques in ML
- Implement basic ML pipeline and automation



## Module Agenda

### Module 1: MLOps and Data

- Introduction to MLOps
- Machine Learning Lifecycle Overview
- MLOps Components and Tools
- Setting Up an ML Project
- Data Management Fundamentals
- Feature Stores
- Model Development
- Implementing a Basic ML Pipeline

Module 2: Development, Experimentation, and Evaluation

Module 3: Deployment, Monitoring, and Pipelines

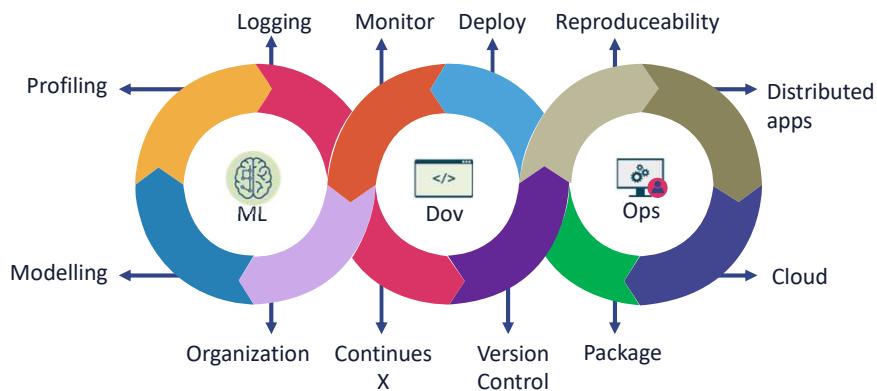


### Topic 1: Introduction to MLOps

- What is MLOps?
- Importance of MLOps
- Evolution of MLOps from DevOps
- DevOps vs. DataOps vs. AIOps
- Key challenges MLOps addresses

## What is MLOps?

**Definition:** MLOps (machine learning operations) is a set of practices that combines machine learning, data engineering, and DevOps to reliably and efficiently deploy ML models in production



- **Scalability:** Ensures seamless scaling of ML models across diverse environments
- **Efficiency:** Streamlines workflows, reducing time and effort required for model deployment
- **Reproducibility:** Facilitates consistent and repeatable ML pipelines
- **Monitoring & maintenance:** Provides tools for monitoring model performance and managing model decay over time

## Importance of MLOps

## MLOps Evolution from DevOps

- **DevOps:** Automates software lifecycle
- **MLOps:** Extends automation to ML model lifecycle
- **Goal:** Integrate ML workflows with operational practices

## DevOps vs. DataOps vs. AIOps

**DevOps:** Combines software development and IT operations to shorten the development lifecycle

**DataOps:** Focuses on data management, ensuring data quality and accessibility for analytics and ML

**AIOps:** Uses AI to enhance IT operations, focusing on automation and proactive problem-solving

## Key Challenges MLOps Addresses

- **Deployment:** Automates model deployment
- **Version control:** Manages model versions
- **Scalability:** Ensures efficient scaling
- **Monitoring:** Tracks performance and data drift

## Knowledge Check

1. What is the primary goal of MLOps?
  - A. Managing the ML model lifecycle
  - B. Developing software applications
  - C. Creating data pipelines
  - D. Monitoring the application performance



## Discussion

### Common pain points in traditional ML Processes

- Data quality & availability
- Model selection & tuning
- Deployment & scalability
- Cost & resource management
- Overfitting/underfitting

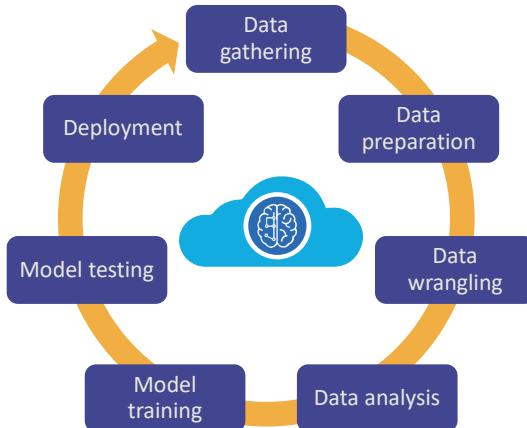


## Topic 2: Machine Learning Lifecycle Overview

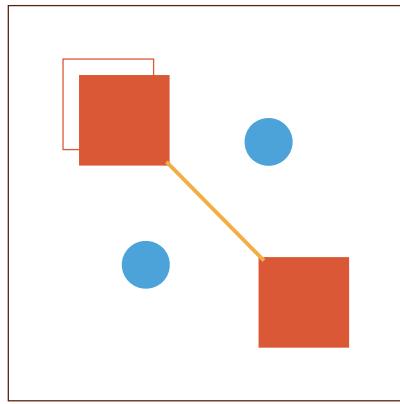
- Lifecycle overview
- Data collection
- Model training
- Model deployment
- Monitoring and maintenance

## ML Lifecycle Overview

**Definition:** Series of stages involved in developing, deploying, and maintaining machine learning models



**Definition:** Process of feeding data to the algorithm to learn patterns so that the model generalizes well to the new, unseen data



## Model Training

**Definition:** Integrating the trained model into production systems to make predictions on real-world data, enabling actionable insights



## Model Deployment

**Definition:** Continuously tracking model performance and retraining as needed to detect performance degradation and data drift to maintain model accuracy



## Monitoring and Maintenance

- **Definition:** Gathering relevant data from various sources for model training
- **Importance:** High-quality, diverse data is crucial for model accuracy and performance



## Data Collection

### Knowledge Check

2. MLOps is solely concerned with the deployment of machine learning models.
  - A. True
  - B. False



## Topic 3: MLOps Components and Tools

- MLOps components overview
- Data versioning and management
- Model versioning and experiment tracking
- Automated testing
- CI/CD pipeline for ML
- Containerization and orchestration
- Monitoring and logging
- Infrastructure as Code (IaC)
- Model governance and security
- Continuous training and feature store

## MLOps Components Overview

What are the components of MLOps architecture?

- 1 Data preparation in the pipeline
- 2 CI/CD model deployment pipeline
- 3 Model evaluation and feature repository
- 4 Model deployment and version control
- 5 Automated model monitoring

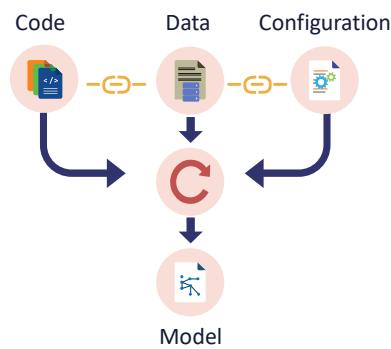
# Data Versioning and Management

**Definition:** Managing different versions of datasets used for model training

- **Tools:** DVC, Pachyderm
- **Benefits:** Ensures reproducibility and traceability of experiments

**Definition:** Keeping track of different model versions and experiment results

- **Tools:** MLflow, Neptune
- **Benefits:** facilitates collaboration and reproducibility of model results



# Model Versioning and Experiment Tracking

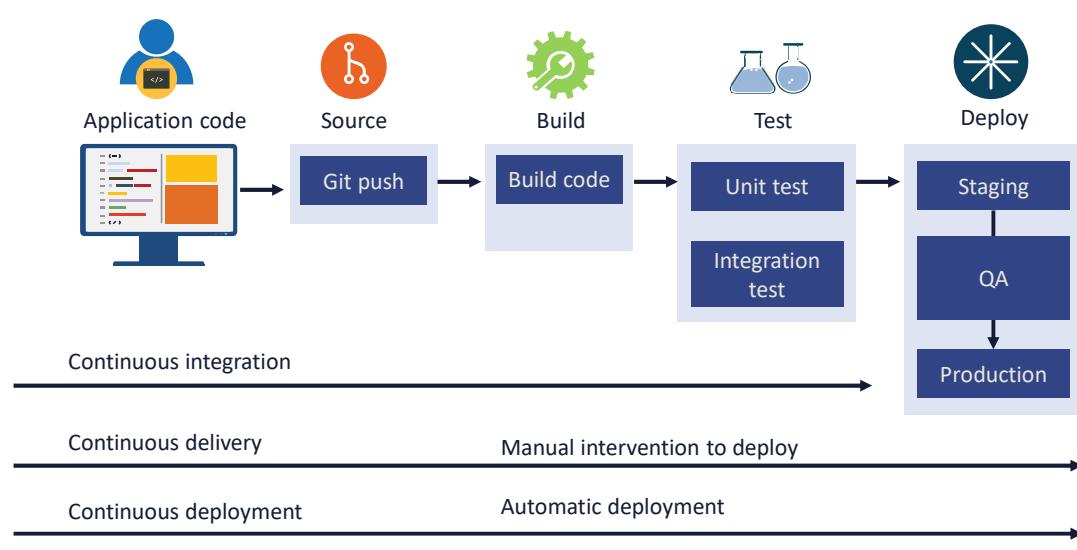
## Automated Testing

**Definition:** Testing models and data pipelines automatically to catch errors early

- **Tools:** PyTest, Great Expectations
- **Benefits:** ensures model integrity and data quality before deployment

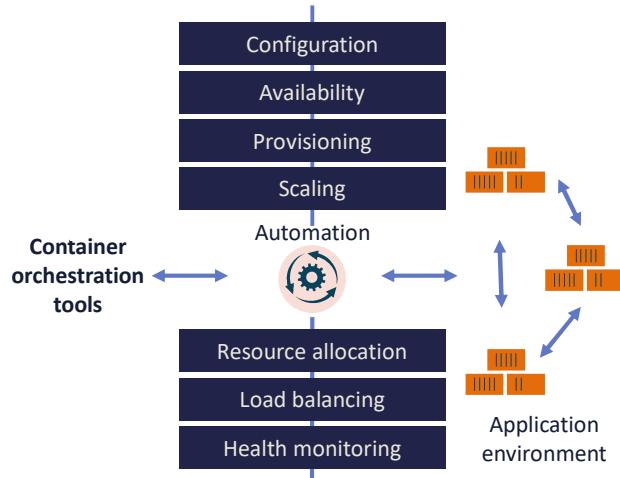


## CI/CD Pipeline for ML



# Containerization and Orchestration

- **Definition:** Packaging models with dependencies for consistent deployment
    - **Tools:** Dockers, Kubernetes
    - **Benefits:** simplifies model deployment across different platforms



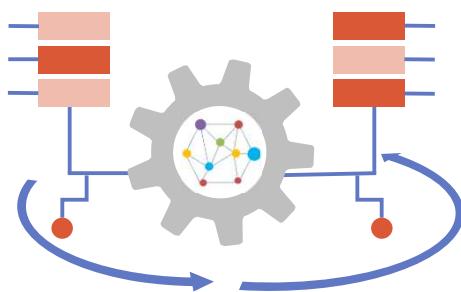
**Definition:** Tracking model performance and capturing logs in real-time

- **Tools:** Prometheus, ELK (Elasticsearch, Logstash, Kibana) Stack
  - **Benefits:** Helps detect performance degradation and ensures system reliability

# Monitoring and Logging

**Definition:** Ensures compliance, security, and ethical use of ML models

- **Tools:** Seldon, TensorFlow Privacy
- **Benefits:** Protects sensitive data and maintains regulatory compliance

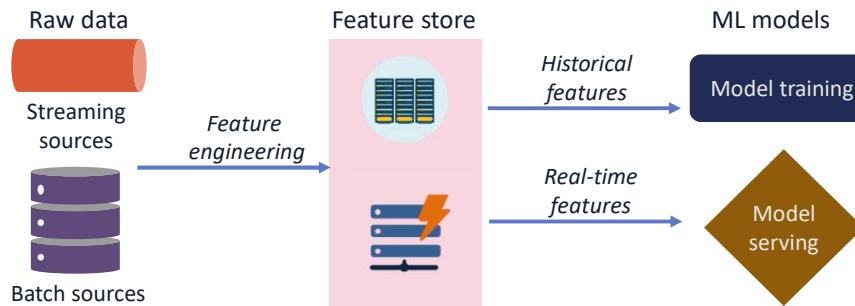


## Model Governance and Security

## Continuous Training (CT) and Feature Store

**Definition:** Regularly retraining models to adapt to new data

- **Tools:** TFX, Feast
- **Benefits:** Keeps models updated with changing data and business needs



## Knowledge Check

3. Tools like MLFlow and Kubeflow help automate various stages of the ML lifecycle.
  - A. True
  - B. False

## Knowledge Check

4. Which of the following tools is used for experiment tracking in MLOps?
  - A. Dockers
  - B. MLFlow
  - C. Kubernetes
  - D. Terraform

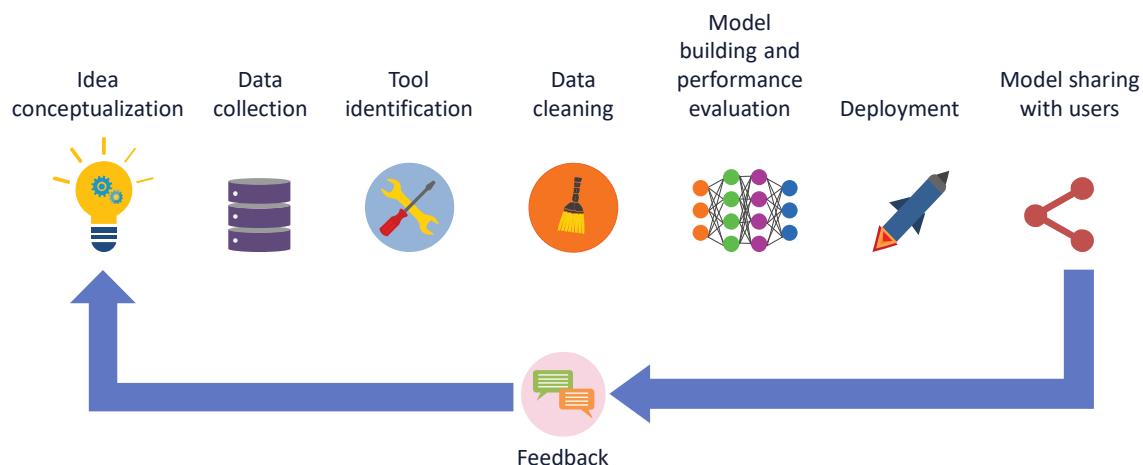


## Topic 4: Setting Up an ML Project

- Introduction to setting up an ML project
- Git and GitHub setup
- Creating a GitHub repository
- Understanding version control in ML projects
- Importance of virtual environments
- Creating and managing virtual environments
- What are pre-commit hooks?
- Setting up pre-commit hooks
- Benefits of using pre-commit hooks

## Introduction to Setting up an ML Project

### Guide to building a machine learning model



- **Git**: Version control system to track changes
- **GitHub**: Platform to host and share repositories
- **Command**: `git init` to initialize a local repository

## Git and GitHub Setup

### Creating a GitHub Repository

Navigate to GitHub and create a new repository

- Click on a “New Repository”
- Enter repository name and description
- Choose visibility and initialize with README

## Version Controls in ML Projects

- Track changes in code and experiments
- Allow collaboration without conflicts
- Maintain a history of model versions

- Isolate project dependencies
- Avoid conflict between different projects
- Manage dependencies effectively

## Virtual Environments

## Creating and Managing Virtual Environments

- **Tools:** venv, conda
- **Commands:**
  - `python -m venv myenv`
  - `conda create --name myenv`

## Benefits of Virtual Environments

**Isolation:** Keeps project dependencies separate, preventing conflicts between different projects

**Reproducibility:** Ensures that the environment can be recreated easily on different systems

**Version control:** Allows using different library versions across projects without interference

## What are Pre-Commit Hooks?

- Scripts that run before commits
- Enforce code quality and standards
- Prevent bad code from being committed

- Install pre-commit tool
- Create `.pre-commit-config.yaml`
- Add desired hooks and install them

## Setting Up Pre-Commit Hooks

## Infrastructure as Code (IaC)

**Definition:** Managing and provisioning infrastructure through code

- **Tools:** Terraform, AWS CloudFormation
- **Benefits:** Enables reproducible and scalable infrastructure deployment

## Benefits of Using Pre-Commit Hooks

- Ensures code quality
- Catches errors early
- Enforces coding standards

## Knowledge Check

5. Which of the following tools is used for managing a virtual environment in Python?
  - A. Git
  - B. Docker
  - C. venv
  - D. Prometheus

## Lab 1: Initialize an ML Project with GitHub, Virtual Environments, and Pre-Commit Hooks

### Instructions

1. Create and set up the GitHub repository
2. Create and activate a Python virtual environment
3. Configure pre-commit hooks for linting and testing
4. Push changes to GitHub with Git tracking

Estimated completion time: 30 minutes



## Topic 5: Data Management Fundamentals

- Data management fundamentals
- Understanding data lifecycles
- Data collection
- Data processing
- Data storage
- Data archiving
- Importance of data versioning
- Techniques for data versioning
- Data governance
- Compliance and regulations



## Topic 5: Data Management Fundamentals

- Data management practices
- Data storage solutions
- Cloud vs. on-premises storage
- Selecting the right solution
- Ensuring data quality and integrity

## Data Management Fundamentals

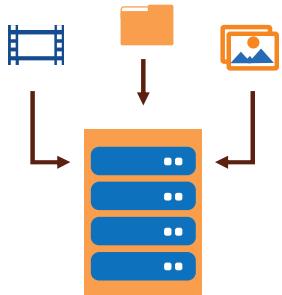


- **Definition:** The process of collecting, storing, organizing and maintaining data efficiently and securely throughout its lifecycle
- **Goals:** Ensure data quality, accessibility, compliance, and integrity for effective use in machine learning projects



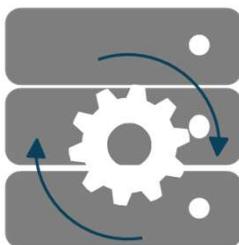
## Data Collection

- **Sources:** APIs, databases, web scraping, sensors
- **Considerations:** Data relevance, volume, and quality



## Data Processing

- **Task:** Cleaning, transformation, feature engineering
- **Goal:** Prepare data for model training



## Data Storage

- **Options:** Cloud storage, on-premises storage
- **Factors:** Scalability, cost, and security



## Data Archiving

- **Purpose:** Long-term storage of data no longer in active use
- **Strategies:** Cold storage, backup solutions



- **Definition:** Manages changes to data over time
- **Benefits:** Ensures reproducibility and traceability

## Importance of Data Versioning

## Techniques for Data Versioning

**Version control system:**  
Implement systems like Git for code and DVC for data

**Data lineage tracking:**  
Record the flow and transformations of data from source to model input

**Data partitioning:**  
Segment data into meaningful versions based on time, source, or feature set

## Data Governance



- **Definition:** Framework for managing data quality and security
- **Importance:** Ensures data integrity, compliance, and ethical use

## Compliance and Regulations



- **Regulations:** GDPR, CCPA, HIPPA
- **Focus:** Data privacy, user consent, security

## Data Management Best Practices

- **Practices:** Data quality checks, documentation, access control
- **Goal:** Ensure reliable, secure, and high-quality data

## Data Storage Solutions

- **Cloud storage:** Scalable and cost-effective (e.g., AWS S3)
- **On-premises storage:** Secure and controlled local storage
- **Hybrid storage:** Combines cloud and on-premises benefits
- **Data lakes:** Stores raw, large-scale data for analytics

## Cloud vs. On-Premises Storage

- Flexible
- Scalable
- Remote access
- Low hardware cost
- Efficient backup and recovery
- Security compliance

Cloud-based



On-premises

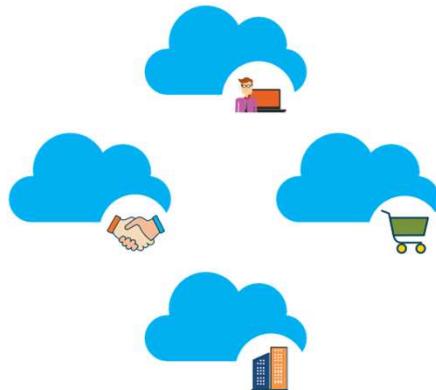


vs.

- Controlled
- Secure
- Cost-effective for large datasets
- High hardware and maintenance costs
- Medium backup cost

## Selecting the Right Storage Solution

- **Considerations:** Data size, accessibility, cost, security
- **Recommendations:** Match storage solutions to project requirements and scalability



## Ensuring Data Quality and Integrity

- **Data quality checks:** Implement validation rules, handle missing values, and remove duplicates
- **Data integrity:** Maintain data consistency and accuracy through robust data management practices
- **Tools:** Great Expectations, Pandas Profiling



## Knowledge Check

6. Data versioning is primarily used to track changes in code.
  - A. True
  - B. False

## Knowledge Check

7. Which of the following tools is commonly used for data versioning in ML projects?
- A. GitHub
  - B. Docker
  - C. DVC
  - D. Kubernetes



## Discussion

### Importance of EDA and Preprocessing in ML projects

- Exploratory data analysis
- Data cleaning
- Feature engineering
- Preferred tools

## Lab 2: EDA, Feature Engineering, and Data Cleaning

### Instructions

1. Perform exploratory data analysis using pandas
2. Visualize data distributions using matplotlib and seaborn (histograms, scatter plots, box plots, pair plots, etc.)
3. Create new features for better model performance
4. Clean the dataset by handling missing values and removing outliers

Estimated completion time: 30 minutes



## Topic 6: Feature Stores

- What is a feature store?
- Importance of feature stores in MLOps
- Key benefits of using feature stores
- Overview of feature store types
- Online feature stores
- Offline feature stores
- Feature extraction and transformation
- Feature storage and serving

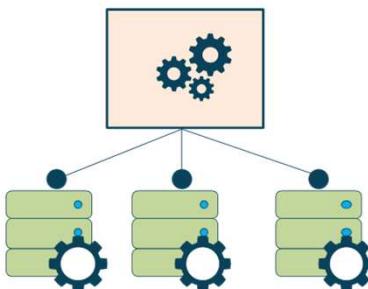


## Topic 6: Feature Stores

- Data consistency across training and serving
- Ensuring feature versioning
- Reuse and sharing of features
- Monitoring and managing features
- Technical challenges
- Data quality and governance
- Operational challenges

### What Is a Feature Store?

**Definition:** A centralized repository for storing, managing, and serving features for machine learning models



## Importance of Feature Stores in MLOps



- **Consistency:** Ensures the same features are used in training and serving, reducing data drift
- **Scalability:** Supports efficient reuse of features across different models and projects
- **Collaboration:** Enables teams to share and access features easily, improving productivity

## Key Benefits of Using Feature Stores



- **Centralized feature management:** All features are stored in one place, reducing redundancy
- **Improved efficiency:** Pre-computed features speed up model training and deployment
- **Versioning and lineage:** Track changes and maintain a history of feature versions

## Overview of the Feature Store Types



- **Online feature stores:** Used for real-time feature serving with low latency
- **Offline feature stores:** Used for batch processing and model training
- **Use cases:** Choose the right type based on the use case (e.g., real-time vs. batch predictions)

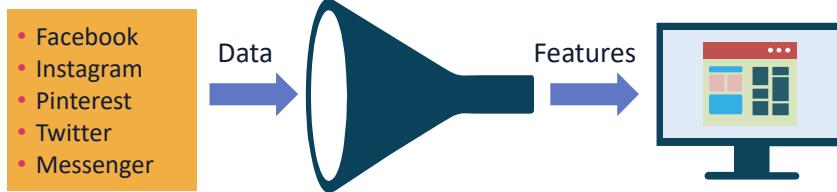
- **Definition:** Feature stores that serve features in real-time with low latency
- **Use cases:** Suitable for applications requiring real-time predictions, like fraud detection or recommendation systems
- **Benefits:** Reduces response time for predictions and supports dynamic features

### Online Feature Stores

## Offline Feature Stores

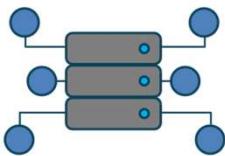
- **Definition:** Feature stores used for batch processing, model training, and historical data analysis
- **Use cases:** Suitable for periodic model retraining, historical data analysis, and batch predictions
- **Benefits:** Supports large-scale data processing and ensures consistency in model training

- **Feature extraction:** Process of converting raw data into meaningful features
- **Feature transformation:** Modify features to improve model performance (e.g., scaling, encoding)
- **Tools:** Pandas, Spark, Scikit-learn for transforming features



## Feature Extraction and Transformation

## Feature Storage and Serving



- **Storage:** How features are stored in the feature store (e.g., key-value stores, databases)
- **Serving:** Retrieval of features for both training and real-time predictions
- **Efficiency:** Ensures fast access to features, reducing prediction latency



## Data Consistency Across Training and Serving

- **Training-serving skew:** Differences between features used during training and serving can degrade model performance
- **Solutions:** Ensure that the same features and transformations are applied in both environments
- **Tools:** Feature stores help maintain consistency and avoid training-serving skew

## Ensuring Feature Versioning



- **Definition:** Keeps track of different versions of features as they evolve over time
- **Importance:** Ensures reproducibility and traceability of experiments
- **Tools:** Uses version control systems to manage and track feature changes

## Reuse and Sharing of Features

- **Reusable features:** Design features that can be used across multiple models
- **Collaborative engineering:** Share features within teams to reduce duplication and improve efficiency
- **Centralized management:** Use feature stores to centralize feature access and management



## Monitoring and Managing Features



- **Monitoring:** Track feature performance and quality over time
- **Management:** Organize features based on projects, teams, and usage
- **Tools:** Use monitoring tools to alert on data drift and feature degradation

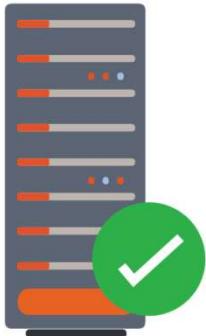
**Scalability:**  
Handling large-scale data efficiently

**Real-time serving:**  
Ensuring low-latency feature serving for online applications

**Integration:**  
Integrating feature stores with existing data infrastructure

## Technical Challenges

## Data Quality and Governance



- **Data quality:** Ensure the accuracy and completeness of features
- **Governance:** Establish policies for data access, usage, and compliance
- **Tools:** Use governance frameworks and tools to maintain data quality and integrity

## Operational Challenges



## Knowledge Check

8. Which of the following is a primary benefit of using feature stores in MLOps?
  - A. Reducing model training time by avoiding data preprocessing
  - B. Ensuring consistency of features across training and serving
  - C. Eliminating the need for version control in ML projects
  - D. Automating the model deployment process

## Knowledge Check

9. What is the main purpose of feature versioning in the feature store?
  - A. Improve the speed of model prediction
  - B. Track changes in features and maintain reproducibility
  - C. Automate data preprocessing and transformation
  - D. Eliminate the need for feature engineering



## Topic 7: Model Development

- Introduction to model development
- Step in developing ML models
- Data preparations and preprocessing
- Factors to consider when choosing algorithm
- Overview of common algorithms
- Selecting the best algorithm
- Understanding the training process
- Importance of validation
- Cross validation techniques



## Topic 7: Model Development

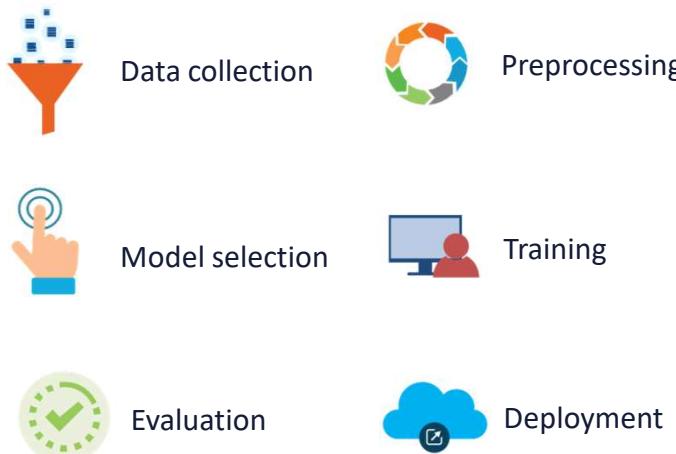
- What is overfitting?
- Techniques to prevent overfitting
- Balancing bias and variance
- Introduction to evaluation metrics
- Classification metrics
- Regression metrics

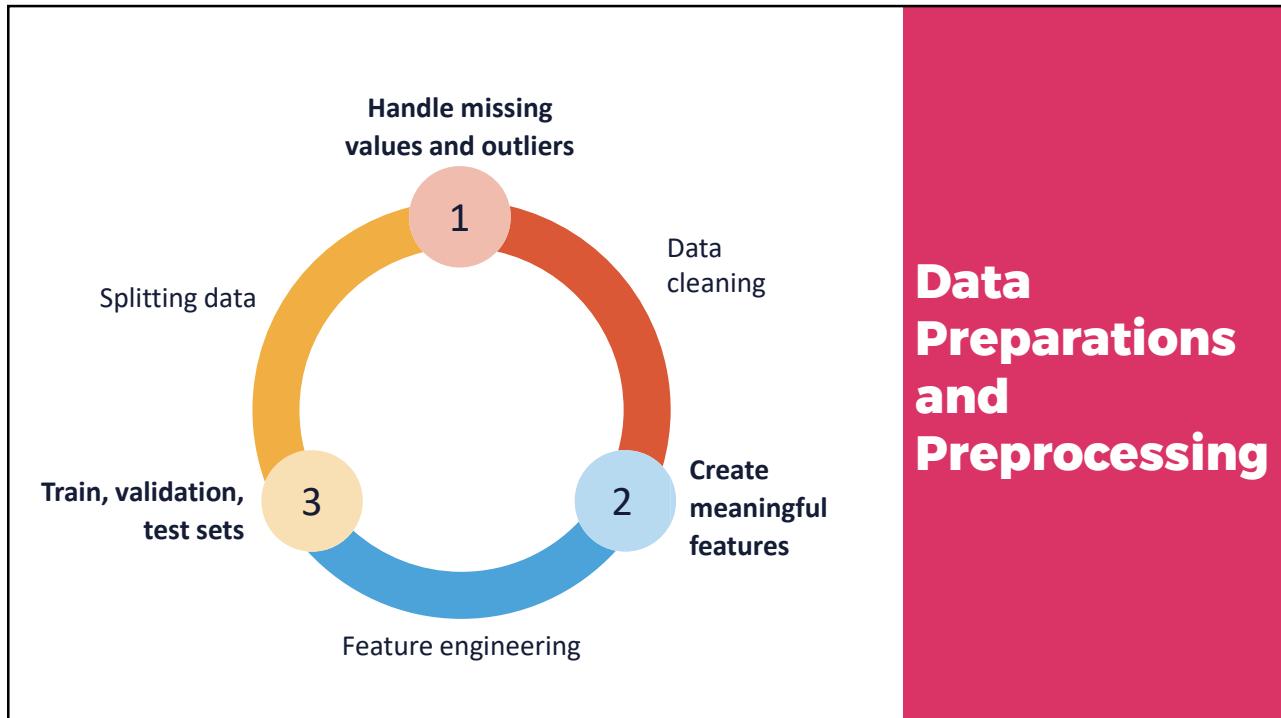
- **Definition:** The process of building, training, and validating ML models
- **Lifecycle stages:** Data preparation, model selection, training, validation, evaluation, and deployment



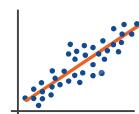
## Introduction to Model Development

### Steps in Developing ML Models

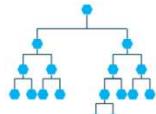




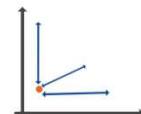
## Overview of Common Algorithms



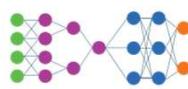
Linear regression



Decision tree



Support vector  
machine



Neural networks



K-nearest neighbor

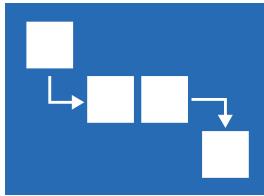
## Selecting the Best Algorithm

- **Problem type:** Classification, regression, clustering
- **Trade-offs:** Accuracy vs. interpretability



## Understanding the Training Process

- **Training data:** Learning from labelled examples
- **Loss function:** Measures prediction error



- **Purpose:** Prevent overfitting
- **Techniques:** Train-test split, cross-validation



**Importance  
of Validation**

## Cross-Validation Techniques

1. K-fold cross-validation
2. Leave-one-out cross-validation
3. Stratified cross-validation
4. Time series cross-validation

- **Definition:** Model fits noise in training data
- **Symptoms:** High training accuracy, low test accuracy



**What Is Overfitting?**

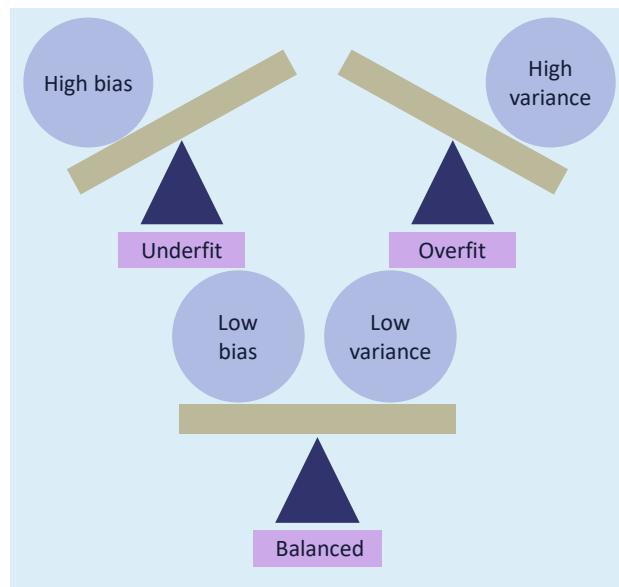
## Techniques to Prevent Overfitting

- Regularization
- Early stopping
- Dropout
- Data augmentation



## Balancing Bias and Variance

**Trade-off:** Optimal balance for generalization

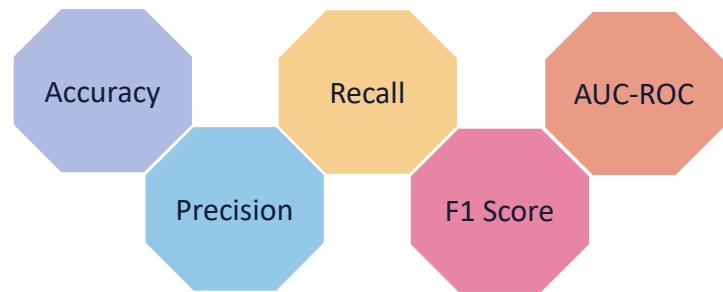


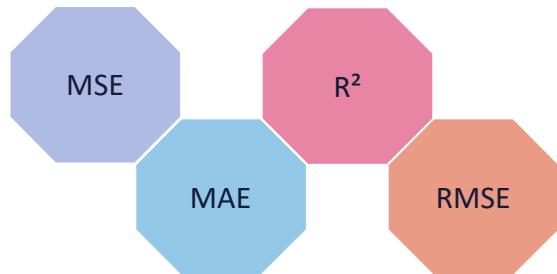
**Purpose:** Quantify model performance



## Introduction to Evaluation Metrics

### Classification Metrics





## Regression Metrics

### Knowledge Check

10. Cross-validation is used to prevent overfitting by ensuring that the model performs well on different subsets of the training data.
- A. True
  - B. False

## Knowledge Check

11. Which of the following metrics is most suitable for evaluating the performance of a regression model?
- A. Precision
  - B. Mean Squared Error (MSE)
  - C. Accuracy
  - D. F1 Score



## Topic 8: Implementing a Basic ML Pipeline

- Introduction to ML pipeline
- Components of ML pipeline
- Setting up the environment
- Defining the pipeline structure
- Step-by-step guide to building a pipeline
- Incorporating data preprocessing steps
- Handling missing data



## Topic 8: Implementing a Basic ML Pipeline

- Feature engineering in pipeline
- Model development within the pipeline
- Hyperparameter tuning
- Splitting data for training and testing
- Training the pipeline
- Evaluating pipeline performance
- Cross-validation of the pipeline

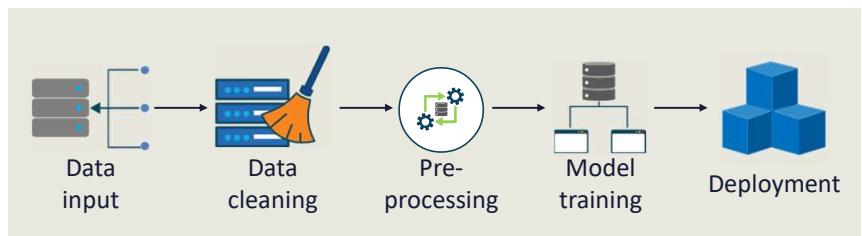


## Topic 8: Implementing a Basic ML Pipeline

- Handling imbalanced data in pipeline
- Basics of pipeline automation
- Continuous integration and development (CI/CD)
- Automating training pipelines
- Model monitoring and logging
- Best practices for pipeline automation

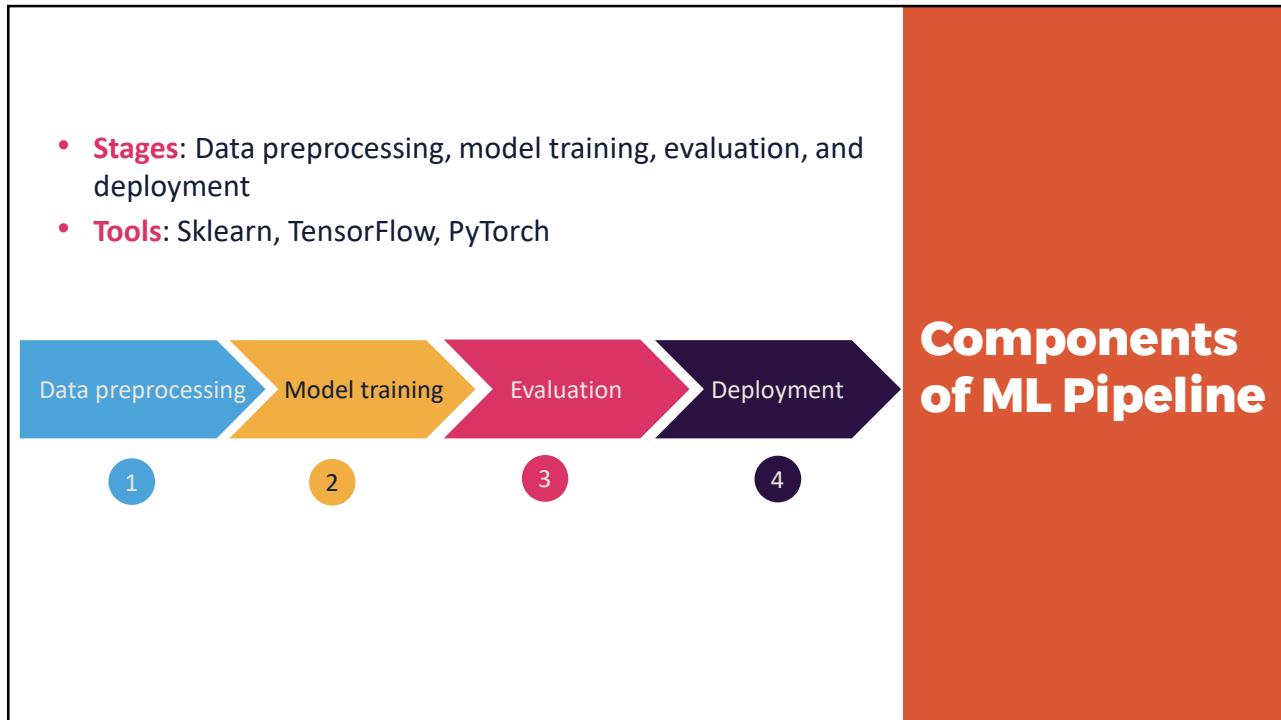
## Introduction to ML Pipelines

- **Definition:** A sequence of steps for data processing and model training
- **Importance:** Ensures repeatability, scalability, and efficiency in ML workflows



## Setting up the Environment

- **Tools required:** Python, Sklearn, Pandas, NumPy
- **Basic steps:** Import necessary libraries



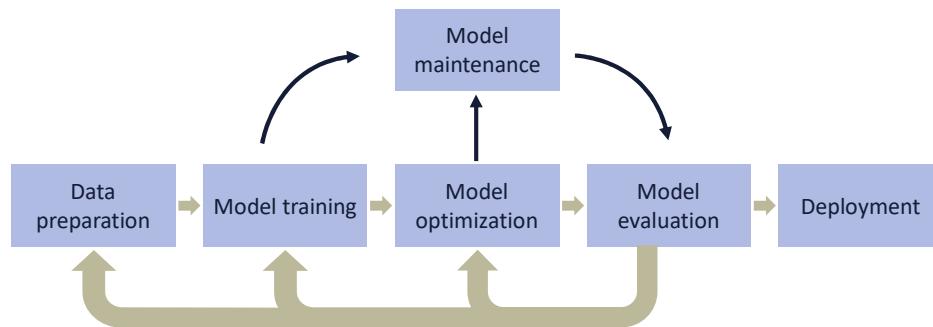
## Defining the Pipeline

### Initialization

- **Sklearn pipeline:** Use Sklearn's pipeline class to define steps
- **Components:** Include steps like StandardScalar PCA, LogisticRegression

## Step by Step Guide to Building a Pipeline

**Create a pipeline:** Use Sklearn to define a simple pipeline



## Incorporating Data Preprocessing Steps

- **Preprocessing components:** Scaling, encoding, feature selection
- **Sklearn preprocessing:** Using StandardScaler, OneHotEncoder



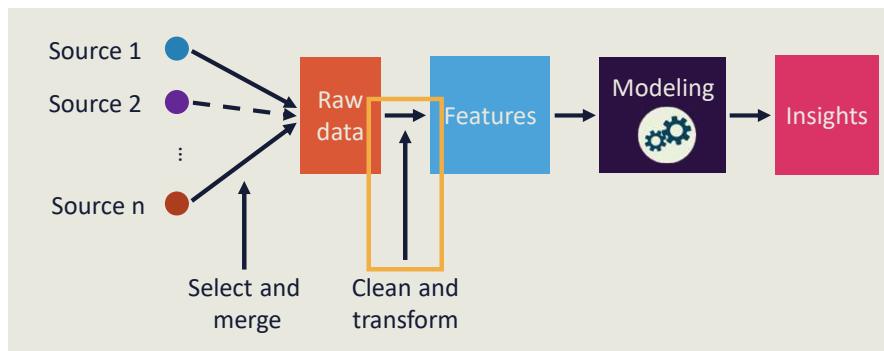
## Handling Missing Data

- **Strategies:** Imputation, removal
- **Sklearn imputer:** Using SimpleImputer to handle missing data



## Feature Engineering in Pipeline

- **Adding features:** Create new features within the pipeline
- **FeatureUnion:** Combine multiple feature processing steps



## Model Development within Pipeline

- **Adding models:** Integrate models like LogisticRegression, SVM
- **Code example:** As presented on the right side

```

1 from sklearn.pipeline import Pipeline
2 from sklearn.preprocessing import StandardScaler
3 from sklearn.linear_model import LogisticRegression
4 from sklearn.svm import SVC
5 from sklearn.datasets import load_iris
6 from sklearn.model_selection import train_test_split
7 # Load dataset and split into train and test sets
8 X, y = load_iris(return_X_y=True)
9 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)
10 # Create a pipeline with preprocessing and a model (Logistic Regression)
11 pipeline = Pipeline([
12     ('scaler', StandardScaler()),
13     ('classifier', LogisticRegression()) ])
14 # Train and predict with Logistic Regression
15 pipeline.fit(X_train, y_train) print("Logistic Regression predictions:", pipeline.predict(X_test))
16 # Update the classifier to SVM and retrain
17 pipeline.set_params(classifier=SVC(kernel='linear')).fit(X_train, y_train)
18 print("SVM predictions:", pipeline.predict(X_test))

```

## Hyperparameter Tuning

- **Grid search:** Using `GridSearchCV` with pipelines
- **Code example:** As presented on the right side

```

1 from sklearn.pipeline import Pipeline
2 from sklearn.preprocessing import StandardScaler
3 from sklearn.svm import SVC
4 from sklearn.datasets import load_iris
5 from sklearn.model_selection import GridSearchCV, train_test_split
6 # Load dataset and split into train and test sets
7 X, y = load_iris(return_X_y=True)
8 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)
9 # Create a pipeline with preprocessing and a model (SVM)
10 pipeline = Pipeline([
11     ('scaler', StandardScaler()),
12     ('svc', SVC()) ])
13 # Define the hyperparameters to be tuned
14 param_grid = {
15     'svc__C': [0.1, 1, 10], # Regularization parameter
16     'svc__kernel': ['linear', 'rbf'] # Kernel type
17 }
18 # Use GridSearchCV to find the best combination of parameters
19 grid_search = GridSearchCV(pipeline, param_grid, cv=5)
20 grid_search.fit(X_train, y_train)
21 # Output the best parameters
22 print("Best parameters found:", grid_search.best_params_)

```

## Code Explanation

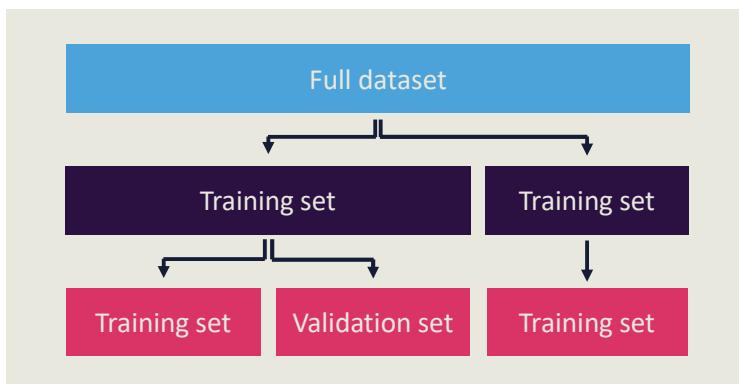
- **Pipeline setup:** The pipeline includes **StandardScaler** for data scaling and SVC as the model
- **Hyperparameter grid:** A dictionary (**param\_grid**) is created to specify the hyperparameters of SVC to be tuned
- **GridSearchCV:** **GridSearchCV** searches for the specified hyperparameters using cross-validation
- **Best parameters:** After fitting, the best parameters found are printed
- This example shows how to efficiently integrate hyperparameter tuning into a pipeline using **GridSearchCV**

## Knowledge Check

12. A Sklearn pipeline can include both data preprocessing steps and model training steps in a single workflow.
- True
  - False

## Train-Test Split

- **Train-test split:** Use `train_test_split` from Sklearn
- **Importance:** Ensures fair evaluation of the pipeline



## Splitting Data for Training and Testing

## Training the Pipeline

- **Fit method:** Train the pipeline using the `fit` method
- The pipeline consists of **StandardScaler** for scaling and **LogisticRegression** for classification
- The `fit` method trains the pipeline, applying all steps sequentially to the training data

```
1 from sklearn.pipeline import Pipeline
2 from sklearn.preprocessing import StandardScaler
3 from sklearn.svm import SVC
4 from sklearn.datasets import load_iris
5 from sklearn.model_selection import GridSearchCV, train_test_split
6 # Load dataset and split into train and test sets
7 X, y = load_iris(return_X_y=True)
8 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)
9 # Create a simple pipeline
10 pipeline = Pipeline([
11     ('scaler', StandardScaler()),
12     ('classifier', LogisticRegression())
13 ]) # Train the pipeline using the fit method
14 pipeline.fit(X_train, y_train)
```

## Evaluating the Pipeline

- **Metrics:** Use **accuracy**, **precision**, **recall**, and **F1 score**
- **Pipeline training:** The pipeline is trained using **fit** on the training data
- **Evaluation:** Predictions are made of the test training data, and evaluation metrics like accuracy, precision, recall and F1 score are calculated and printed

```

● ● ●
1 from sklearn.pipeline import Pipeline
2 from sklearn.preprocessing import StandardScaler
3 from sklearn.linear_model import LogisticRegression
4 from sklearn.datasets import load_iris
5 from sklearn.model_selection import train_test_split
6 from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score
7 # Load dataset and split into train and test sets
8 X, y = load_iris(return_X_y=True)
9 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)
10 # Create a simple pipeline
11 pipeline = Pipeline([
12     ('scaler', StandardScaler()),
13     ('classifier', LogisticRegression())
14 ]) # Train the pipeline using the fit method
15 pipeline.fit(X_train, y_train)
16 # Make predictions on test data
17 y_pred = pipeline.predict(X_test)
18 # Evaluate the pipeline performance
19 accuracy = accuracy_score(y_test, y_pred)
20 precision = precision_score(y_test, y_pred, average='macro')
21 recall = recall_score(y_test, y_pred, average='macro')
22 f1 = f1_score(y_test, y_pred, average='macro')
23 #Print evaluation metrics
24 print(f'Accuracy: {accuracy:.2f}, Precision: {precision:.2f}, Recall: {recall:.2f}, F1 Score: {f1:.2f}')

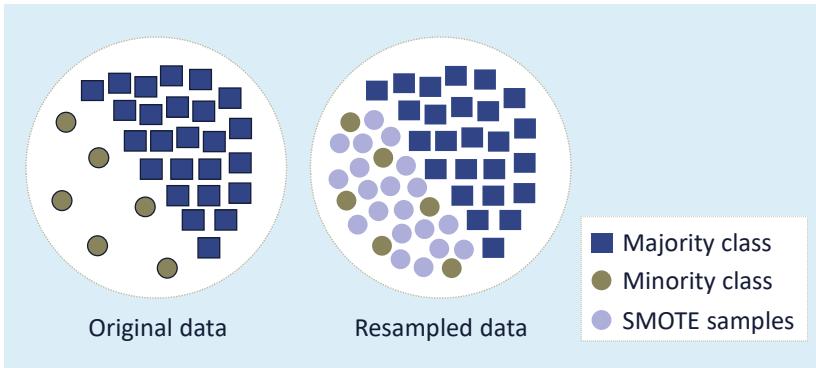
```

## Validating the Pipeline

- **Cross-validation:** Use **cross\_val\_score** for robust evaluation
- **Benefits:** Provides more accurate performance estimates



## Handling Imbalanced Data in Pipeline

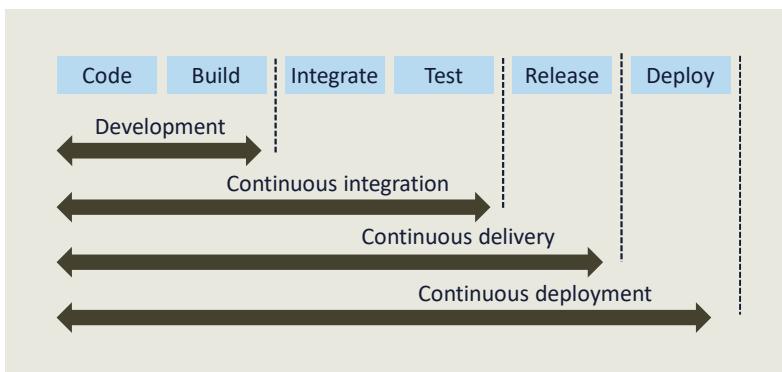


## Basics of Pipeline Automation

- **Definition:** Automating repetitive tasks in the ML lifecycle
- **Importance:** Increase efficiency and reduce human error



**Tools:** Jenkins, GitHub Actions, Azure DevOps



## Continuous Integration and Development

## Automating A Training Pipeline

- **Scheduling training:** Use tools like Airflow for scheduling
- **DAG definition:** The Directed Acyclic Graph (DAG) defines the workflow for automating the training process
- **Data loading and model training:** The `train_model` function loads data, creates a pipeline, trains a model, and saves it

```

1  from airflow import DAG
2  from airflow.operators.python_operator import PythonOperator
3  from datetime import datetime, timedelta
4  from sklearn.datasets import load_iris
5  from sklearn.model_selection import train_test_split
6  from sklearn.linear_model import LogisticRegression
7  from sklearn.pipeline import Pipeline
8  from sklearn.preprocessing import StandardScaler
9  import joblib
10 # Function to load and preprocess data
11 def load_data():
12     X, y = load_iris(return_X_y=True)
13     X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)
14     return X_train, X_test, y_train, y_test
15 # Function to train and save the model
16 def train_model():
17     X_train, X_test, y_train, y_test = load_data()
18     pipeline = Pipeline([
19         ('scaler', StandardScaler()),
20         ('classifier', LogisticRegression())
21     ])
    
```

## Automating A Training Pipeline

**Scheduling:** The DAG is scheduled to run daily using the  
`'schedule_interval=@daily'`

```
1     pipeline.fit(X_train, y_train)
2     joblib.dump(pipeline, '/path/to/save/model.pkl') # Save the trained model
3 # Define the default arguments for the DAG
4 default_args = {
5     'owner': 'airflow',
6     'start_date': datetime(2023, 9, 20),
7     'retries': 1,
8     'retry_delay': timedelta(minutes=5)
9 # Define the DAG
10 dag = DAG(
11     'train_model_pipeline',
12     default_args=default_args,
13     schedule_interval= '@daily', # Schedule to run daily
14     catchup=False
15 )
16 # Define the training task
17 train_task = PythonOperator(
18     task_id='train_model',
19     python_callable=train_model,
20     dag=dag
21 ) # Set the task dependencies train_task
22 train_task
```

## Model Monitoring and Logging

- **Monitoring:** Track model performance over time
- **Logging:** Record metrics and model outputs for analysis



## Best Practices for Pipeline Automation



- **Tips:** Version control, modular code, clear documentation
- **Resources:** Links to automation tools and documentation

### Knowledge Check

13. Which of the following tools is commonly used for automating the scheduling of training pipelines in machine learning?
- A. GitHub Actions
  - B. Jenkins
  - C. Apache Airflow
  - D. Dockers

## Lab 3: Building and Automating a Basic ML Pipeline

### Instructions

1. Automate data preprocessing and feature selection
2. Train and evaluate an ML model
3. Use metrics to evaluate model performance
4. Save the model using joblib

Estimated completion time: 30 minutes



## Reflection Activity

### Lesson Learned from Demos and Lab

- What did you find most valuable in today's activities?
- How will you apply these techniques to real-world projects?
- What challenges did you encounter, and how did you solve them?

## Challenge Lab 1: Create an End-to-End MLOps Pipeline – From Preprocessing to GitHub Automation with the Iris Dataset

### Instructions

1. Set up an ML project with GitHub
2. Perform data exploration, preprocessing, and feature engineering
3. Automate the ML pipeline using GitHub Actions
4. Train and evaluate an ML model within the Pipeline

Estimated completion time: 45 minutes

## Agenda Review

### Module 1: MLOps and Data

- ✓ Introduction to MLOps
- ✓ Machine Learning Lifecycle Overview
- ✓ MLOps Components and Tools
- ✓ Setting Up an ML Project
- ✓ Data Management Fundamentals
- ✓ Feature Stores
- ✓ Model Development
- ✓ Implementing a Basic ML Pipeline

### Module 2: Development, Experimentation, and Evaluation

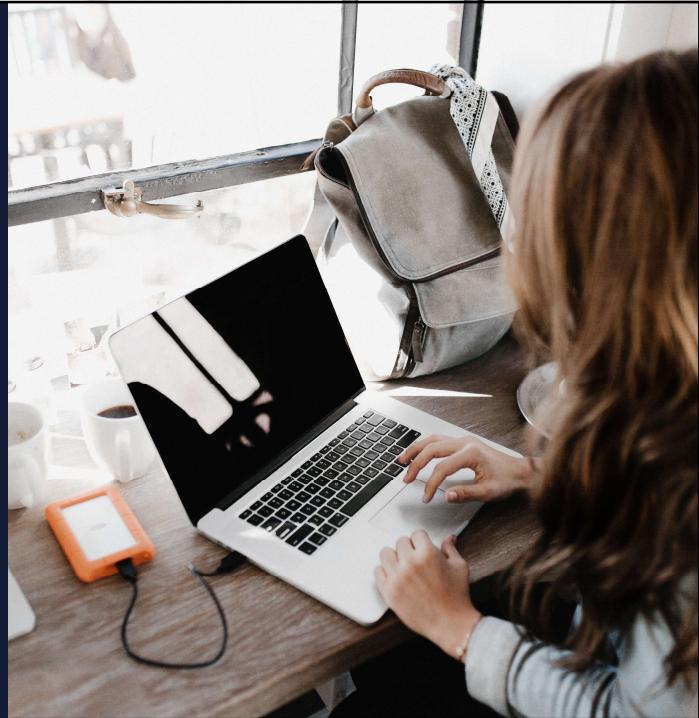
### Module 3: Deployment, Monitoring, and Pipelines



## Module 2: Development, Experimentation and Evaluation

After today's module, you will be able to

- Develop a comprehensive understanding of model development and experiment tracking
- Master techniques for model interpretability and explainability



## Module 2: Development, Experimentation and Evaluation

After today's module, you will be able to

- Implement and evaluate machine learning algorithms with key performance metrics
- Optimize model performance through effective hyperparameter tuning



## Module Agenda

### Module 1: MLOps and Data

### Module 2: Development, Experimentation, and Evaluation

- Model Development Strategies
- ML Model Interpretability and Explainability
- Implementing Algorithms
- Experiment Tracking and Model Evaluation
- Setting Up MLflow for Experiment Tracking
- Evaluating Models
- Hyperparameter Tuning Techniques
- Automated Hyperparameter Tuning

### Module 3: Deployment, Monitoring, and Pipelines



## Topic 1: Model Development Strategies

- Traditional vs. modern approaches to model development
- Iterative nature of model development
- Role of MLOps in modern model development
- Data-centric approach
- Model-centric approach
- Comparing data-centric and model centric approaches
- Importance of experimentation in model development
- Running controlled experiments



## Topic 1: Model Development Strategies

- Using cross-validation for experimentation
- Grid search for hyperparameter tuning
- Role of version control in model development
- Collaborative tools in MLOps
- Importance of documentation in collaborative development
- Best practices for collaborative model development
- Future trends in collaborative model development

## Traditional vs. Modern Approaches to Model Development

### Traditional approach



● Static dataset



● Manual feature engineering



● Linear development

### Modern approach



● Iterative process



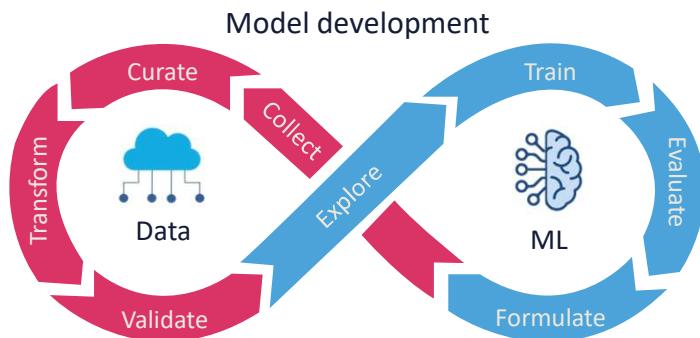
● Automated pipeline



● Rapid development and deployment

## Iterative Nature of Model Development in MLOps

- **Continuous improvement:** models are refined through multiple iterations
- **Feedback loop:** monitoring and evaluation feed back into the development cycle



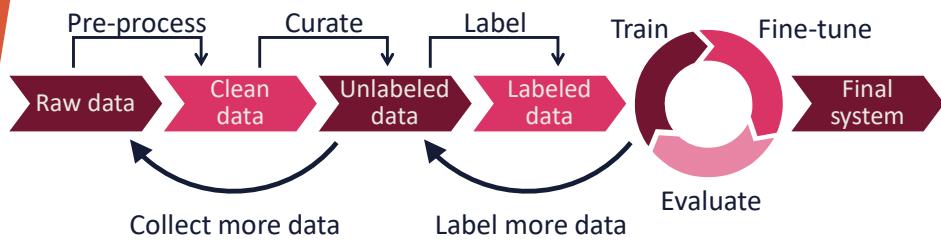
- **Automation:** automates data pipelines, model training, and deployment
- **Scalability:** supports scalable model development and deployment



## Role of MLOps in Modern Model Development

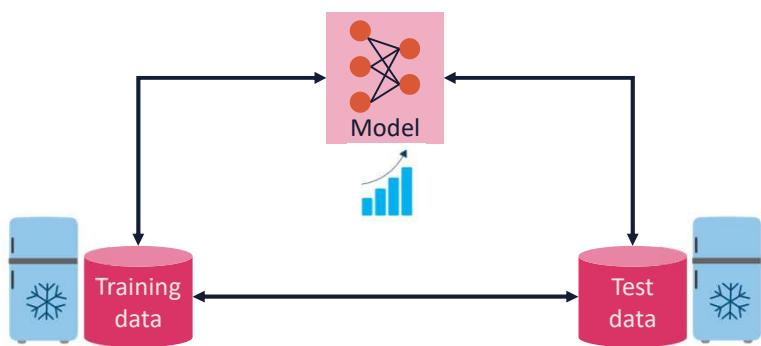
## Data-Centric Approach

- **Definition:** focuses on improving data quality and quantity
- **Key practices:** data cleaning, augmentation, and labeling

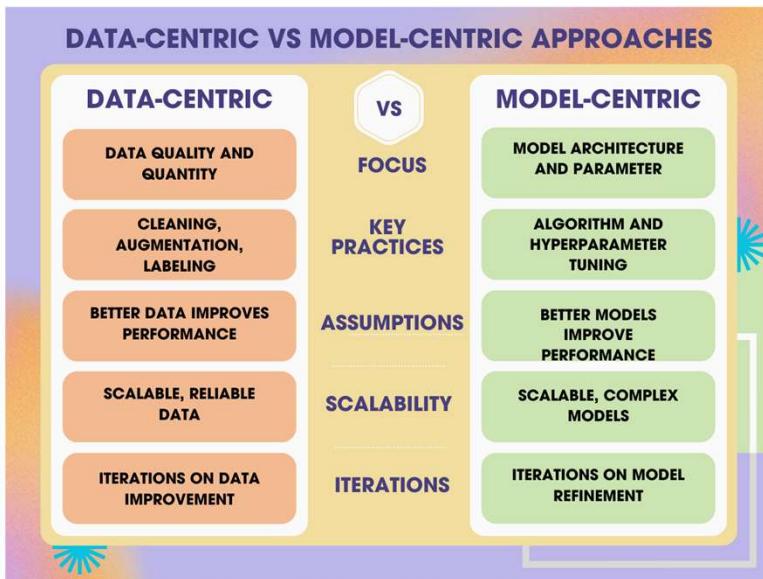


## Model-Centric Approach

- **Definition:** focuses on selecting the best algorithms and tuning model parameters
- **Key practices:** hyperparameter tuning, algorithm selection, and model architecture



## Comparison



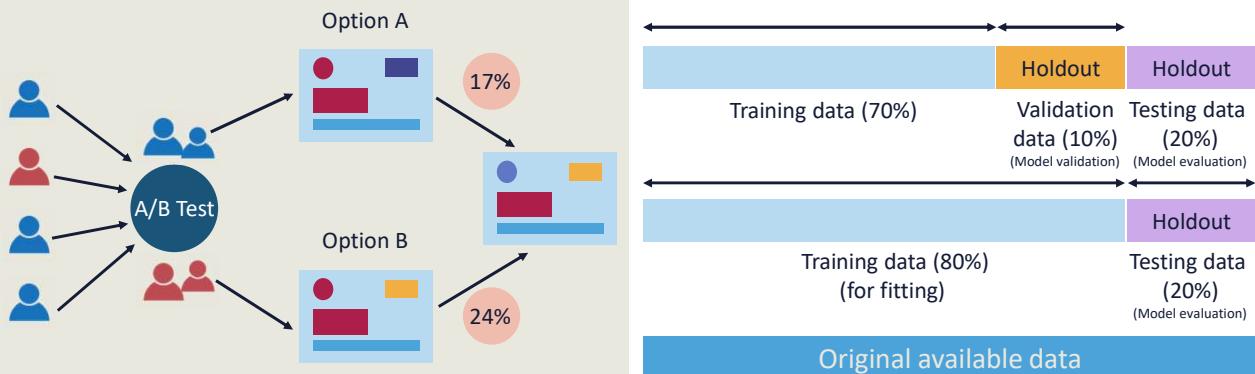
## Comparing Data-Centric and Model-Centric Approaches

## Importance of Experimentation in Model Development

- **Controlled experiments:** compare different models and parameters
- **Objective evaluation:** use metrics to assess model performance

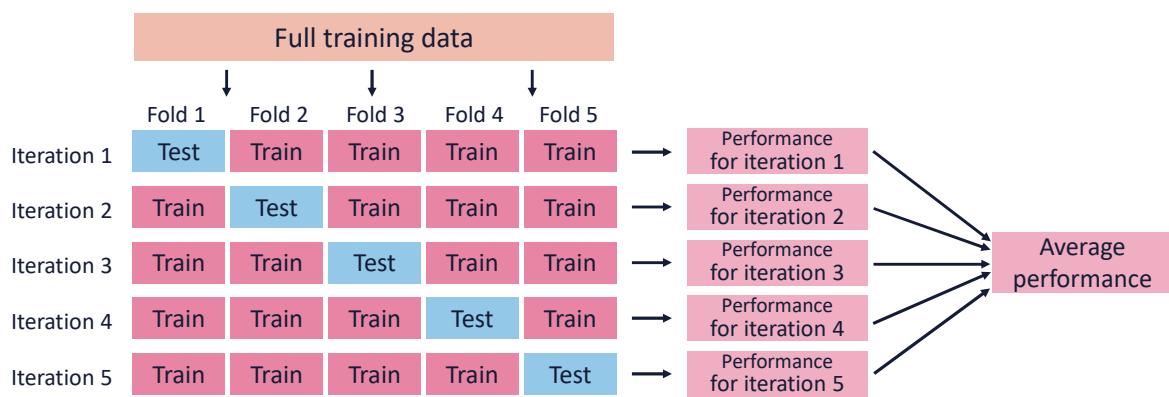
## Running Controlled Experiments

- **A/B testing:** compare two models on the same dataset



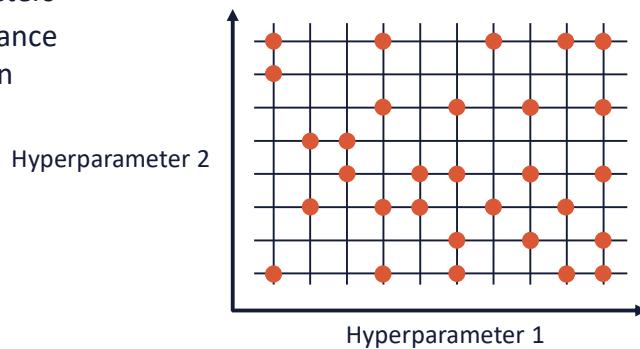
## Using Cross-Validation for Experimentation

- **Definition:** a technique for assessing model performance by splitting data into multiple folds
- **Benefits:** provides a more accurate measure of model generalization



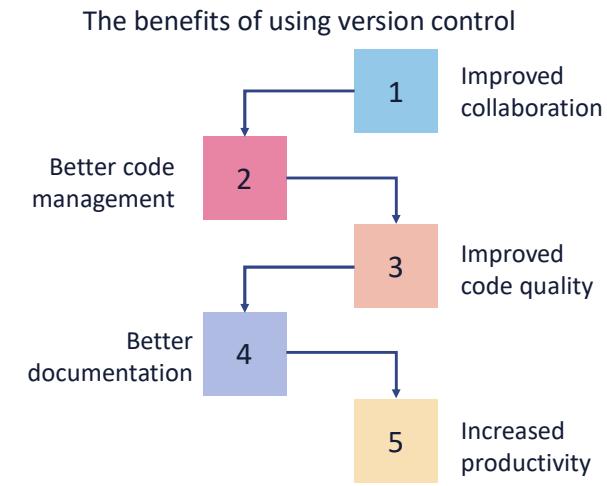
## Grid Search for Hyperparameter Tuning

- **Definition:** a method for exhaustively searching for the best hyperparameters
- **Benefits:** optimizes model performance through systematic experimentation



## Role of Version Control in Model Development

**Definition:** system for managing changes to code and models



## Collaborative Tools in MLOps

**Benefits:** streamline model development and deployment

Git

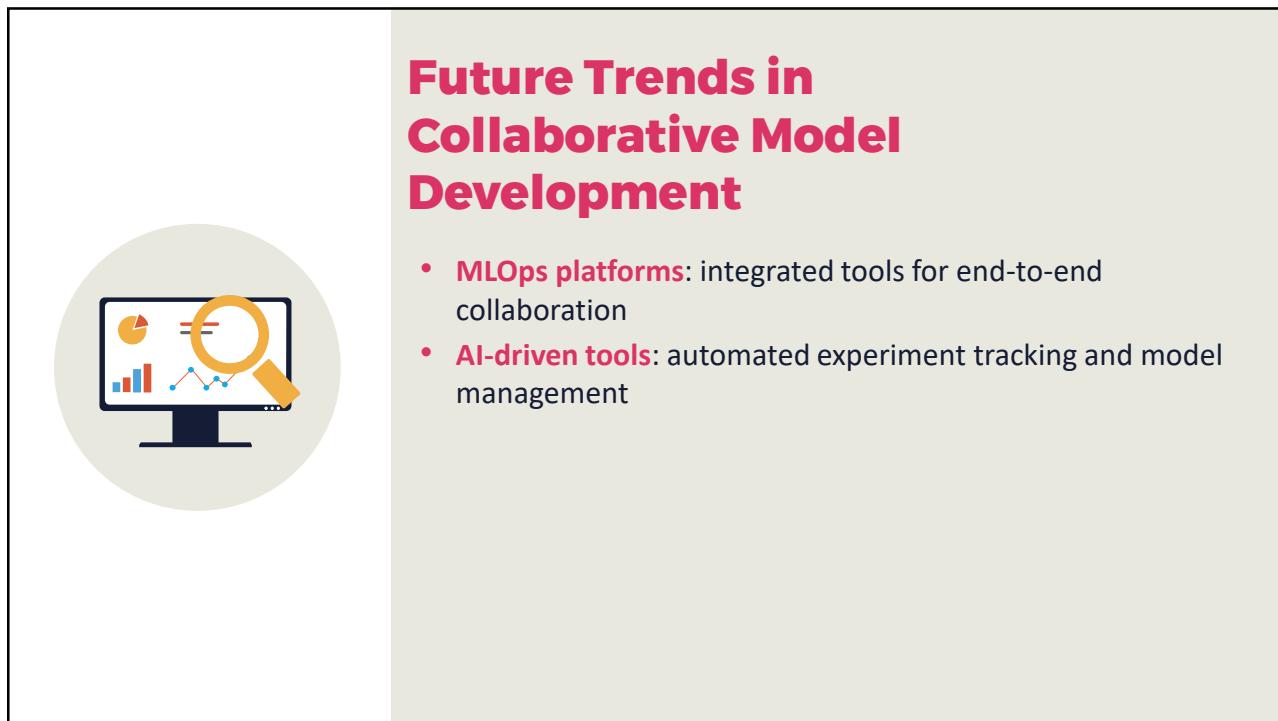
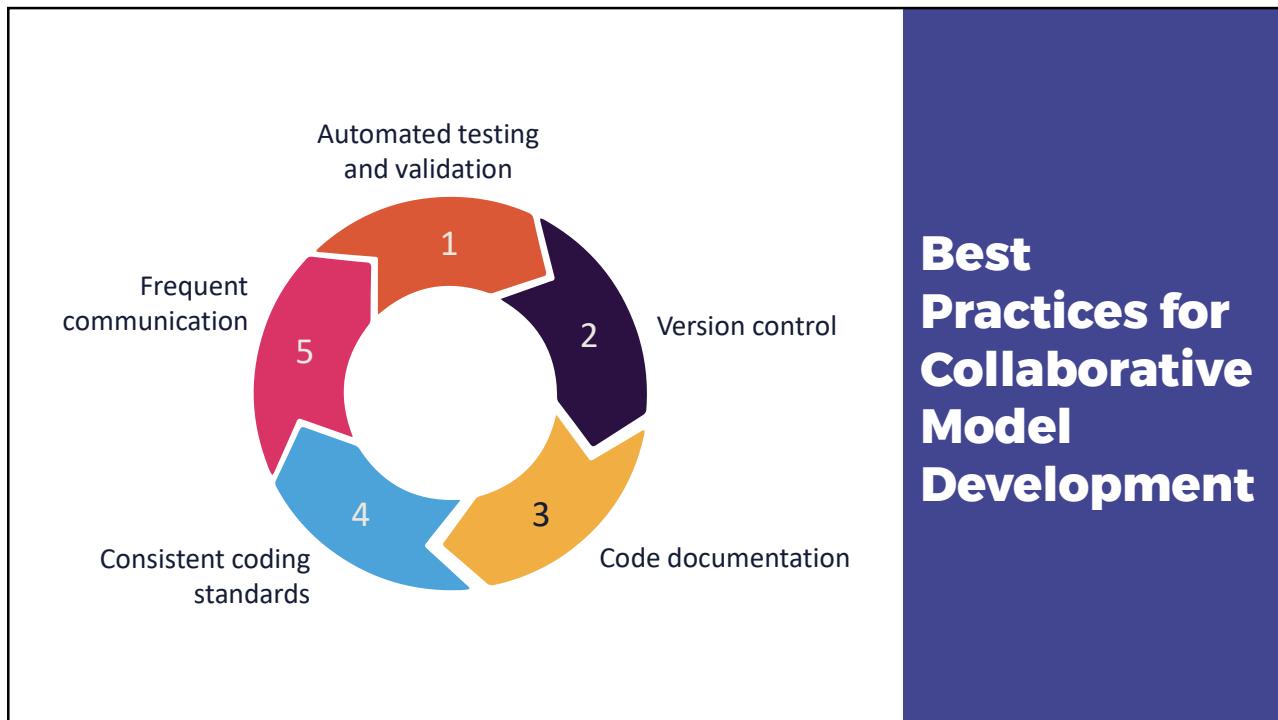
DVC

MLflow

## Importance of Documentation in Collaborative Development



- **Documentation:** captures details about experiments, data, and models
- **Benefits:** facilitates knowledge sharing and reproducibility



## Knowledge Check

1. Modern approaches in MLOps emphasize an iterative process with automated pipelines and continuous integration.
  - A. True
  - B. False

## Knowledge Check

2. What is the primary focus of the traditional approach to model development?
  - A. Continuous integration and deployment
  - B. Iterative experimentation
  - C. Static datasets and manual feature engineering
  - D. Automated pipelines



## Topic 2: ML Model Interpretability and Explainability

- Why model interpretability is crucial in MLOps
- Key concepts: global vs. local interpretability
- Regulatory requirements for explainability
- Feature importance
- SHAP (SHapley Additive exPlanations)
- LIME (Local Interpretable Model-agnostic Explanations)
- Explaining decision trees

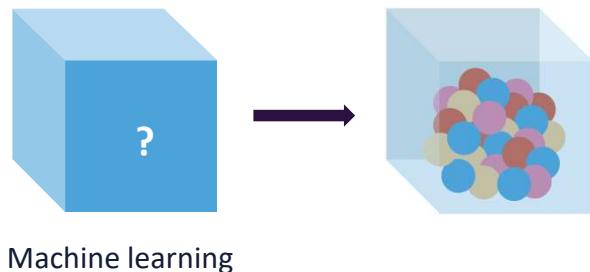


## Topic 2: ML Model Interpretability and Explainability

- Explaining deep learning models
- Comparing interpretability across model types
- SHAP library
- LIME library
- Integrated gradients
- Trade-offs between model complexity and interpretability
- Regulatory and ethical challenges
- Future directions in model interpretability

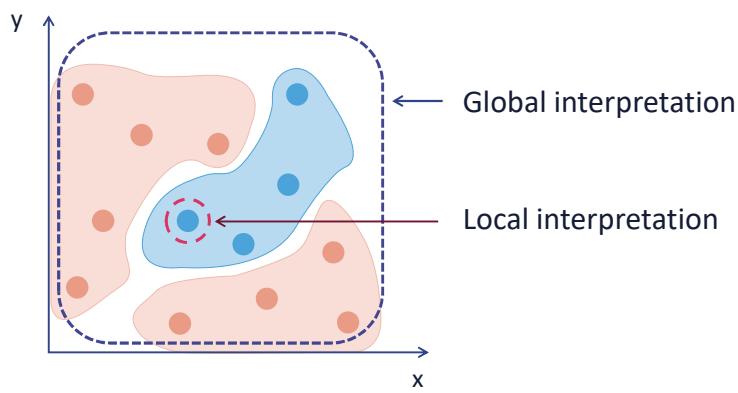
## Why Model Interpretability Is Important in MLOps

- **Definition:** ability to understand and trust model predictions
- **Importance:** ensures transparency, trust, and compliance in ML Models



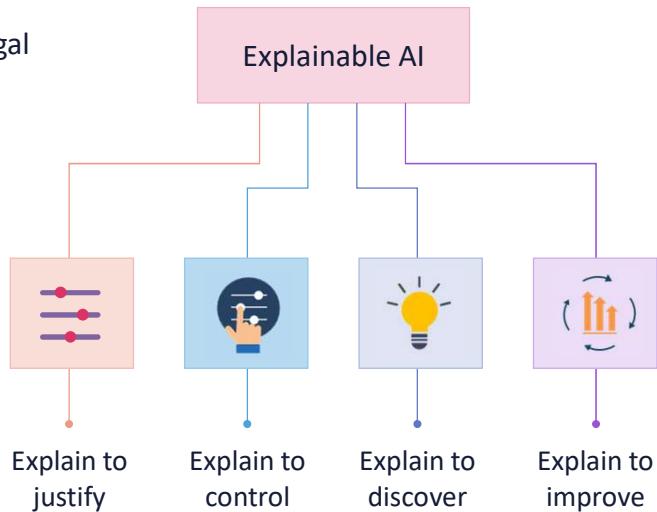
## Key Concepts: Global vs. Local Interpretability

- **Global interpretability:** understand overall model behavior
- **Local interpretability:** explain individual predictions



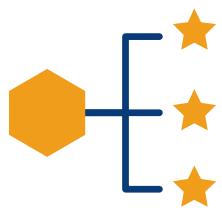
## Regulatory Requirements for Explainability

- **Industries:** finance, healthcare, and legal
- **Compliance:** ensuring models meet transparency and accountability standards



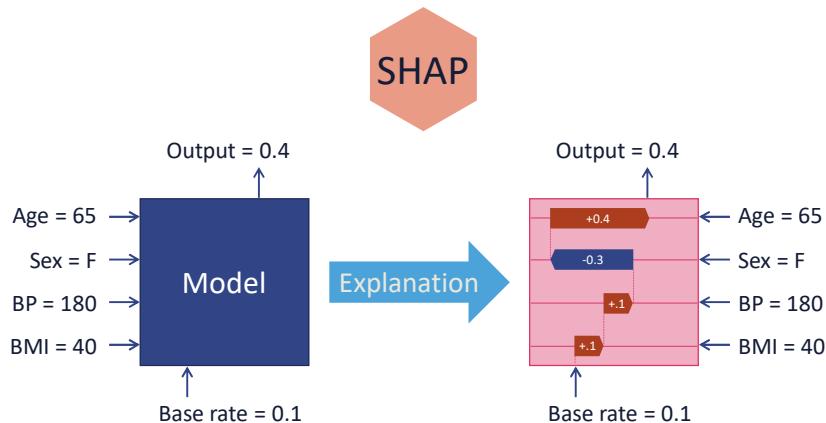
## Feature Importance

- **Definition:** measures the contribution of each feature to the model's predictions
- **Techniques:** permutation importance, Gini importance



## SHAP

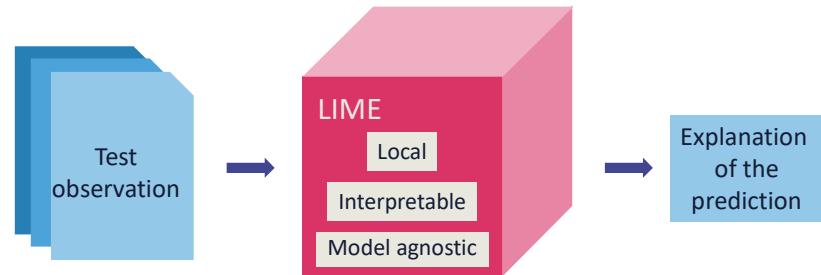
- **Definition:** method based on the Shapley values from game theory
- **Use:** provides both global and local interpretability



**SHAP  
(SHapley  
Additive  
exPlanations)**

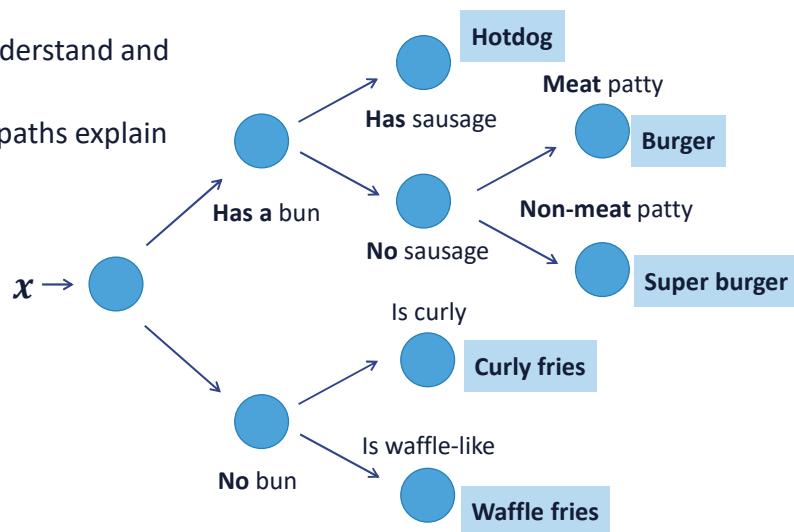
## LIME (Local Interpretable Model Explanations)

- **Definition:** perturbs data locally and fits interpretable models
- **Use:** explain individual predictions for any black-box model



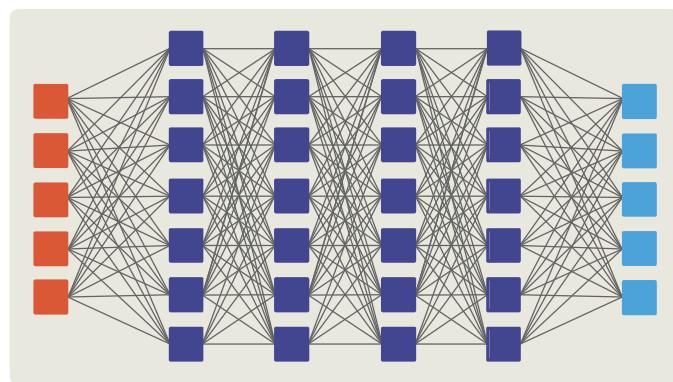
## Explaining Decision Trees

- **Tree structure:** easy to understand and visualize
- **Interpretability:** decision paths explain model decisions



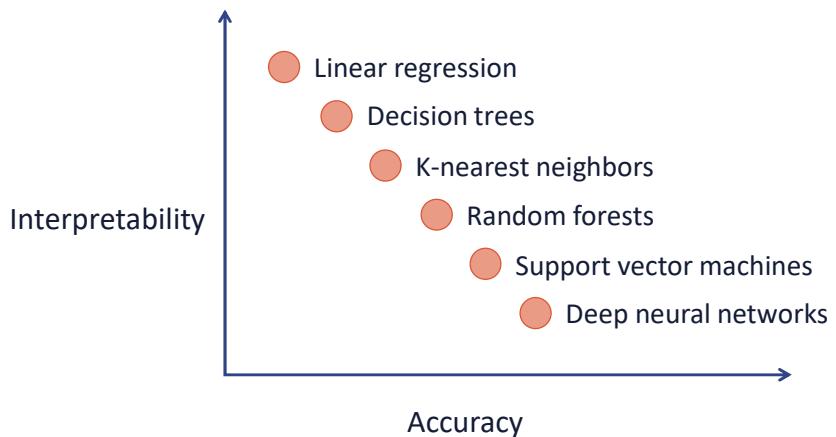
## Explaining Deep Learning Models

- **Challenges:** high complexity, many parameters
- **Techniques:** layer-wise relevance propagation, integrated gradients



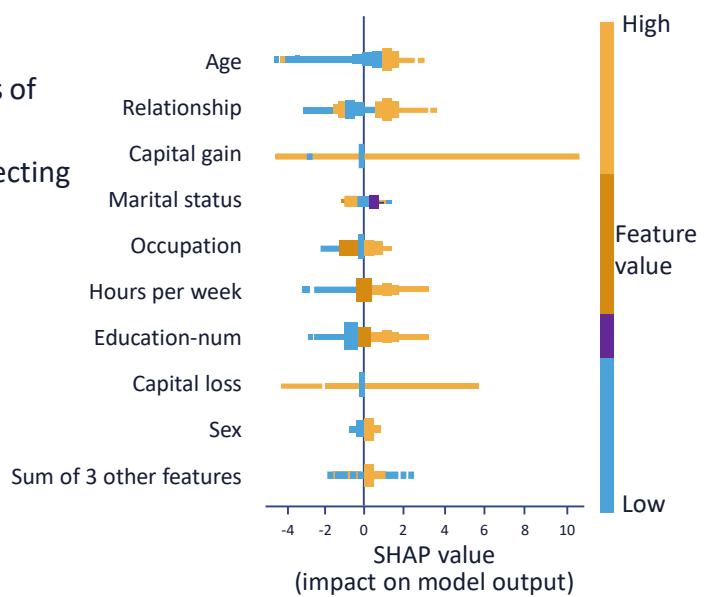
## Comparing Interpretability Across Model Types

### Comparison



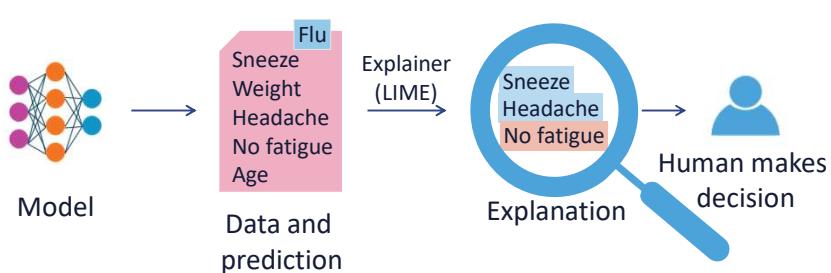
## SHAP Library

- **Functionality:** explains predictions of any model using SHAP values
- **Use case:** identify key features affecting model decisions



## Explainer

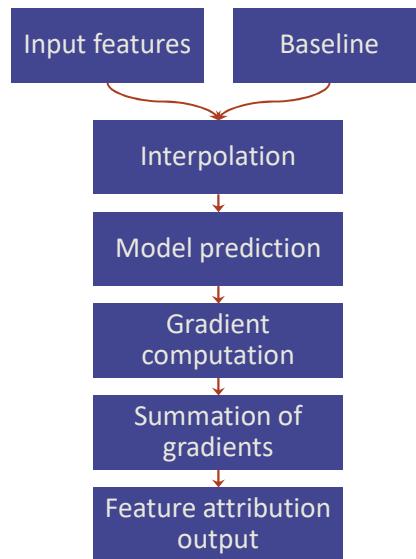
- **Functionality:** provides local explanations by approximating the decision boundary
- **Use case:** explains individual predictions for complex models



## LIME Library

## Integrated Gradients

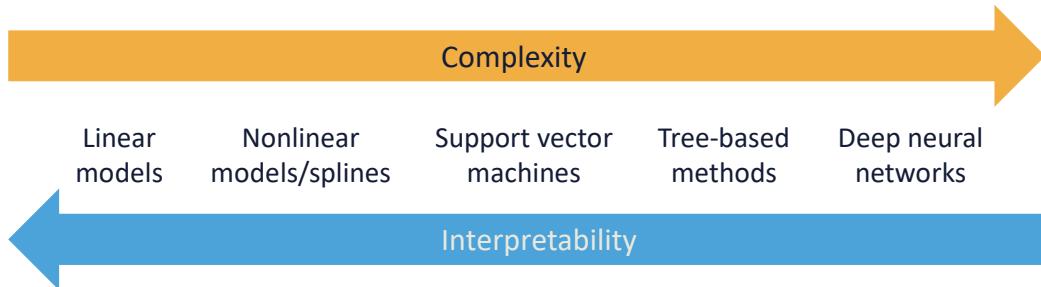
- **Functionality:** assigns attribution scores to input features
- **Use case:** explains deep learning model predictions, especially in image and text models



## Trade-offs Between Model Complexity and Interpretability

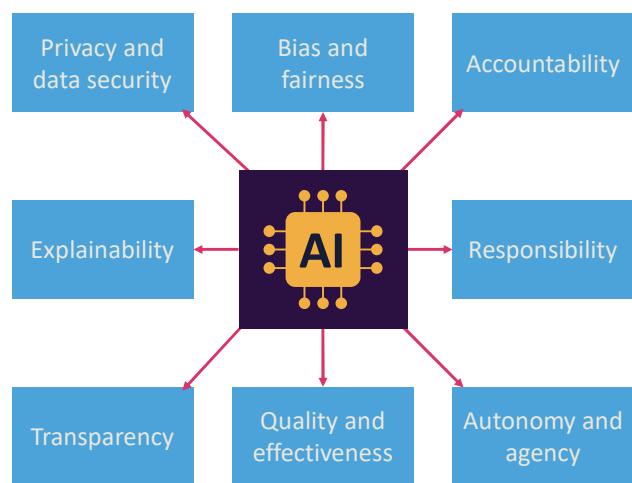
### Trade-off: Complexity and Interpretability

- **Complex models:** offer higher accuracy but lower interpretability
- **Simple models:** easier to interpret but may lack predictive power



## Regulatory and Ethical Challenges

- **Compliance:** models must meet transparency requirements (e.g., GDPR)
- **Bias and fairness:** ensuring models do not discriminate





Challenges  
ahead

## Future Directions in Model Interpretability

- **Advanced techniques:** research on better interpretability methods (e.g., counterfactual explanations)
- **Interpretable AI:** developing inherently interpretable models

### Knowledge Check

3. Local interpretability explains how a machine learning model makes decisions across the entire dataset.
  - A. True
  - B. False

## Knowledge Check

4. Which of the following techniques provides both global and local interpretability?
- A. LIME
  - B. SHAP
  - C. Gini Importance
  - D. Integrated Ingredients



## Discussion

### Interpretability in Real-World Scenarios

- Importance in ML
- Challenges of interpretability in practical applications
- Approaches to enhance interpretability for end-users



## Topic 3: Implementing Algorithms

- Criteria for choosing the right algorithm
- Algorithm selection based on domain
- Implementing a classification algorithm
- Training the model
- Implementing a regression algorithm
- Metrics for evaluating model performance

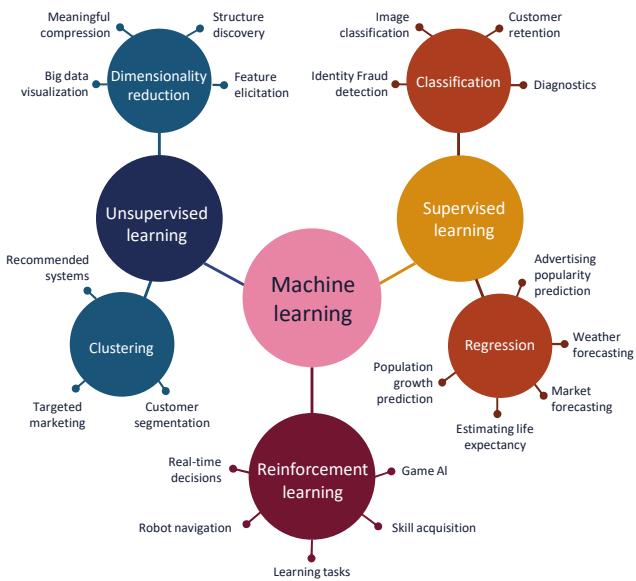


## Topic 3: Implementing Algorithms

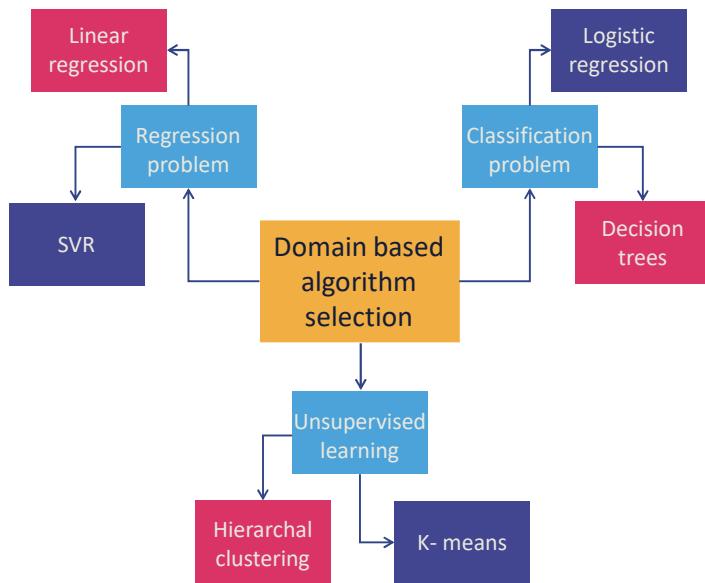
- Visualizing model performance
- Implementing and comparing two algorithms
- Visualizing model comparisons
- Visualizing model comparisons

## Criteria for Choosing the Right Algorithm

- **Problem type:** classification, regression, clustering
- **Data characterization:** size, structure, and distribution
- **Model complexity:** trade-off between accuracy and interpretability



## Algorithm Selection Based on Domain

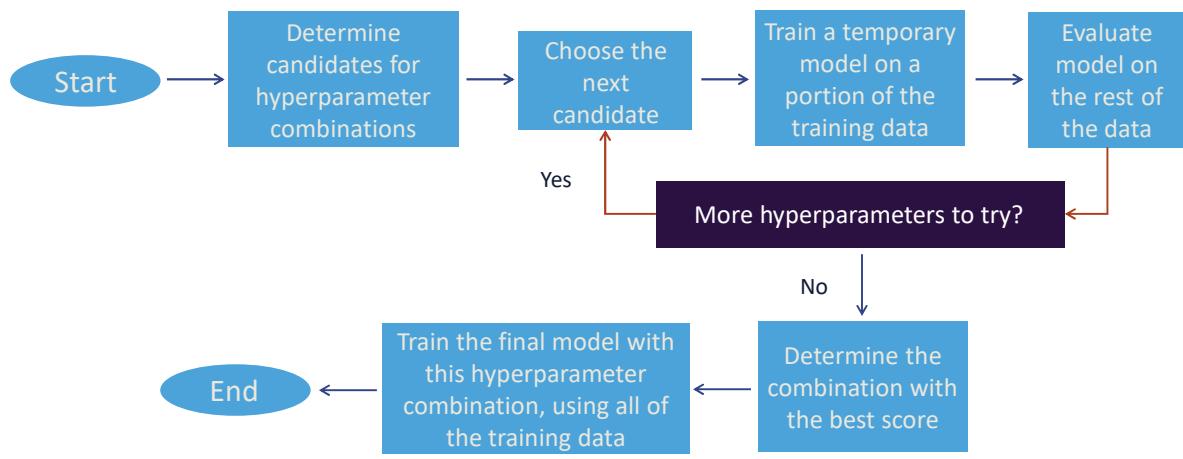


## Implementing a Classification Algorithm

```
1 from sklearn.model_selection import train_test_split
2 from sklearn.linear_model import LogisticRegression
3 from sklearn.datasets import load_iris
4
5 # Load data and split
6 X, y = load_iris(return_X_y=True)
7 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)
8
9 # Train logistic regression model
10 model = LogisticRegression()
11 model.fit(X_train, y_train)
12
13 # Make predictions
14 predictions = model.predict(X_test)
15
```

## Training the Model

- **Training process:** fit the model on training data
- **Hyperparameter tuning:** adjust parameters to optimize performance



# Implementing a Regression Algorithm

```
● ● ●  
1 from sklearn.linear_model import LinearRegression  
2  
3 # Train linear regression model  
4 regressor = LinearRegression()  
5 regressor.fit(X_train, y_train)  
6  
7 # Make predictions  
8 regression_predictions = regressor.predict(X_test)  
9
```

## Metrics for Evaluating Model Performance

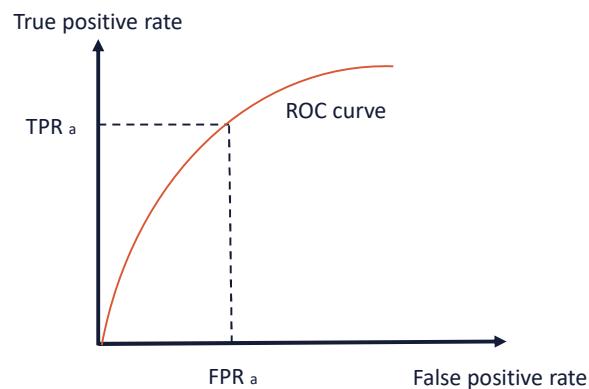
Classification

Regression

## Visualizing Model Performance

- **Confusion matrix:** visualizes classification results

		Predicted	
		Positive	Negative
Actual	Positive	True positive	False negative
	Negative	False positive	True negative



## Implementing and Comparing Two Algorithms

- **Example:** logistic regression vs. decision trees
- **Dataset:** Iris
- **Performance metrics:** accuracy and F1 Score
- **Comparison:** after training the models, predictions are made on the test set, and their metrics are printed side by side for easy comparison

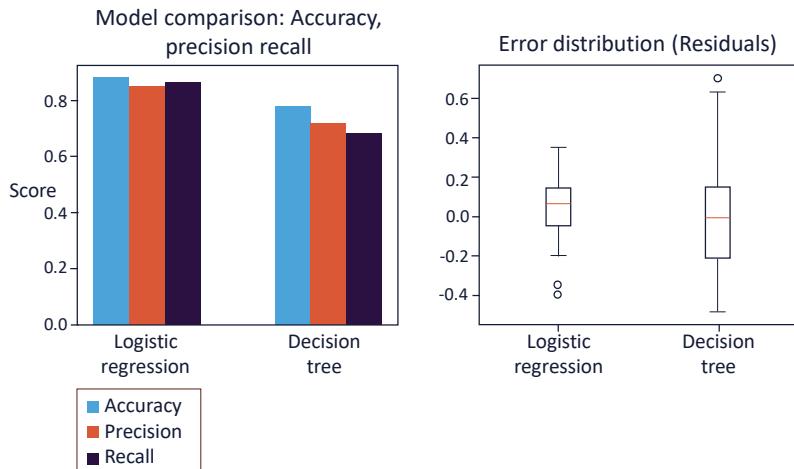
```

1 # Import necessary libraries
2 from sklearn.model_selection import train_test_split
3 from sklearn.metrics import accuracy_score, f1_score
4 from sklearn.linear_model import LogisticRegression
5 from sklearn.tree import DecisionTreeClassifier
6 from sklearn.datasets import load_iris
7
8 # Load the dataset (Iris dataset in this example)
9 data = load_iris()
10 X = data.data # Features
11 y = data.target # Labels
12
13 # Split the data into training and testing sets
14 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)
15
16 # Initialize the models
17 log_reg = LogisticRegression(max_iter=200)
18 decision_tree = DecisionTreeClassifier()
19
20 # Train both models
21 log_reg.fit(X_train, y_train)
22 decision_tree.fit(X_train, y_train)
23
24 # Make predictions
25 log_reg_pred = log_reg.predict(X_test)
26 decision_tree_pred = decision_tree.predict(X_test)
27
28 # Evaluate the models
29 log_reg_accuracy = accuracy_score(y_test, log_reg_pred)
30 decision_tree_accuracy = accuracy_score(y_test, decision_tree_pred)
31
32 log_reg_f1 = f1_score(y_test, log_reg_pred, average='weighted')
33 decision_tree_f1 = f1_score(y_test, decision_tree_pred, average='weighted')
34
35 # Print the results
36 print("Logistic Regression - Accuracy: {:.2f}, F1 Score: {:.2f}".format(log_reg_accuracy, log_reg_f1))
37 print("Decision Tree - Accuracy: {:.2f}, F1 Score: {:.2f}".format(decision_tree_accuracy, decision_tree_f1))
38

```

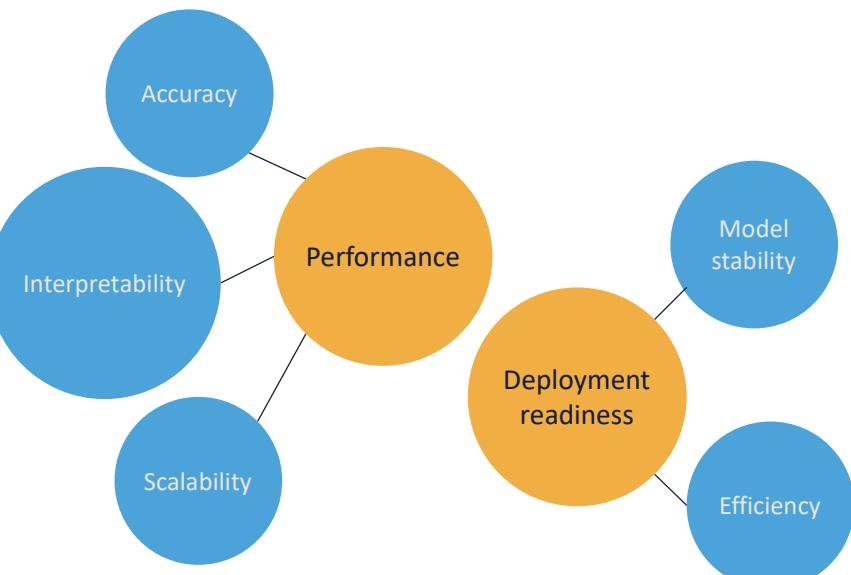
## Visualizing Model Comparisons

- **Bar chart:** compares accuracy, precision, recall
- **Box plot:** visualize the distribution of error



## Selecting the Best Algorithm for Deployment

### Algorithm Selection



## Knowledge Check

5. Which of the following is a key criterion for selecting an algorithm in machine learning?
  - A. Problem type
  - B. Data size
  - C. Algorithm popularity
  - D. All of the above

## Lab 4: Select, Implement, and Evaluate Algorithms

### Instructions

1. Explore the dataset, preprocessing, and feature engineering
2. Implement multiple machine learning algorithms
3. Evaluate the models using various metrics
4. Compare the result to identify the best-performing model

Estimated completion time: 30 minutes



## Topic 4: Experiment Tracking and Model Evaluation

- Why is experiment tracking crucial in MLOps?
- Overview of experiment tracking tools
- Integrating experiment tracking into the pipeline
- Setting up MLflow for experiment tracking
- Setting up weights & biases for experiment tracking
- Key metrics for model evaluation

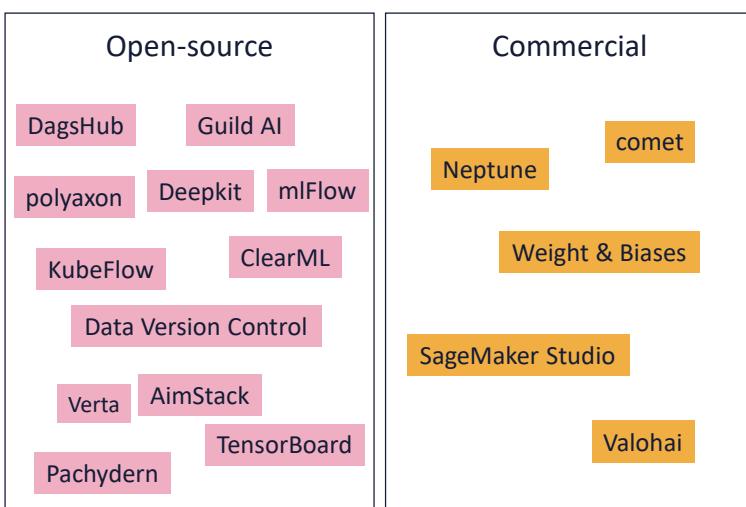
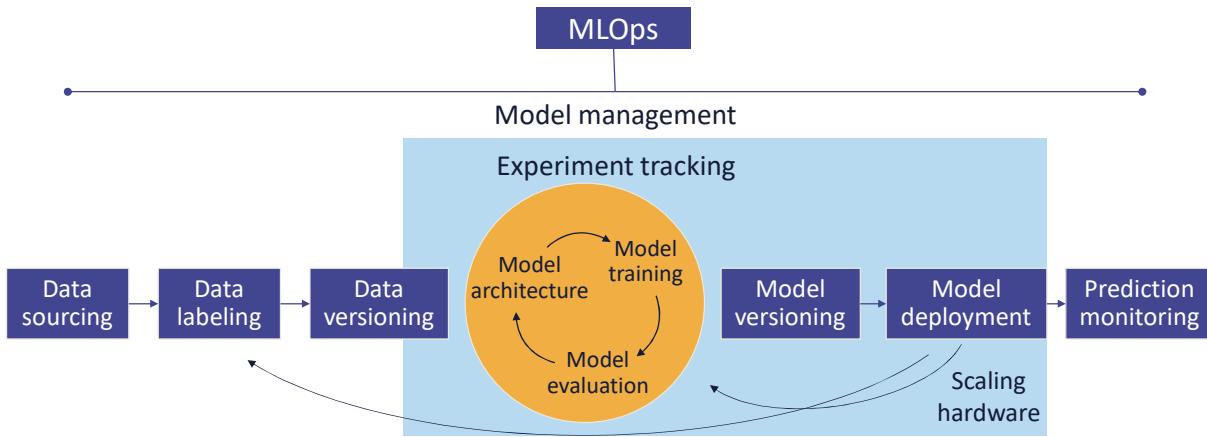


## Topic 4: Experiment Tracking and Model Evaluation

- Importance of post-deployment monitoring
- Tracking experiments with multiple models
- Visualizing metrics over time
- Confusion matrix: beyond accuracy
- Using ROC curves and precision-recall curves
- Visualizing feature importance

## Why Is Experiment Tracking Crucial?

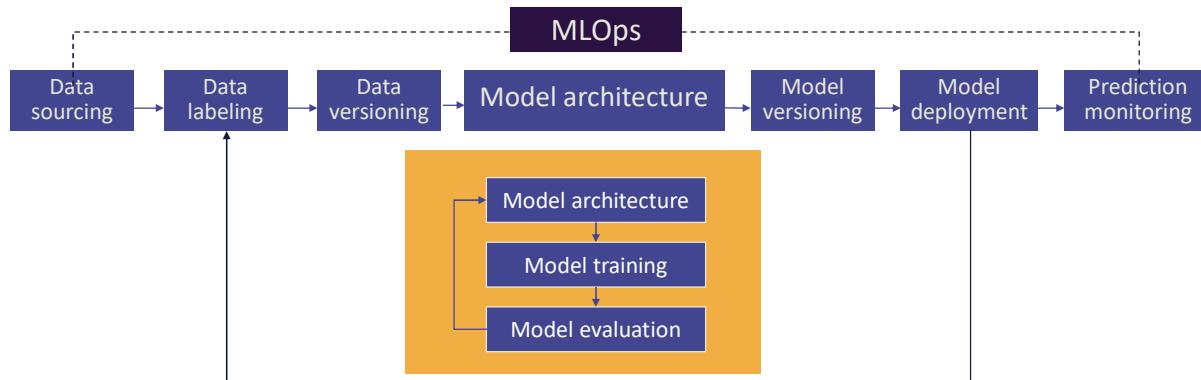
- **Definition:** documenting all experiments, models, and metrics
- **Benefits:** enables reproducibility, comparison, and collaboration



## Overview of Experiment Tracking Tools

## Integrating Experiment Tracking into the Pipeline

- **Automation:** automate logging of experiments and metrics
- **Code integration:** incorporate tracking code into existing scripts



## Setting Up MLflow for Experiment Tracking

- **Installation:** install MLflow using pip.
- **Basic setup:** create a new experiment and start logging

```

● ● ●
1 # Install MLflow
2 !pip install mlflow
3
4 # Create a new experiment
5 import mlflow
6 mlflow.create_experiment("Experiment Tracking")
7
8 # Start logging
9 mlflow.log_param("param_name", value)
10 mlflow.log_metric("metric_name", value)
11
  
```

## Setting Up Weights & Biases for Experiment Tracking

- **Installation:** install weights & biases using pip
- **Basic setup:** initialize a new project and start logging

```
● ● ●
1 # Install Weights & Biases
2 !pip install wandb
3
4 # Initialize a new project
5 import wandb
6 wandb.init(project="experiment-tracking")
7
8 # Log parameters and metrics
9 wandb.log({"accuracy": 0.95, "loss": 0.05})
```

## Key Metrics for Model Evaluation

### Classification Metrics

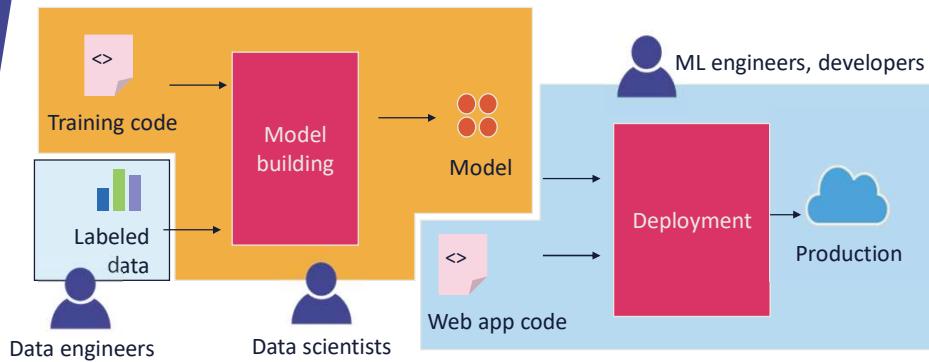
- **Precision** =  $\frac{TP}{(TP+FP)}$
- **Recall** =  $\frac{TP}{(TP+FN)}$
- **F1** =  $\frac{2 \times Precision \times Recall}{(Precision+Recall)}$
- **ROC-AUC** =  $\int_0^1 TPR(FPR^{-1}(x))dx$

### Regression Metrics

- **MAE** =  $\frac{\sum_{i=1}^n |y_i - \hat{x}_i|}{n}$
- **MSE** =  $\frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2$
- **R2** =  $1 - \frac{\sum_i (y_i - \hat{y}_i)^2}{\sum_i (y_i - \bar{y}_i)^2}$

## Importance of Post-Deploy Monitoring

- **Model drift:** monitor for changes in data distribution
- **Performance degradation:** track metrics over time to identify issues

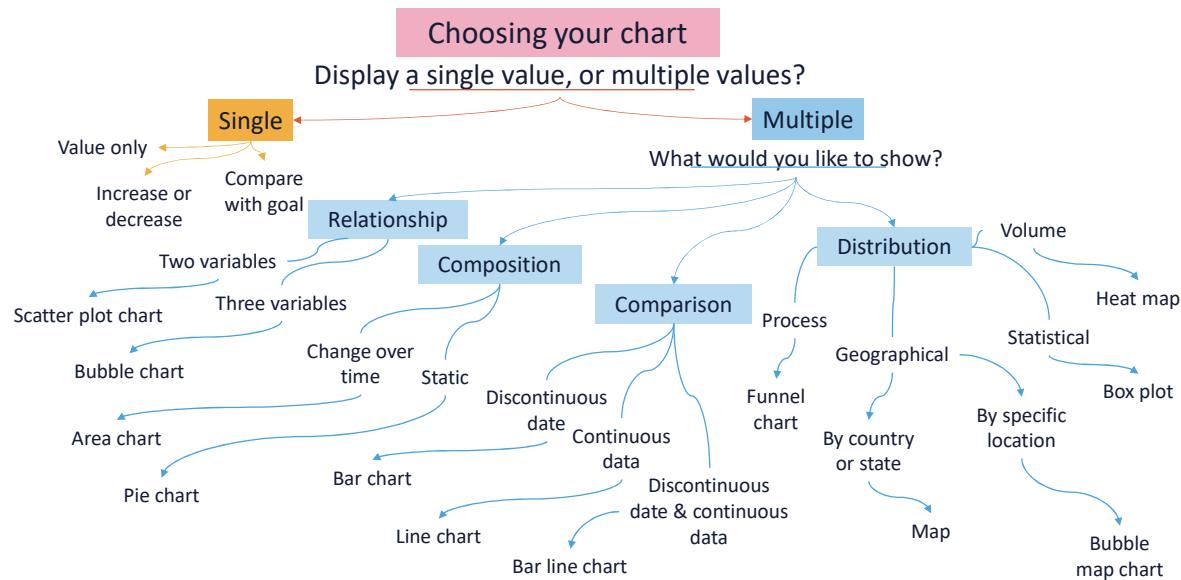


## Tracking Experiments with Multiple Models

- **Multi-model tracking:** compare performance across different models
- **Versioning:** keep track of different model versions



## Visualization Metrics Over Time



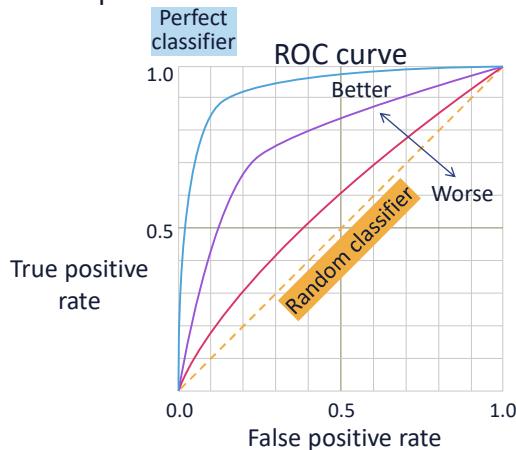
## Confusion Matrix: Beyond Accuracy

- Understanding model error:** includes correct predictions and tells where the model goes wrong
- Error analysis:** focus on false positives and false negatives
- Insight:** helps improve model performance through targeted corrections

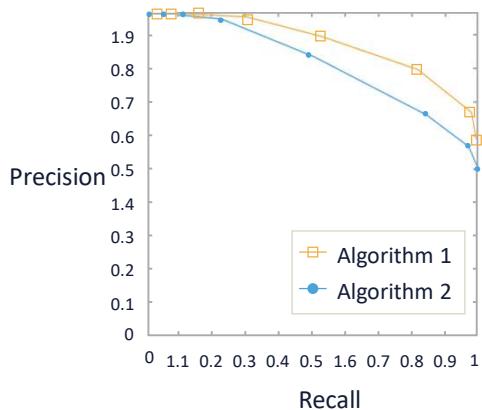
		Predicted values	
		Positive	Negative
Actual values	Positive	True positive (TP)	False negative (FN) (Type II Error)
	Negative	False positive (FP) (Type I Error)	True negative (TN)

## Using ROC Curves and Precision-Recall Curves

- **ROC Curve:** plots true positive rate vs. false positive rate

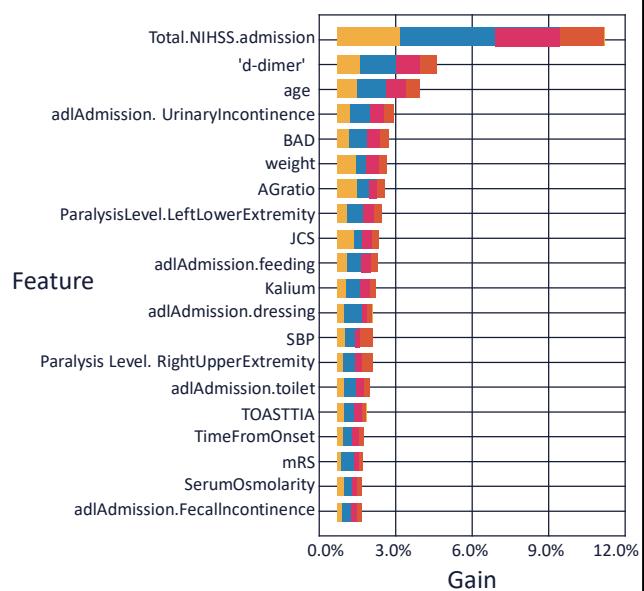


- **Precision-Recall Curve:** useful for imbalanced datasets



## Visualizing Feature Importance

- **Feature importance plot:** rank features by their impact on the model
- **SHAP values:** provides both global and local explanations



## Knowledge Check

6. What is the primary benefit of experiment tracking in MLOps?
  - A. Faster model training
  - B. Improved data preprocessing
  - C. Enables reproducibility and collaboration
  - D. Automates data collection

## Knowledge Check

7. Which tool is commonly used for tracking experiments and visualizing metrics in machine learning projects?
  - A. PyTorch
  - B. TensorFlow
  - C. MLflow
  - D. Jupyter



## Topic 5: Setting Up MLflow for Experiment Tracking

- What is MLflow?
- Installing and setting up MLflow
- Logging parameters, metrics, and artifacts
- Integrating MLflow into the model development pipeline
- Using MLflow tracking UI
- Analyzing metrics to determine the best model



## Topic 5: Setting Up MLflow for Experiment Tracking

- Comparing different runs of an experiment
- Storing models in MLflow
- Retrieving and using stored models
- Best practices for using MLflow in production

## What Is MLflow?

- **Definition:** MLflow is an open-source platform for managing the end-to-end machine learning lifecycle
- **Capabilities:** include experiment tracking, project packaging, model management, and a central model registry
- **Components:** MLflow Tracking, MLflow Projects, MLflow Models, and MLflow Registry

## Installing and Setting up MLflow



- **Installation:** install MLflow using `pip install mlflow`
- **Setup:** start MLflow server with `mlflow ui`
- **Configuration:** set the tracking **URI** (uniform resource identifier) to the server's address

## Logging Parameters, Metrics, and Artifacts

- **Parameters**: log hyperparameters like learning rate and number of layers
- **Metrics**: record evaluation metrics like accuracy and loss
- **Artifacts**: save artifacts like model files, plots, and datasets

```
● ● ●  
1 import mlflow  
2  
3 # Log parameters  
4 mlflow.log_param("learning_rate", 0.01)  
5 mlflow.log_param("n_layers", 3)  
6  
7 # Log metrics  
8 mlflow.log_metric("accuracy", 0.95)  
9 mlflow.log_metric("loss", 0.05)  
10  
11 # Log artifact  
12 mlflow.log_artifact("model.pkl")
```

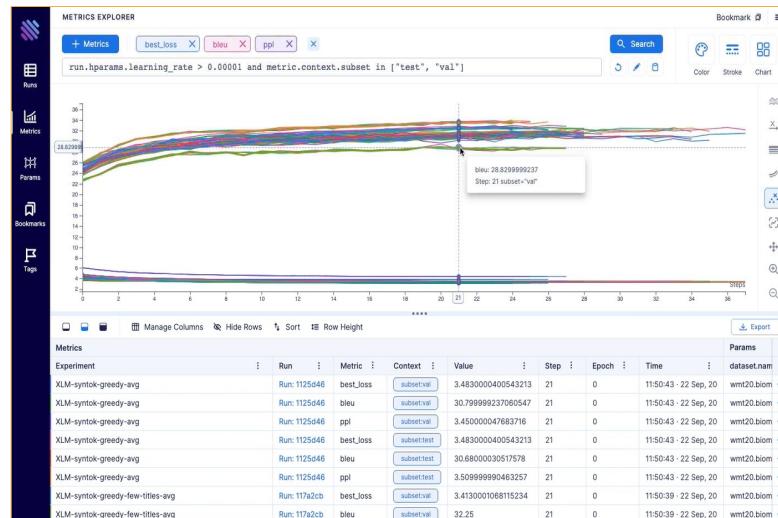
## Integrating MLflow into the Model Development Pipeline

- **Integration points**: use MLflow in any stage of the model development pipeline
- **Logging from code**: use `mlflow.start_run()` to begin tracking an experiment
- **Automatic logging**: enable auto-logging for popular frameworks like TensorFlow and PyTorch

```
● ● ●  
1 import mlflow  
2  
3 # Start a new run  
4 with mlflow.start_run():  
5     # Log parameters and metrics here  
6     mlflow.log_param("learning_rate", 0.01)  
7     mlflow.log_metric("accuracy", 0.95)
```

## Using MLflow Tracking UI

- Experiment dashboard:** view all experiments, runs, and their results
- Run comparison:** compare parameters, metrics, and artifacts across multiple runs
- UI customization:** filter and sort experiments to analyze performance

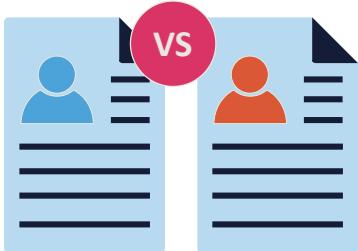


## Analyzing Metrics to Determine Best Model



- Key metrics:** focus on metrics like accuracy, precision, recall, and F1 score
- Visualization:** use MLflow's built-in visualization tools to plot metrics
- Model selection:** identify the best model based on metric performance

## Comparing Different Runs of an Experiment



- **Run comparison:** select multiple runs to compare side-by-side
- **Parameter analysis:** evaluate the impact of different parameters on performance
- **Artifact review:** compare artifacts like model files, confusion matrices, and plots

## Storing Models in MLflow

- **Model logging:** use `mlflow.log_model()` to save models during training
- **Model registry:** organize and manage different model versions in the registry
- **Version control:** keep track of model versions and their associated metadata

```
● ○ ●
1 import mlflow
2
3 # Save the model
4 mlflow.log_model(model, "model")
5
6 # Register the model in the registry
7 mlflow.register_model("model", "my_model")
```

## Retrieving and Using Stored Models

- **Model retrieval:** use `mlflow.load_model()` to load a stored model
- **Deployment:** deploy the model using MLflow's deployment tools
- **Prediction:** use the loaded model to make predictions on new data

```
● ● ●  
1 import mlflow  
2  
3 # Load the model  
4 model = mlflow.load_model("model")  
5  
6 # Make predictions  
7 predictions = model.predict(X_new)
```

## Best Practices for Using MLflow in Production



Best practice

- **Model versioning:** use the model registry to track different versions
- **Automated logging:** automate parameter and metric logging in training scripts
- **Scalability:** use cloud storage and remote tracking servers for large-scale experiments

## Knowledge Check

8. Which component of MLflow is responsible for tracking and logging experiments, parameters, and metrics?
- A. MLflow Registry
  - B. MLflow Projects
  - C. MLflow Tracking
  - D. MLflow Models



## Interactive Activity

### Interpretability in Real-World Scenarios

- Develop a strategy for tracking experiments in large ML projects
- Identify essential metrics to monitor model performance
- Explore best practices for managing multiple experiments with MLflow

## Lab 5: Experiment Tracking and Model Management with MLflow

### Instructions

1. Set up MLflow for tracking experiments
2. Track metrics and parameters for different models
3. Compare the results of multiple model runs and visualize them in MLflow's UI
4. Save and retrieve models for deployment

Estimated completion time: 30 minutes



## Topic 6: Evaluating Models

- Setting up the evaluation environment
- Implementing key evaluation metrics
- Visualizing evaluation results
- Comparing models using evaluation metrics
- Interpreting the comparison results

## Setting Up Evaluation Environment



- **Environment setup:** install necessary libraries (e.g., scikit-learn, matplotlib)
- **Data preparation:** load and preprocess the evaluation dataset
- **Model loading:** load re-trained models for evaluation

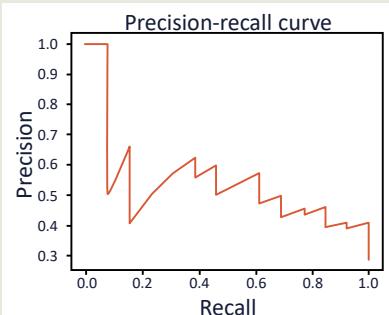
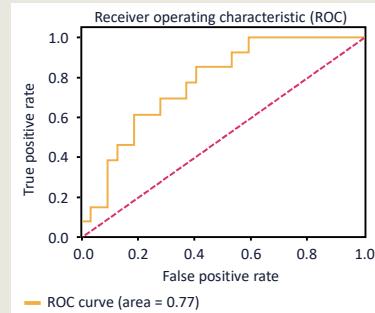
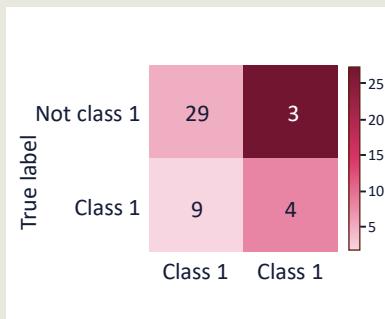
## Implementing Key Evaluation Metrics

- **Confusion matrix:** evaluates true vs. predicted labels
- **ROC curve:** assesses the true positive rate vs. false positive rate
- **Precision-recall curve:** measures precision and recall trade-off

```
1 from sklearn.metrics import confusion_matrix, roc_curve, precision_recall_curve
2 import matplotlib.pyplot as plt
3
4 # Confusion Matrix
5 cm = confusion_matrix(y_true, y_pred)
6 print("Confusion Matrix:\n", cm)
7
8 # ROC Curve
9 fpr, tpr, _ = roc_curve(y_true, y_score)
10 plt.plot(fpr, tpr)
11 plt.title("ROC Curve")
12 plt.show()
13
14 # Precision-Recall Curve
15 precision, recall, _ = precision_recall_curve(y_true, y_score)
16 plt.plot(recall, precision)
17 plt.title("Precision-Recall Curve")
18 plt.show()
```

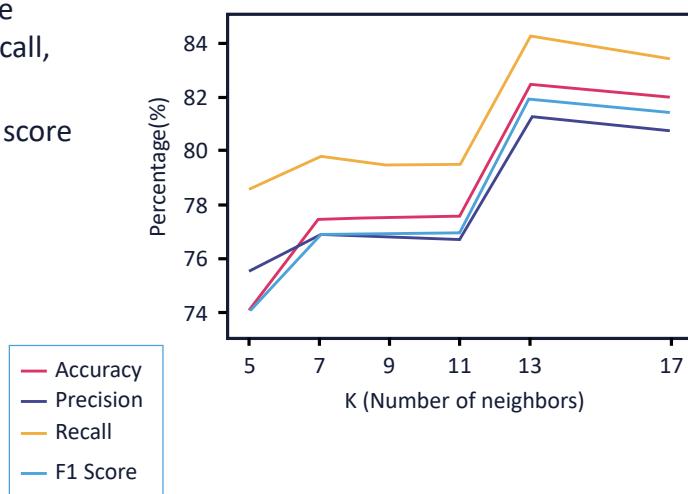
## Visualizing Evaluation Results

- **Confusion matrix plot:** visualizes true positives, false positives, etc.
- **ROC curve plot:** shows the trade-off between sensitivity and specificity
- **Precision-recall curve plot:** evaluates imbalanced datasets



## Comparison Models Using Evaluation Metrics

- **Side-by-side comparison:** compare metrics like accuracy, precision, recall, etc.
- **Aggregate score:** use a composite score to rank models



## Setting Up Evaluation Environment



- **Model selection:** choose the best model based on performance and interpretability
- **Actionable insight:** identify areas for improvement and optimization

## Lab 6: Model Performance Evaluation and Comparison

### Instructions

1. Set up an environment for comprehensive model evaluation and comparison
2. Implement a range of evaluation metrics
3. Visualize model performance with plots and charts
4. Interpret results to analyze trade-offs between models for better decision-making

Estimated completion time: 25 minutes



## Topic 7: Hyperparameter Tuning Techniques

- What are hyperparameters and why they matter
- Hyperparameters vs. parameters
- Overview of grid search technique
- Overview of random search technique
- Grid search vs. random search: comparison
- Introduction to Bayesian optimization

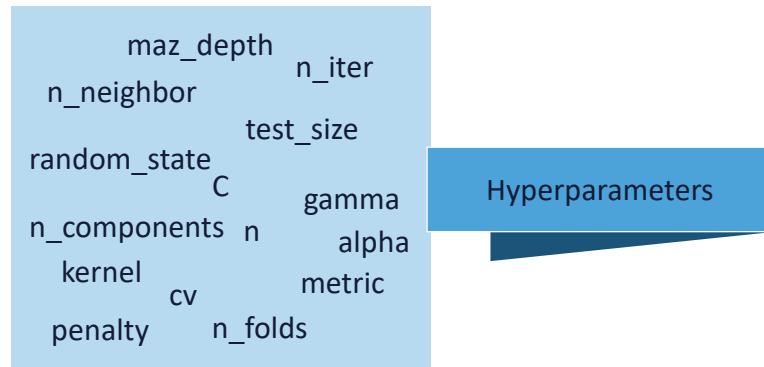


## Topic 7: Hyperparameter Tuning Techniques

- Bayesian optimization: how it works
- Time and resource constraints
- Tips for efficient hyperparameter tuning
- Hyperparameter tuning in large-scale systems

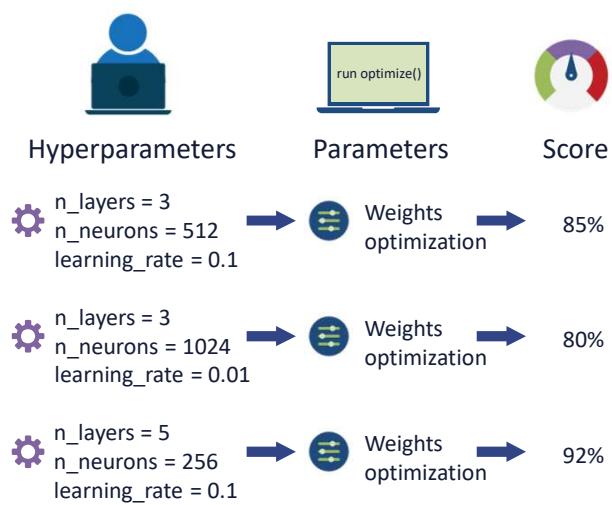
## What Are Hyperparameters?

- **Definition:** hyperparameters are configurations set before the training process
- **Importance:** control model behavior and performance
- **Examples:** learning rate, number of layers, tree depth, etc.



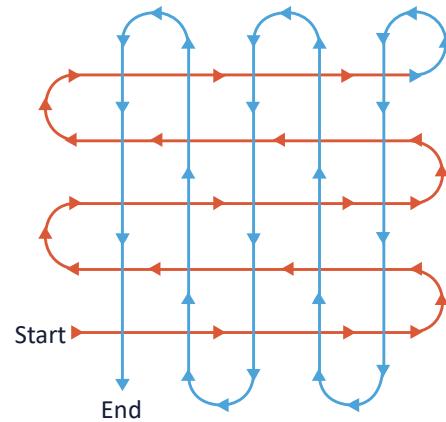
## Hyperparameters vs. Parameters

- **Hyperparameters:** set before training, influence the learning process
- **Parameters:** learned from data, represent model coefficients
- **Examples:** in linear regression, hyperparameters might include the learning rate, while parameters are the weights and biases



## Overview of the Grid Search Technique

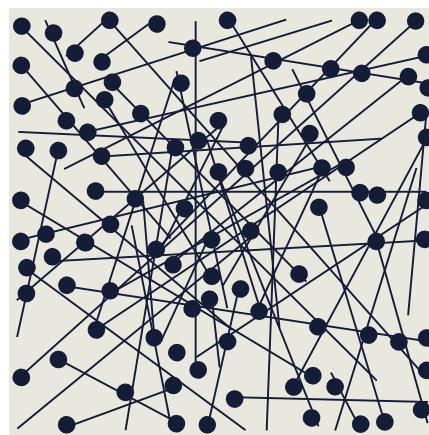
- **Definition:** exhaustive search over a predefined hyperparameter space
- **Advantage:** finds the best combination of hyperparameters
- **Limitation:** computationally expensive, time-consuming



Procedure for grid search pattern technique

## Overview of the Random Search Technique

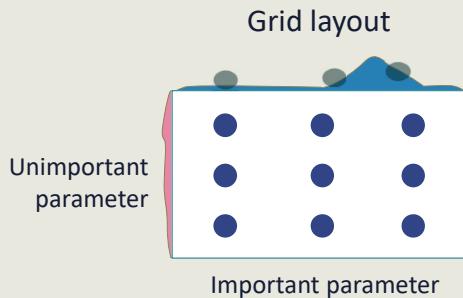
- **Definition:** randomly sample hyperparameter combinations from a distribution
- **Advantage:** more efficient than grid search, covers larger search space
- **Limitation:** may not find the absolute best combination



## Grid Search vs. Random Search: Comparison

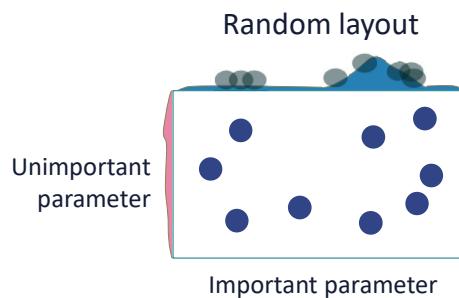
### Grid search

- Exhaustive
- Suitable for small datasets
- Time consuming



### Random search

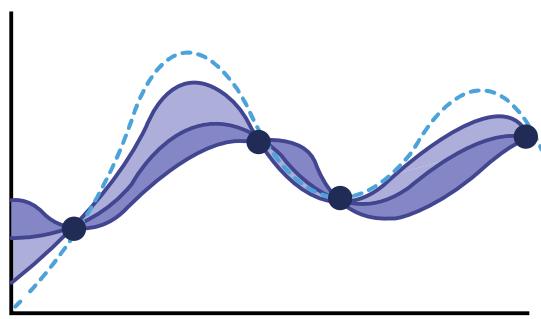
- Efficient
- Scalable
- Good for large datasets



**Recommendation:** use random search when computational resources are limited

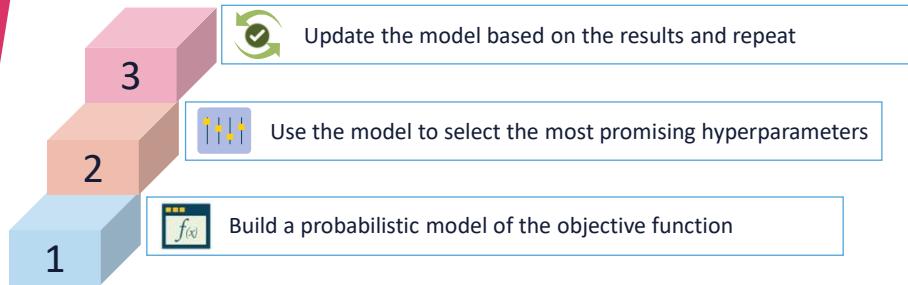
## Introduction to Bayesian Optimization

- **Definition:** uses probabilistic models to predict the performance of hyperparameters
- **Advantage:** efficient search, focuses on promising areas
- **Application:** suitable for high-dimensional and expensive search spaces



## Bayesian Optimization: How Does It Work?

### Workflow



## Time and Resource Constraints

- **Time complexity:** grid search is exhaustive, while Bayesian is efficient
- **Resource management:** prioritize resources based on dataset size and model complexity
- **Scalability:** use distributed search for large-scale problems



## Tips for Hyperparameter Tuning



- **Start simple:** begin with a small search space and simple methods
- **Use random search first:** explore a broad range of hyperparameters
- **Leverage early stopping:** avoid overfitting and save time

- **Distributed search:** use tools like Ray Tune for distributed tuning
- **Parallel processing:** speed up tuning with parallel evaluations
- **Automated tuning:** consider AutoML framework for complex tasks

**Hyperparameter Tuning in Large Scale Systems**

## Knowledge Check

9. Bayesian optimization focuses the search on the least promising areas of the hyperparameter space.
- A. True
  - B. False

## Knowledge Check

10. Which of the following is a disadvantage of Grid Search for hyperparameter tuning?
- A. It randomly samples hyperparameters
  - B. It is computationally expensive and time-consuming
  - C. It uses probabilistic models to predict performance
  - D. It is efficient for large-scale problems



## Topic 8: Automated Hyperparameter Tuning

- Overview of automated hyperparameter tuning
- Tools for automated hyperparameter tuning
- Grid search and random search in scikit-learn
- Setting up Bayesian optimization with Optuna
- Running the hyperparameter tuning process

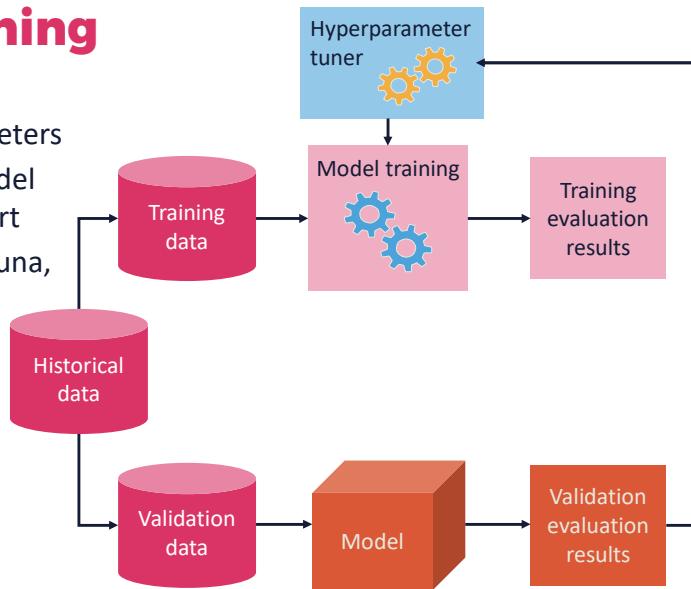


## Topic 8: Automated Hyperparameter Tuning

- Logging results with Optuna
- Running hyperparameter tuning with Hyperopt
- Analyzing the best hyperparameters
- Visualizing the results
- Selecting the optimal hyperparameter set for deployment

## Overview of Automated Hyperparameter Tuning

- **Definition:** automated process for searching for the best hyperparameters
- **Benefits:** saves time, improves model performance, reduces manual effort
- **Automated tuning:** Hyperopt, Optuna, Auto-Sklearn



## Tools for Automated Hyperparameter Tuning

### Hyperopt

Uses tree-structured Parzen estimators for efficient search

### Optuna

Offers advanced features like pruning and parallel execution

### Auto-Sklearn

Combines automated hyperparameter tuning with model selection

## Grid Search and Random Search in Scikit-Learn

- **Grid search:** exhaustively searches over a predefined parameter grid
- **Random search:** randomly samples parameter combinations
- **Implementation:** use `GridSearchCV` and `RandomizedSearchCV` from scikit-learn

```

1  from sklearn.model_selection import GridSearchCV, RandomizedSearchCV
2
3  # Define parameter grid
4  param_grid = {'C': [0.1, 1, 10], 'kernel': ['linear', 'rbf']}
5  model = SVC()
6
7  # Grid Search
8  grid_search = GridSearchCV(model, param_grid, cv=5)
9  grid_search.fit(X_train, y_train)
10
11 # Random Search
12 random_search = RandomizedSearchCV(model, param_grid, n_iter=10, cv=5)
13 random_search.fit(X_train, y_train)

```

## Setting Up Bayesian Optimization with Optuna

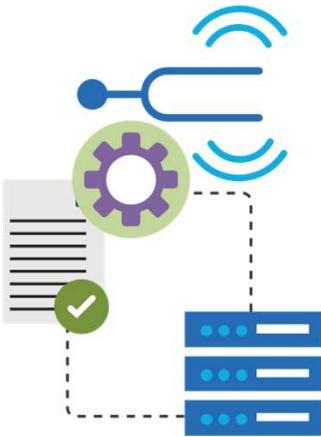
- **Installation:** install Optuna using pip
- **Basic setup:** define the objective function and search space
- **Running trials:** optimize the objective function over multiple trials

```

1  !pip install optuna
2
3  import optuna
4
5  # Define objective function
6  def objective(trial):
7      C = trial.suggest_loguniform('C', 1e-5, 1e2)
8      kernel = trial.suggest_categorical('kernel', ['linear', 'rbf'])
9      model = SVC(C=C, kernel=kernel)
10     return cross_val_score(model, X_train, y_train, cv=5).mean()
11
12 # Create a study and optimize
13 study = optuna.create_study(direction='maximize')
14 study.optimize(objective, n_trials=50)

```

## Running the Hyperparameter Tuning Process



- **Start simple:** begin with a small search space and simple methods
- **Use random search first:** explore a broad range of hyperparameters
- **Leverage early stopping:** avoid overfitting and save time

## Logging Results with Optuna

- **Built-in logging:** Optuna logs each trial's parameters and metrics
- **Visualization:** Use `plot_optimization_history()` to visualize progress
- **Tracking best parameters:** easily access the best parameters and score

```
● ● ●  
1 # Visualize optimization history  
2 optuna.visualization.plot_optimization_history(study)  
3  
4 # Retrieve the best parameters and score  
5 best_params = study.best_params  
6 best_score = study.best_value  
7 print(f"Best Parameters: {best_params}")  
8 print(f"Best Score: {best_score}")
```

# Running Hyperparameter Tuning with Hyperopt

- **Hyperopt overview:** use tree structured Parzen estimators for optimization
- **Basic setup:** define objective function and search space
- **Running for search:** use `fmin()` to run the optimization

```
● ● ●
1 from hyperopt import fmin, tpe, hp, Trials
2
3 # Define the objective function
4 def objective(params):
5     model = SVC(**params)
6     return -cross_val_score(model, X_train, y_train, cv=5).mean()
7
8 # Define the search space
9 space = {
10     'C': hp.loguniform('C', -5, 2),
11     'kernel': hp.choice('kernel', ['linear', 'rbf'])
12 }
13
14 # Run the optimization
15 trials = Trials()
16 best = fmin(objective, space, algo=tpe.suggest, max_evals=50, trials=trials)
```

# Analyzing the Best Hyperparameters



- **Best parameters:** review the top-performing hyperparameters
- **Model performance:** check how these parameters impact model accuracy
- **Interpretation:** understand why certain configurations work better

## Visualizing the Results



- **Hyperparameter importance:** visualize impact of each hyperparameter
- **Performance plots:** compare performance among different configurations
- **Optuna plots:** use Optuna's built-in plots for analysis



## Selecting the Optimal Hyperparameter Set for Deployment

- **Best model:** select the model with the highest evaluation score
- **Final training:** retain the model using the best hyperparameters
- **Deployment readiness:** ensure the model is stable and scalable

## Knowledge Check

11. Automated hyperparameter tuning always guarantees finding the optimal hyperparameter configuration.
- A. True
  - B. False

## Lab 7: Implementing Automated Hyperparameter Tuning

### Instructions

1. Set up a hyperparameter tuning job using various automated tuning methods
2. Run the tuning process to identify the best hyperparameter for the model
3. Analyze and interpret the tuning results to determine the best model configuration

Estimated completion time: 30 minutes

## Challenge Lab 2: Experiment Tracking and Hyperparameter Optimization in MLOps

### Instructions

1. Set up MLflow for tracking model training experiments
2. Log model parameters, metrics, and artifacts
3. Train two ML models and track both models' performance using MLflow
4. Implement hyperparameter tuning using Grid Search for Random Forest model
5. Compare the performance of the tuned model to the baseline model and log the final results in MLflow

Estimated completion time: 45 minutes



## Interactive Activity

### Solving Real-world MLOps Challenges

- Identify and address real-world MLOps challenges
- Collaborate on solutions for scaling, and deployment automation
- Present proposed solutions to enhance MLOps practices

## Agenda Review

### Module 1: MLOps and Data

### Module 2: Development, Experimentation, and Evaluation

- ✓ Model Development Strategies
- ✓ ML Model Interpretability and Explainability
- ✓ Implementing Algorithms
- ✓ Experiment Tracking and Model Evaluation
- ✓ Setting Up MLflow for Experiment Tracking
- ✓ Evaluating Models
- ✓ Hyperparameter Tuning Techniques
- ✓ Automated Hyperparameter Tuning

### Module 3: Deployment, Monitoring, and Pipelines



## Module 3: Deployment, Monitoring, and Pipelines

After today's module, you will be able to

- Learn strategies for deploying ML models on cloud platforms
- Address legal and ethical consideration in MLOps
- Design CI/CD workflows and implement monitoring tools
- Explore cloud, edge and federated learning for real-time ML



## Module Agenda

- Module 1: MLOps and Data**
- Module 2: Development, Experimentation, and Evaluation**
- Module 3: Deployment, Monitoring, and Pipelines**
- Model Serving and Deployment Strategies
  - Legal and Compliance Issues in MLOps
  - Containerizing ML Models with Docker
  - Deploying Models to Cloud Platforms
  - Federated Training and Edge Deployments
  - CI/CD for ML
  - Setting Up CI/CD Pipelines for ML
  - Monitoring and Maintaining ML Systems
  - Implementing Monitoring Tools



## Topic 1: Model Serving and Development Strategies

- What is model serving?
- Batch vs. real-time serving
- Key considerations for model serving
- Overview of deployment strategies
- Benefits of A/B testing and blue-green deployment
- Implementing canary deployment
- Benefits of containerization in ML deployment



## Topic 1: Model Serving and Development Strategies

- Introduction to docker for containerization
- Steps to containerize an ML model with Docker
- Serving models with Docker
- Overview of TensorFlow serving
- Introduction to TorchServe
- Overview of Ray Serve

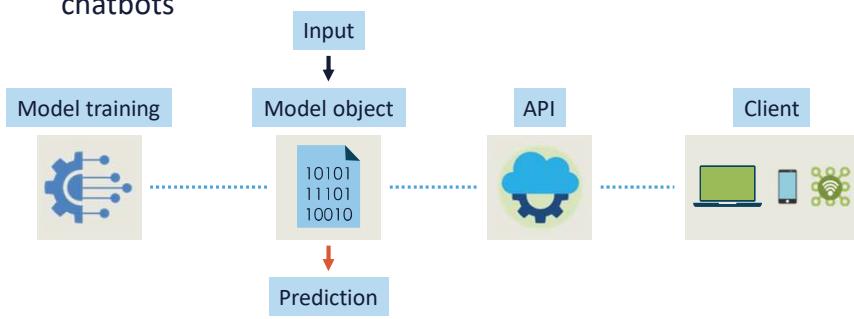


## Topic 1: Model Serving and Development Strategies

- Steps for deploying a model on AWS SageMaker
- Deploying models on Google AI platform
- Deploying models on Azure Machine Learning
- Best practices for model deployment on cloud platforms

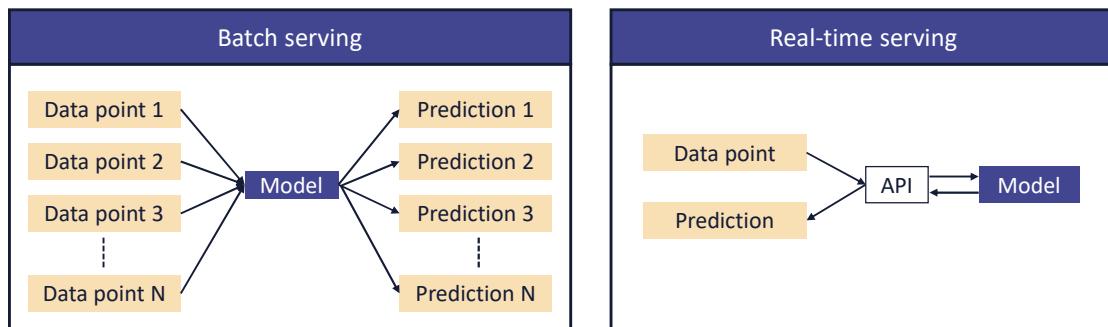
### What Is Model Serving?

- **Definition:** model serving refers to the purpose of making a trained model available for predictions in a production environment
- **Importance:** enables applications to leverage machine learning models for real-time or batch predictions
- **Applications:** fraud detection, recommendation systems, chatbots



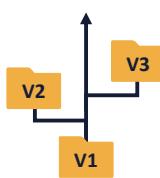
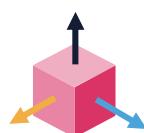
## Batch vs. Real-Time Serving

- **Batch serving:** predictions are made on a large dataset at once, suitable for non-time-critical applications
- **Real-time:** predictions are made instantaneously for each request, essential for time-sensitive applications
- **Use cases:** batch for report generation, real-time for fraud detection



## Key Considerations for Model Serving

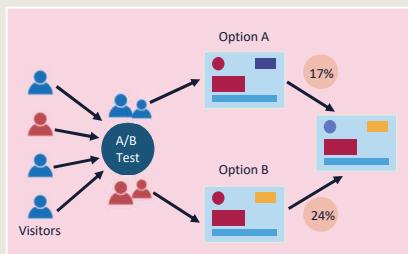
- **Latency:** ensuring low response times for real-time predictions
- **Scalability:** ability to handle an increasing number of requests
- **Model versioning:** managing different versions of a model in production



## Overview of Deployment Strategies

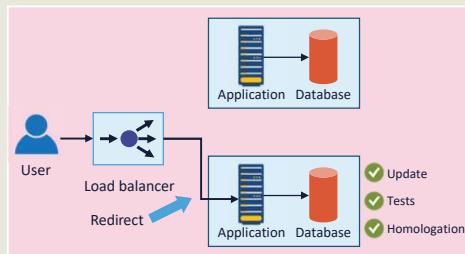
### A/B testing

- Deploying multiple models to compare performance



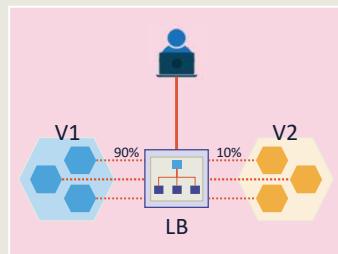
### Blue-green deployment

- Switching between two environments to minimize downtime



### Canary deployment

- Gradual rollout to a subset users to test stability



## Benefits of A/B Testing and Blue-Green Deployment

### A/B testing

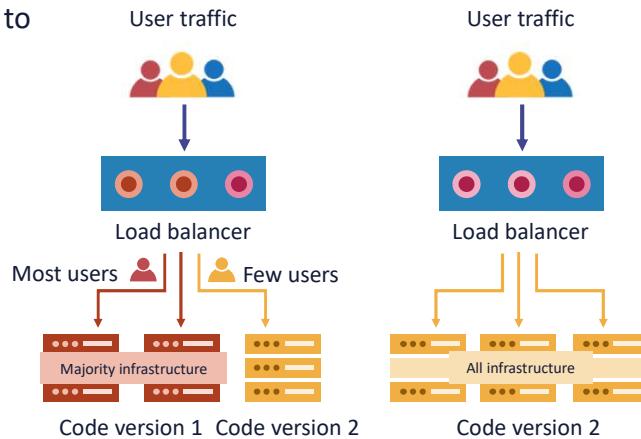
- Performance comparison
  - Evaluate different models simultaneously
- User feedback
  - Collect feedback based on real-world usage
- Risk reduction
  - Test updates on a smaller group before full deployment

### Blue-green deployment

- Zero downtime
  - Seamless transition between environments
- Easy rollback
  - Quickly revert to the previous version if issues arise
- Controlled rollout
  - Gradual deployment to monitor user experience

## Implementing Canary Deployment

- **Definition:** gradual rollout of a model to a small subset of users
- **Steps:** deploy to a small user base, monitor performance, and gradually increase user base
- **Benefits:** reduces risk of failure and allows for real-time monitoring and adjustments



1

**Consistency**

ensures the model runs the same across different environments



2

**Scalability**

easily replicate containers to scale up model serving



3

**Isolation**

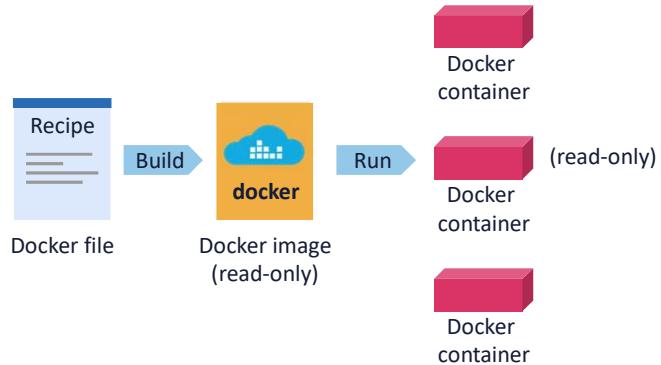
keeps the model and its dependencies isolated from other processes



## Benefits of Containerization in ML Deployment

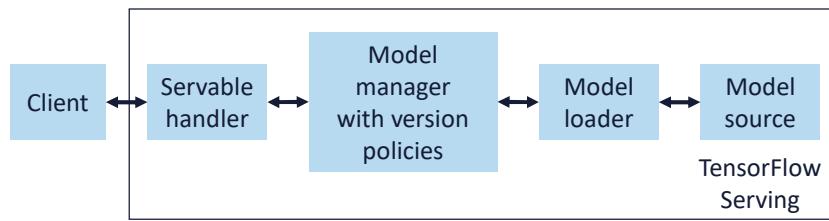
## Introduction to Docker for Containerization

- **What is Docker?** A platform for developing, shipping, and running applications in containers
- **Role in ML deployment:** packages ML models and dependencies in a portable container
- **Advantages:** easy deployment, environment consistency, and scalability



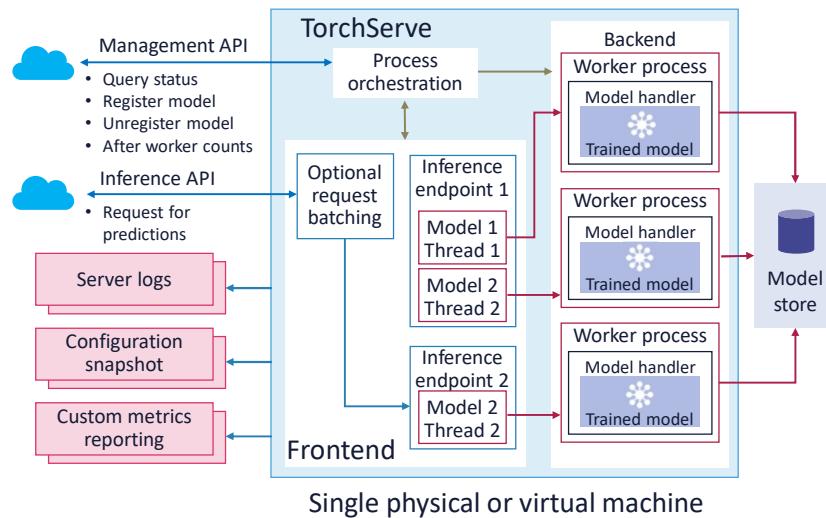
## Overview of TensorFlow Serving

- **Definition:** a flexible, high-performance serving system for TensorFlow models
- **Features:** supports versioning, batching, and GPU acceleration
- **Use cases:** ideal for deploying deep learning models in production



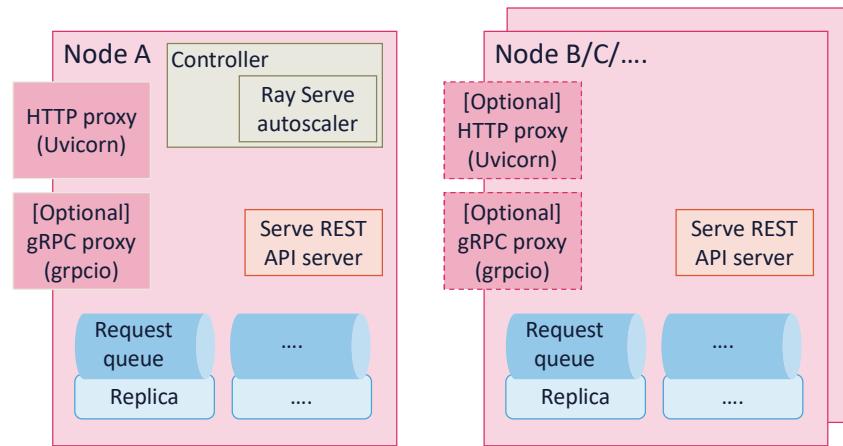
## Introduction to TorchServe

- **Definition:** a model serving framework specifically designed for PyTorch models
- **Features:** supports multi-model serving, model versioning, and logging
- **Use case:** ideal for serving PyTorch models with minimal configuration



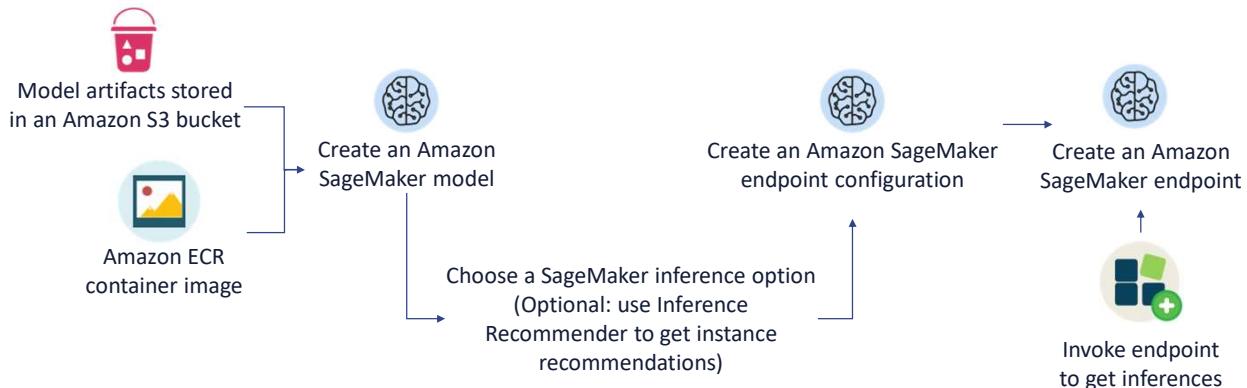
## Introduction to Ray Serve

- **Definition:** a scalable model serving library built on the Ray framework
- **Features:** scalable, low-latency model serving for any framework
- **Use case:** suitable for deploying and scaling ML models in distributed systems



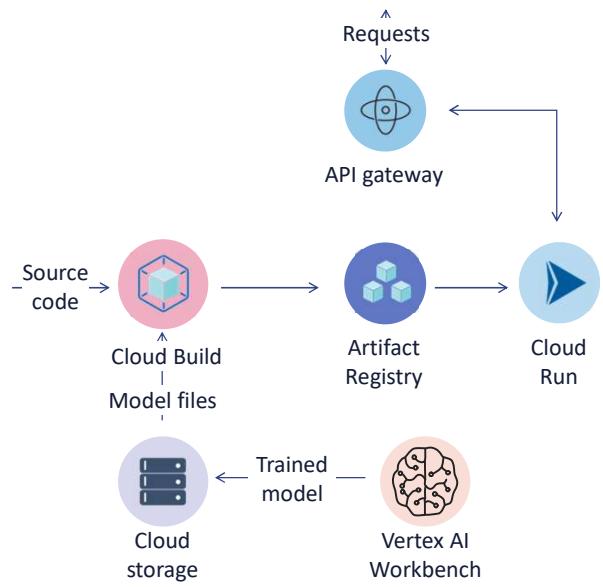
## Steps for Deploying a Model on AWS SageMaker

1. Train and save the model in SageMaker
2. Create a model endpoint for real-time serving
3. Deploy the model and monitor performance



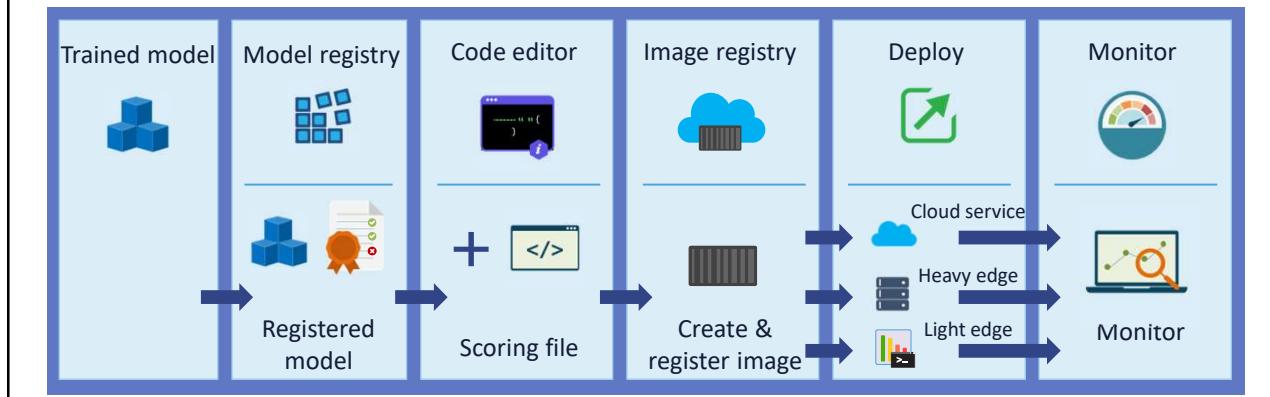
## Deploying Models on Google AI Platform

1. Export the model to Google Cloud Storage
2. Create a model and version in AI Platform
3. Deploy the model to an endpoint for serving



## Deploying Models on Azure Machine Learning

1. Register the model in the Azure ML workspace
2. Create a container image with the model and its dependencies
3. Deploy the model as a web service



**Monitoring**  
set up monitoring to track model performance and usage



**Scalability**  
ensure the deployment can handle varying workloads



**Security**  
implement authentication and authorization for model endpoints



**Best Practices for Model Deployment on Cloud Platforms**

## Knowledge Check

### Matching

1. Match the model serving method with its best use case:

<ul style="list-style-type: none"><li>A. Batch serving</li><li>B. Real-time serving</li><li>C. A/B testing</li><li>D. Blue-green deployment</li><li>E. Canary deployment</li></ul>	<ul style="list-style-type: none"><li>— 1. Minimizing downtime during updates</li><li>— 2. Generating periodic reports</li><li>— 3. Comparing model performance</li><li>— 4. Fraud detection</li><li>— 5. Gradual rollout for new feature</li></ul>
--	---

## Knowledge Check

2. What is the primary advantage of deploying a model using containerization?
  - A. Simplifies hyperparameter tuning
  - B. Ensures consistent environment across deployments
  - C. Reduces the need for feature engineering
  - D. Improves model accuracy automatically



## Discussion

### Ensuring Compliance and Ethical Concerns in MLOps

- Recognize the role of compliance and ethics in MLOps
- Identify risks related to data privacy and fairness
- Explore strategies to upload ethical standards



### Topic 2: Legal and Compliance Issues in MLOps

- Legal risks in deploying machine learning models
- Importance of adhering to industry-specific regulations
- Overview of GDPR and HIPAA
- Industry-specific legal frameworks
- Ensuring fairness and accountability in ML models
- Audit trails and explainability requirements
- Implementing model governance



## Topic 2: Legal and Compliance Issues in MLOps

- Balancing innovation with regulatory restrictions
- Addressing evolving compliance standards
- Best practices for legal and compliance in MLOps

### Legal Risks in Deploying Machine Learning Models

#### Data privacy:

unlawful use or sharing of personal data

#### Bias and discrimination:

unintentional bias leading to unfair outcomes

#### Data privacy:

misuse of proprietary models or data



## Importance of Adhering to Industry-Specific Regulations

### Regulatory compliance:

mandatory for industries like healthcare and finance

### Reputation management:

protect organization from legal and reputational damage

### Ethical AI:

promotes responsible and transparent use of AI technologies

## Overview of GDPR and HIPAA

- **GDPR** (General Data Protection Regulations):
  - **Scope:** applies to data processing in the EU
  - **Requirements:** data subject rights, data protection by design
- **HIPAA** (Health Insurance Portability and Accountability Act):
  - **Scope:** protect information in the US
  - **Requirements:** safeguard for data privacy and security

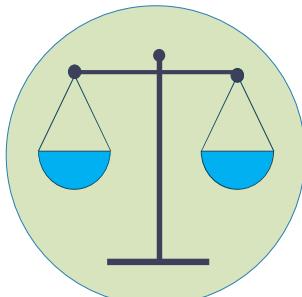


## Industry-Specific Legal Framework



- **Finance:** PCI-DSS for payment data security
- **Healthcare:** CCPA for California-specific data privacy regulations
- **Education:** FERPA for protecting student information

## Ensuring Fairness and Accountability in ML Models



- **Bias mitigation:** techniques to detect and reduce model bias
- **Transparency:** providing clear documentation of model decisions
- **Accountability:** implementing processes for oversight and model auditing

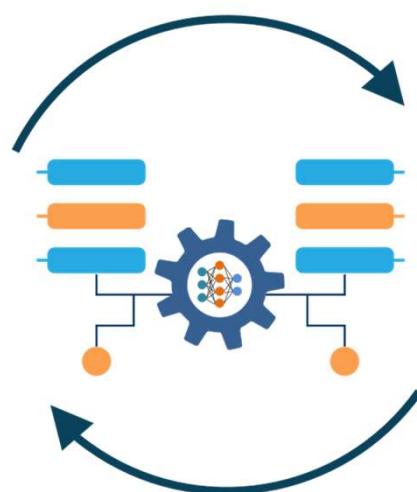
## Audit Trails and Explainability Requirements



- **Audit trail:** maintaining a record of data processing activities and model decisions
- **Explainability:** providing understandable explanations of model outputs
- **Compliance:** required for meeting regulations like GDPR and CCPA

## Implementing Model Governance

- Model lifecycle management: track model development, deployment, and monitoring
- Version control: manage and document different model versions
- Access control: limit access to sensitive models and data



## Balancing Innovation with Regulatory Restrictions



- **Innovation vs. compliance:** navigating the tension between rapid innovation and stringent regulations
- **Risk management:** implementing risk management strategies to innovate responsibly
- **Stakeholder collaboration:** working with legal and compliance teams early in the process

## Addressing Evolving Compliance Standards



- **Continuous monitoring:** stay updated with changes in compliance standards
- **Training and awareness:** regular training for teams on compliance requirements
- **Proactive adaptation:** implement changes proactively to align with new standards

## Best Practices for Legal and Compliance in MLOps

### Data governance:

implement robust data management and security practices



### Ethical AI framework:

establish a framework for ethical AI development and deployment



### Compliance by design:

integrate compliance considerations into the MLOps pipeline



## Knowledge Check

3. Machine learning models deployed in production can unintentionally lead to biased or unfair outcomes.
  - A. True
  - B. False



## Topic 3: Containerizing ML Models with Docker

- What is Docker?
- Key Docker concepts
- Installing Docker
- Basic Docker commands
- Writing a Dockerfile
- Building a Docker image
- Best practices for managing Docker images
- Running a containerized ML model locally

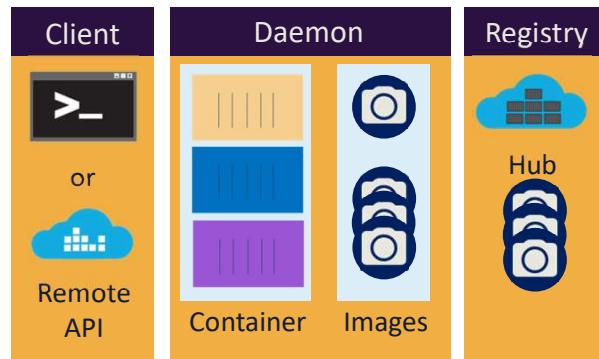


## Topic 3: Containerizing ML Models with Docker

- Pushing Docker images to Docker Hub
- Deploying Docker containers on cloud platforms

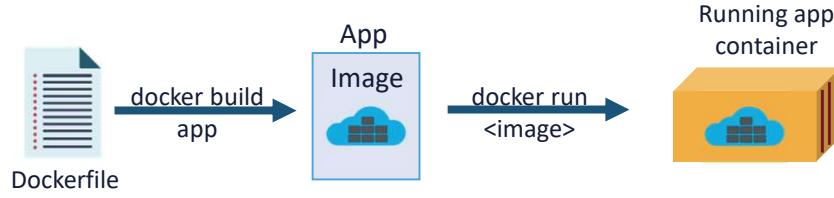
## What Is Docker?

- **Definition:** Docker is a platform that enables developers to package applications and their dependencies into portable containers
- **Purpose:** ensures consistent environment across different systems
- **Benefits:** portability, scalability, and ease of deployment



## Key Docker Concepts

- **Docker container:** a lightweight, standalone, and executable package of software that includes everything needed to run an application
- **Docker image:** a read-only template used to create containers containing instructions to run the application
- **Dockerfile:** a script containing a set of instructions to build a Docker image



1. Download Docker Desktop from the Docker website
2. Follow installation instructions for your OS (Windows, macOS, Linux)
3. Verify installation with the command `docker --version`

## Installing Docker

## Basic Docker Commands

- `docker run`: run a command in a new container
- `docker ps`: list running containers
- `docker build`: build an image from a docker file
- `docker pull`: download an image from registry



## Writing a Dockerfile

- **FROM**: specifies the base image, such as `python:3.8`
- **RUN**: executes commands to install dependencies like libraries
- **COPY**: copies files from the host to the container, such as model files
- **CMD**: specifies the default command to run when the container starts, like `python app.py`

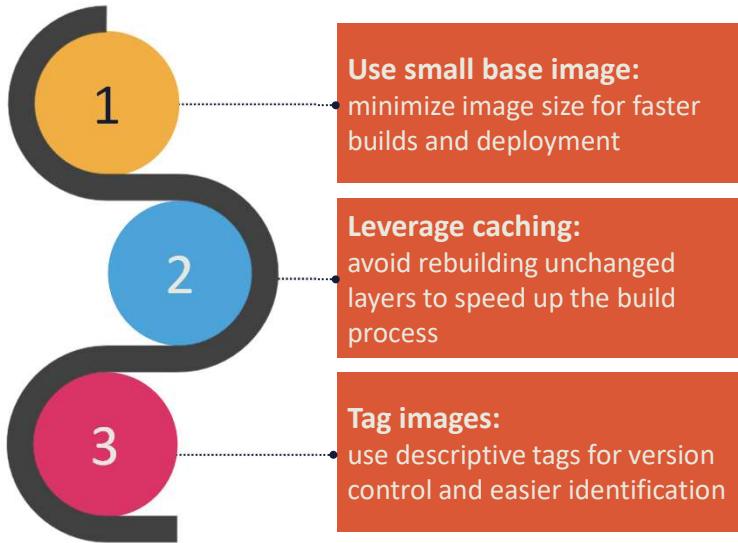
```
● ● ●  
1 # Dockerfile Example  
2 FROM python:3.8  
3 RUN pip install numpy pandas scikit-learn  
4 COPY model.pkl /app/model.pkl  
5 CMD ["python", "app.py"]  
6
```

## Building a Docker Image

1. Navigate to the directory containing the Dockerfile
2. Run `docker build -t my-ml-model .` to build the image
3. Verify the image with `docker images`

```
● ● ●  
1 # Step 1: Navigate to the directory containing the Dockerfile  
2 cd path/to/your/dockerfile-directory  
3  
4 # Step 2: Build the Docker image with a tag  
5 docker build -t my-ml-model .  
6  
7 # Step 3: Verify the image by listing all Docker images  
8 docker images
```

## Best Practices for Managing Docker Images



## Knowledge Check

4. Which of the following best defines Docker?
  - A. A platform for managing large datasets
  - B. A tool for creating virtual machines
  - C. A platform for packaging applications and their dependencies into containers
  - D. A software for direct hardware access

## Knowledge Check

5. Which of the following is **not** a key Docker concept?
  - A. Docker Container
  - B. Docker Image
  - C. Dockerfile
  - D. Docker Host

## Lab 8: Dockerize and Deploy ML Models on Cloud Platforms

### Instructions

1. Build a Dockerfile to containerize an ML model
2. Create and run the Docker container locally
3. Push the Docker image to a repository (e.g., Docker Hub)
4. Deploy the container on a cloud platform (e.g., Azure)

Estimated completion time: 25 minutes



## Topic 4: Deploying Models to Cloud Platforms

- What is cloud deployment?
- Why deploy models to the cloud?
- Steps to prepare a trained model for deployment
- Exporting the model in a deployable format
- Setting up a cloud-based environment
- Configuring networking and security
- Introduction to Ray Serve

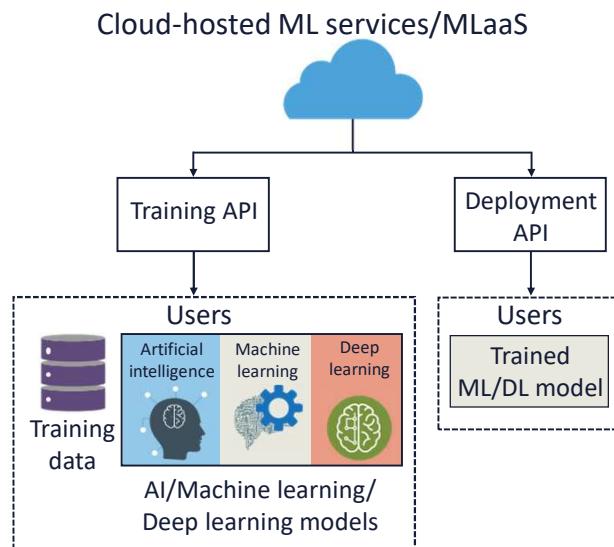


## Topic 4: Deploying Models to Cloud Platforms

- Deploying a model with Ray Serve
- Monitoring and scaling with Ray Serve
- Best practices for using Ray Serve in production

## What Is Cloud Deployment?

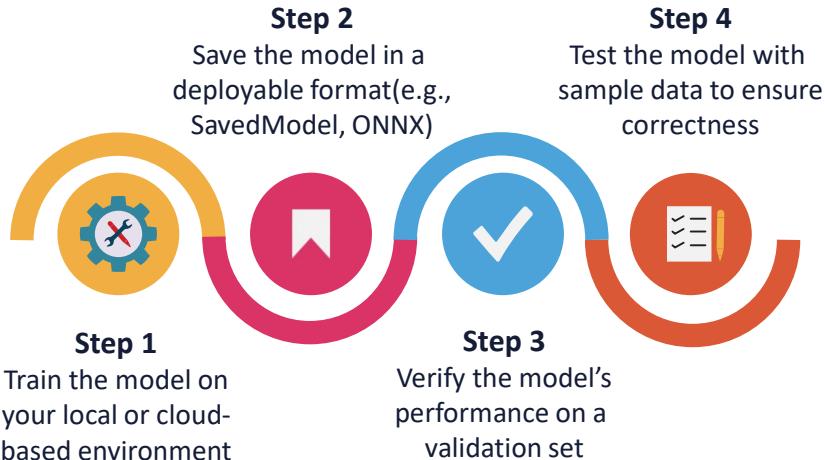
- **Definition:** cloud deployment involves hosting machine learning models on cloud platforms to enable scalable, reliable, and accessible model serving
- **Benefits:** scalability, cost-efficiency, easy integration, and enhanced security
- **Use cases:** real-time predictions, automated retraining, and monitoring



## Why Deploy Models to the Cloud?



## Steps to Prepare a Trained Model for Deployment



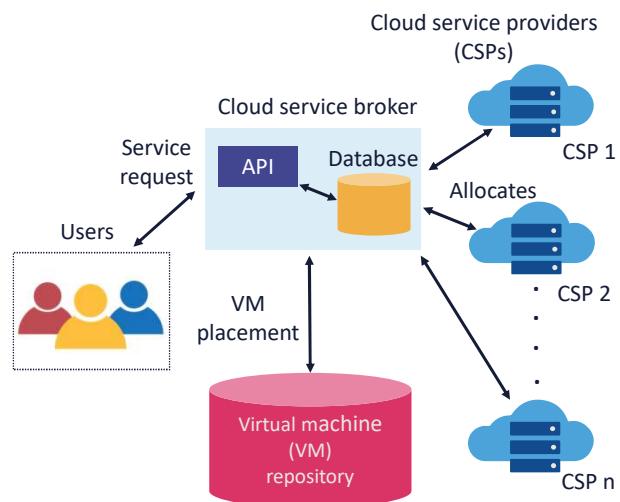
## Exporting the Model in a Deployable Format



- **SavedModel:** a format compatible with TensorFlow serving
- **ONNX:** an open format supported by multiple frameworks like PyTorch and Scikit-Learn
- **PMMI:** a standardized format for sharing models between various tools and platforms

## Setting Up a Cloud-based Environment

- **Select a cloud provider:** choose a platform like AWS, Google Cloud, or Azure
- **Provision resources:** set up virtual machines, storage, and network configurations
- **Install required software:** configure the environment with necessary libraries and tools



## Configuring Networking and Security

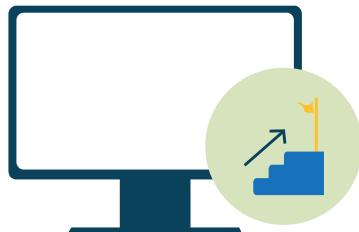
- **Networking:** set up Virtual Private Cloud (VPC) and subnets for secure communication
- **Security groups:** define security rules for inbound and outbound traffic
- **IAM roles:** assign roles and permissions to manage access to resources





## Introduction to Ray Serve

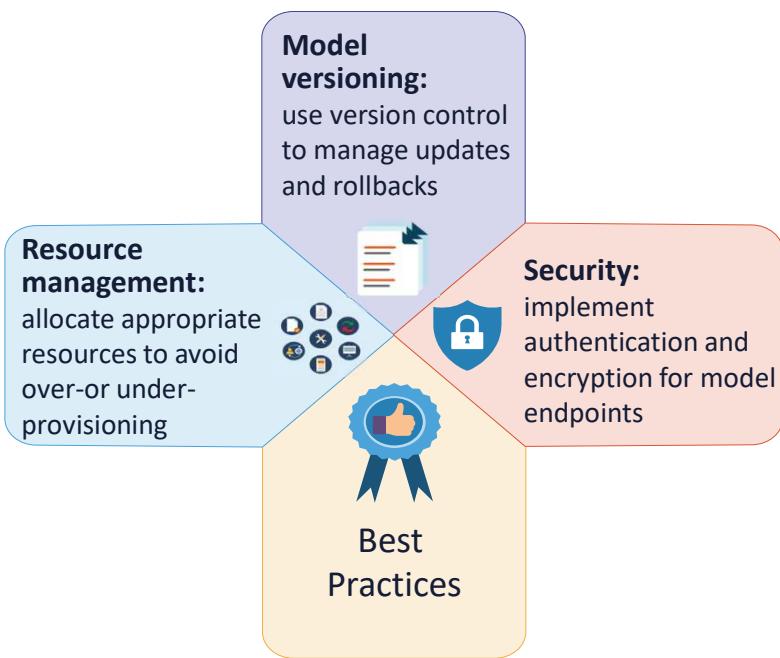
- **Definition:** Ray Serve is a scalable model serving library built on the Ray framework
- **Features:** supports high-throughput, low-latency model serving
- **Use cases:** ideal for deploying machine learning models at scale



## Monitoring and Scaling with Serve

- **Monitoring:** track key metrics such as request latency, error rates, and throughput
- **Scaling:** automatically scale up or down on traffic using Ray Serve's built-in scaling features
- **Logging:** use Ray's logging features to capture and analyze server logs

## Best Practices for Using Ray Serve in Production



## Knowledge Check

6. Cloud deployment enables organizations to scale machine learning models effectively by leveraging cloud platforms' flexibility and resources.
  - A. True
  - B. False



## Topic 5: Federated Training and Edge Deployments

- What is federated learning?
- What is edge computing?
- Federated learning and edge computing integration
- Overview of federated training architecture
- Key components of federated learning
- Advantages of federated training architecture
- Deploying models on edge devices

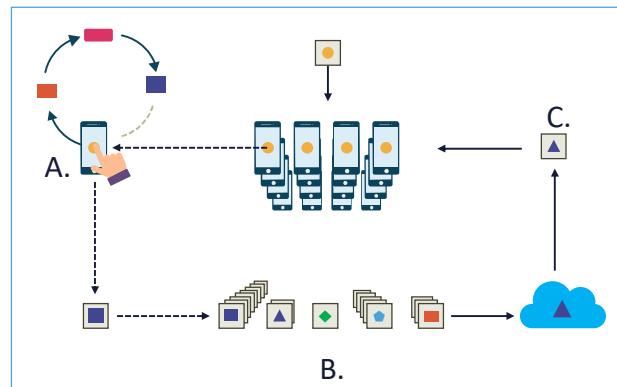


## Topic 5: Federated Training and Edge Deployments

- Considerations for edge deployments
- Edge deployment strategies
- TensorFlow federated for federated learning
- PySyft for Privacy-Preserving Machine Learning
- ONNX for edge model deployment
- Security and privacy challenges
- Technical challenges in federated learning
- Future directions and opportunities

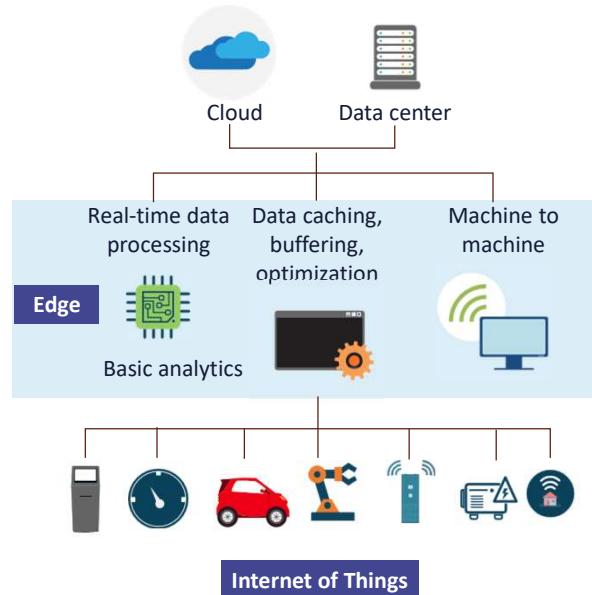
## What Is Federated Learning?

- **Definition:** federated learning is a decentralized approach to machine learning where models are trained across multiple devices without centralizing the data
- **Purpose:** enhances data privacy by keeping data on devices while aggregating only model updates
- **Use cases:** healthcare, finance, and IoT applications where data privacy is critical



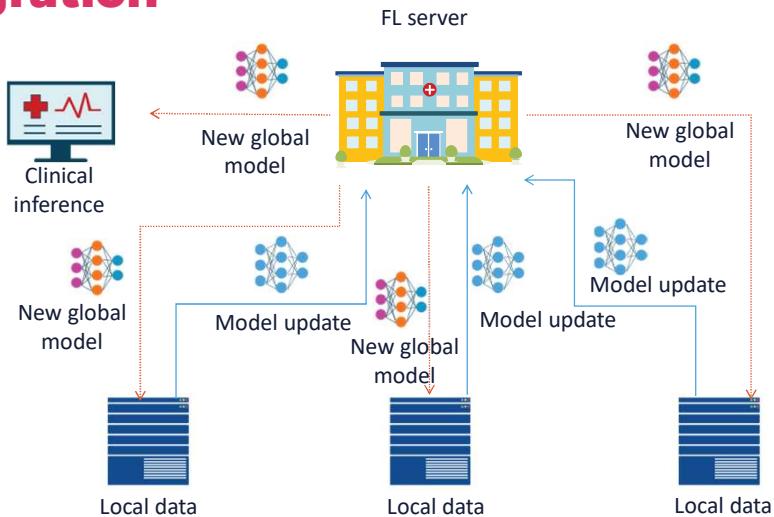
## What Is Edge Computing?

- **Definition:** edge computing involves processing data closer to its source, such as on IoT devices or edge servers
- **Purpose:** reduces latency, saves bandwidth, and enables real-time data processing
- **Use cases:** smart cities, autonomous vehicles, and real-time analytics



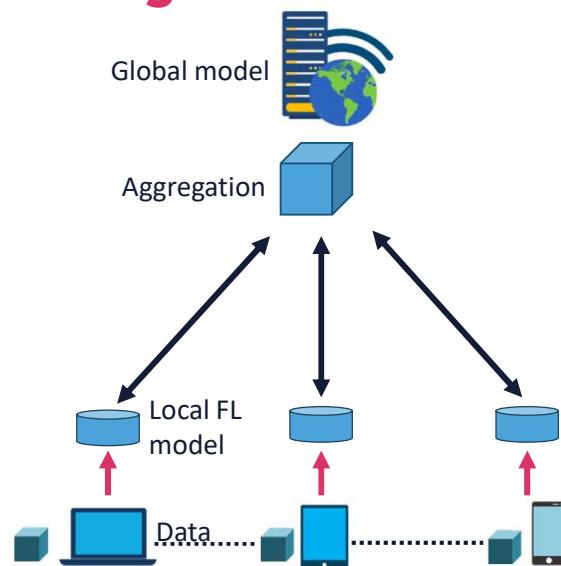
## Federated Learning and Edge Computing Integration

- **Synergy:** federated learning leverages edge computing for on-device model training
- **Data privacy:** ensures data privacy by keeping data on the device
- **Scalability:** scales model training across thousands of edge devices



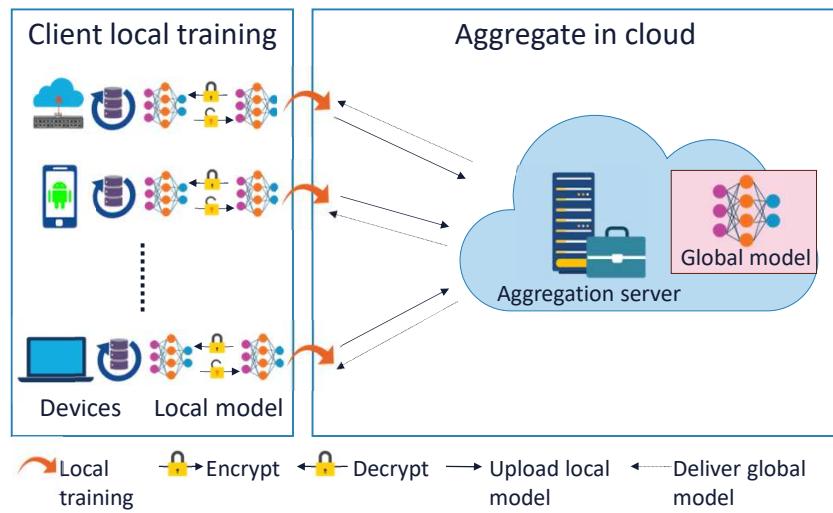
## Overview of Federated Training Architecture

- **Client devices:** train models locally on user devices
- **Server aggregator:** aggregates model updates from all clients
- **Global model:** the aggregated model is shared back with all devices



## Key Components of Federated Learning

- Local model training:** models are trained locally on each device using local data
- Model aggregation:** aggregation server combines local updates to form a global model
- Model distribution:** the global model is distributed back to devices for further training



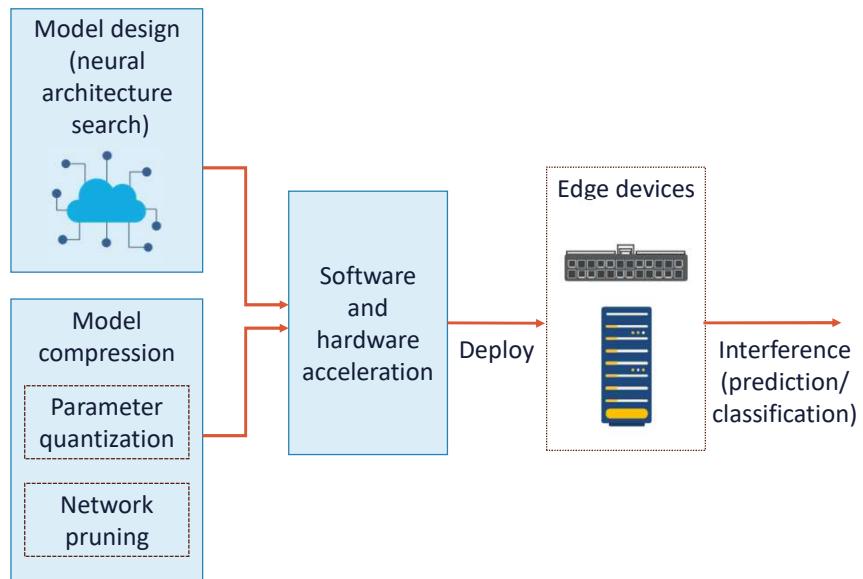
## Advantages of Federated Training Architecture



- Data privacy:** data remains on the device, reducing the risk of data breaches
- Bandwidth efficiency:** only model updates are shared, saving bandwidth
- Scalable learning:** enables model training across a large number of devices

## Key Components of Federated Learning

- **Edge devices:** smartphones, IoT devices, and edge servers
- **Model optimization:** use quantization and pruning to reduce model size
- **Deployment tools:** use frameworks like TensorFlow Lite and ONNX Runtime



## Advantages of Federated Training Architecture

- 
- The icon on the left shows a blue cloud with a rainbow, connected by arrows to a clipboard with a pencil and a network of colored circles representing users or devices.
- **Latency:** ensure low latency for real-time predictions
  - **Resource constraint:** optimize models for limited memory and processing power
  - **Connectivity:** plan for intermittent or no connectivity scenarios

## Edge Deployment Strategies



- **Batch inference:** process multiple data points in batches to optimize resource usage
- **Model caching:** cache models on devices to reduce load times
- **Dynamic model updates:** update models based on changing conditions or new data

- **Features:** provides APIs for implementing federated learning algorithms
- **Use cases:** collaborative model training across multiple clients
- **Integration:** can be integrated with TensorFlow models and data sets

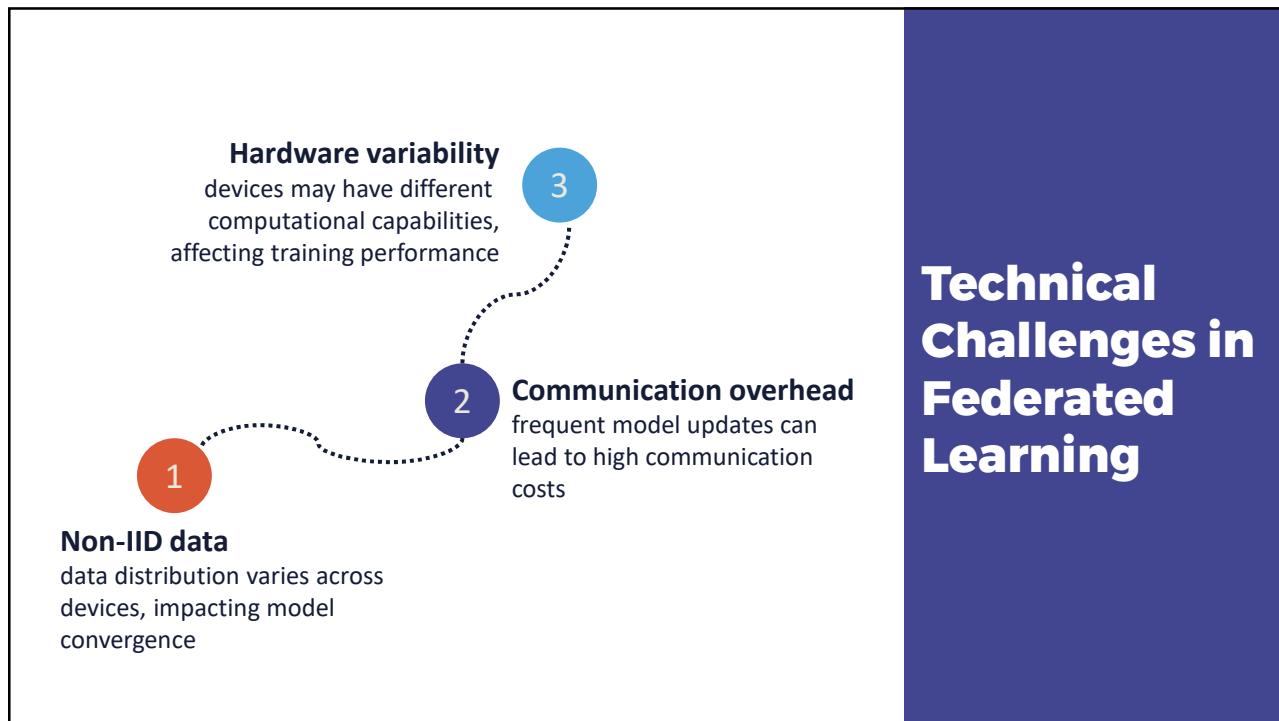
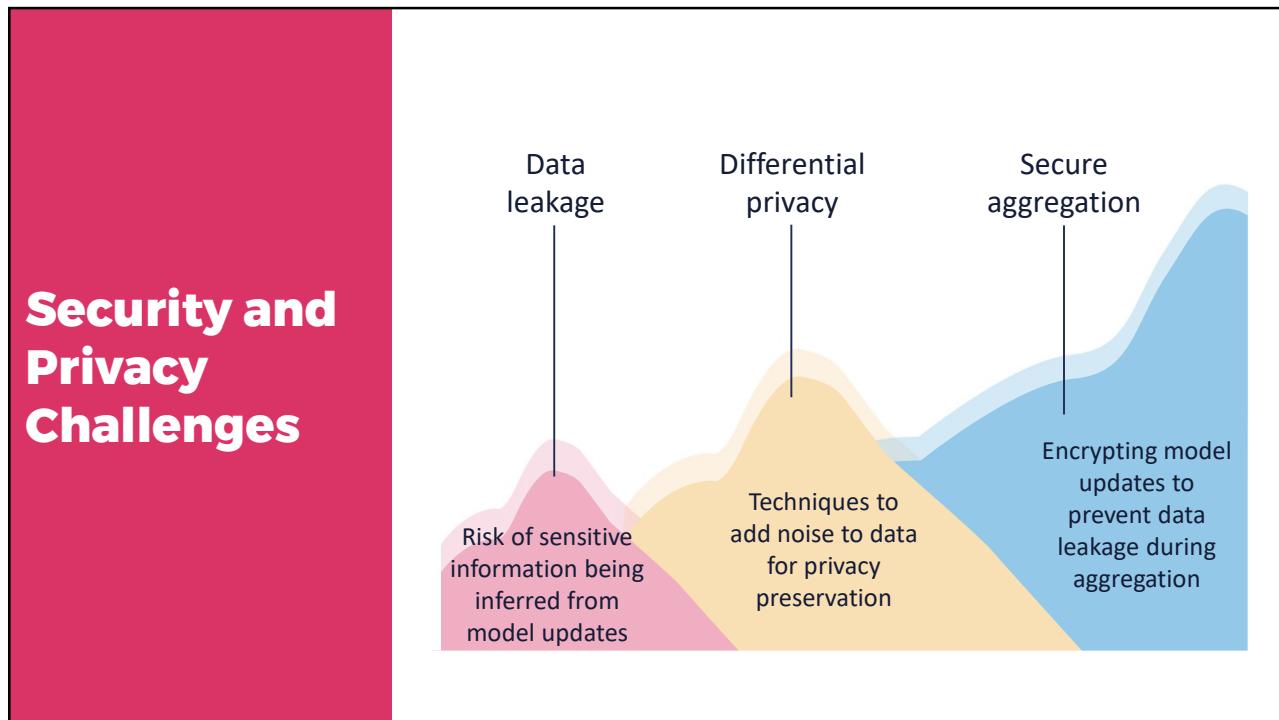
**TensorFlow  
Federated for  
Federated  
Learning**

## PySyft for Privacy-Preserving Machine Learning

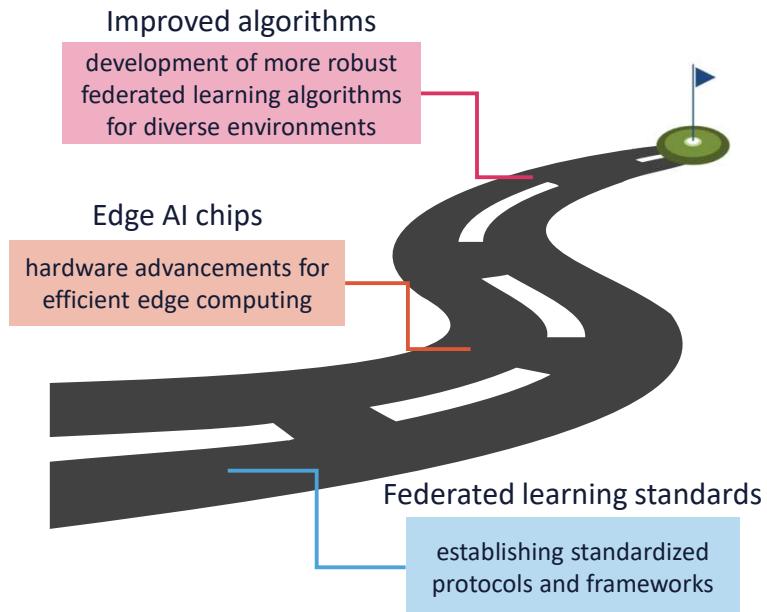
- **Features:** enables federated learning, differential privacy, and encrypted computation
- **Use cases:** secure and privacy-preserving machine learning applications
- **Integration:** works with popular ML frameworks like PyTorch and TensorFlow

- **Features:** provides an open format for AI models, supporting multiple frameworks
- **Use cases:** enables model interoperability and efficient deployment on edge devices
- **Integration:** ONNX Runtime for optimized inference on edge devices

## ONNX for Edge Model Deployment



## Future Directions and Opportunities



### Knowledge Check

7. Which of the following best describes Federated Learning?
  - A. A centralized approach to collect and store data from multiple devices
  - B. A method of training models across devices without centralizing data
  - C. A technique to reduce model complexity
  - D. A system that relies solely on cloud-based data processing

## Knowledge Check

8. What is the primary benefit of Edge Computing in relation to Federated Learning?
- A. Centralizing data processing on the cloud
  - B. Reducing latency by processing data closer to its source
  - C. Decreasing model accuracy
  - D. Increasing the need for high-bandwidth networks



## Topic 6: CI/CD for ML

- What is CI/CD in machine learning?
- Why is CI/CD important in ML?
- Key components of CI/CD in ML
- Tools for CI/CD in ML
- Setting up a CI/CD pipeline for ML
- Best practices for CI/CD pipelines in ML
- Implementing continuous training (CT)
- Linking CI/CD pipelines with experiment tracking
- Benefits of CI/CD and experiment tracking integration

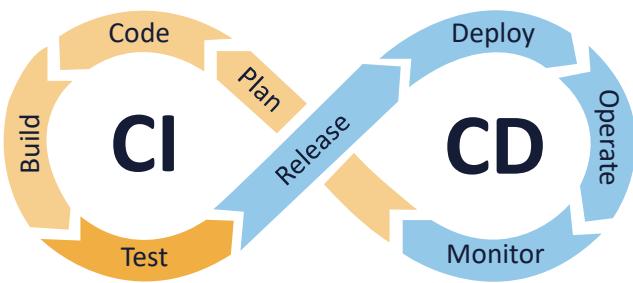


## Topic 6: CI/CD for ML

- Implementing automated experiment tracking in CI/CD
- Setting up automated model validation
- Automating model testing in CI/CD pipelines
- Best practices for automated model validation
- Integrating CI/CD with model validation and testing
- Best practices for CI/CD in ML

### What Is CI/CD in Machine Learning?

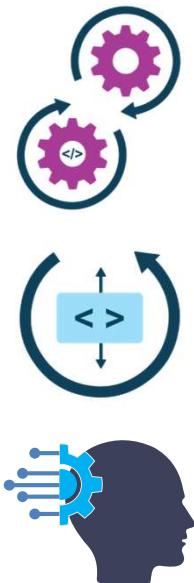
- **Definition:** CI/CD stands for Continuous Integration and Continuous Deployment, a set of practices to automate and streamline the ML lifecycle
- **Purpose:** helps in automating the model building, testing, and deployment processes
- **Benefits:** improves code quality, accelerates development, and reduces manual errors





## Why Is CI/CD Important in ML?

- **Consistency:** ensures consistent model performance across environments
- **Automation:** automates repetitive tasks like testing and deployment
- **Collaboration:** enhances collaboration by providing a common framework for model management

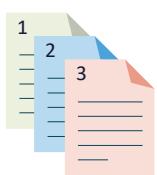


## Key Components of CI/CD in ML

- **Continuous Integration (CI):** automates the process of integrating code changes and testing them
- **Continuous Deployment (CD):** automates the deployment of models to production environments
- **Continuous Training (CT):** retrains models when new data is available

## Tools for CI/CD in ML

- **Jenkins:** an open-source automation server for building and deploying models
- **GitHub Actions:** automates workflows directly from your GitHub repository
- **GitLab CI:** provides CI/CD pipelines integrated with GitLab repositories



### Step 1

Create a version controlled repository for your model code and data



### Step 2

Define CI/CD workflows using a tool like GitHub Actions, or GitLab CI



### Step 3

Configure automated triggers for code commits, model training and testing

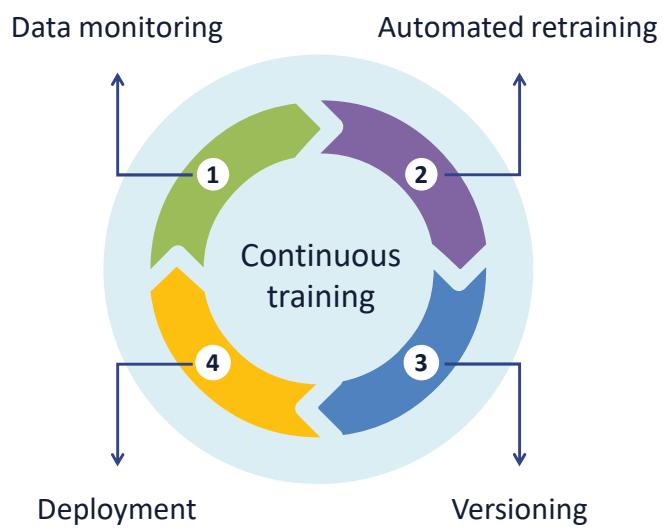
## Setting Up a CI/CD Pipeline for ML

## Best Practices for CI/CD Pipelines in ML



- **Version control:** use version control for code, data, and model artifacts
- **Automated testing:** implement automated tests for model performance and data validation
- **Isolation:** use isolated environments for different stages (e.g., staging and production)

## Implementing Continuous Training

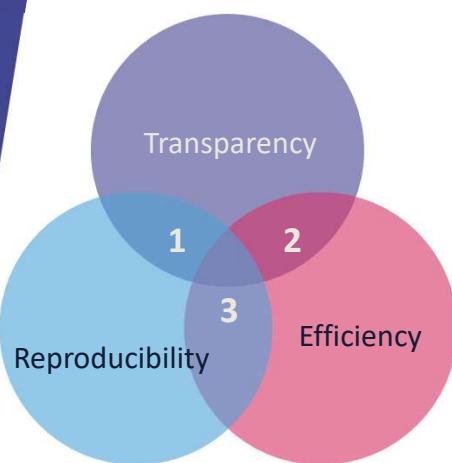


## Linking CI/CD Pipelines with Experiment Tracking



- **Experiment tracking tool:** use tools like MLflow or Weights & Biases for tracking
- **Integration points:** integrate CI/CD pipelines with experiment tracking tools to log metrics, parameters, and model artifacts
- **Benefits:** enables reproducibility, easy comparison of different model versions, and continuous monitoring

### Benefits of CI/CD and Experiment Tracking Integration



1 Ensures experiments can be reproduced with the same conditions and data

2 Provides a clear record of all changes and their impact on model performance

3 Streamlines the process of selecting and deploying the best model version

## Implementing Automated Experiment Tracking in CI/CD



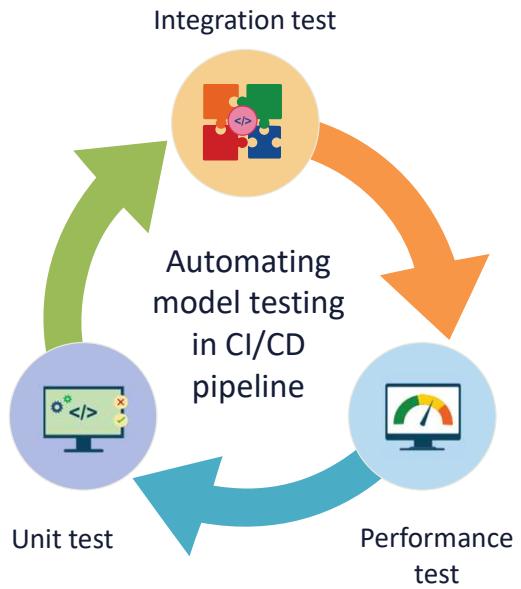
- **Automation:** use CI/CD pipelines to automate the logging of experiments in tools like MLflow
- **Tracking metrics:** automatically track metrics, parameters, and artifacts during model training
- **Integration:** integrate experiment tracking directly into CI/CD workflows for seamless logging

## Setting Up Automated Model Validation

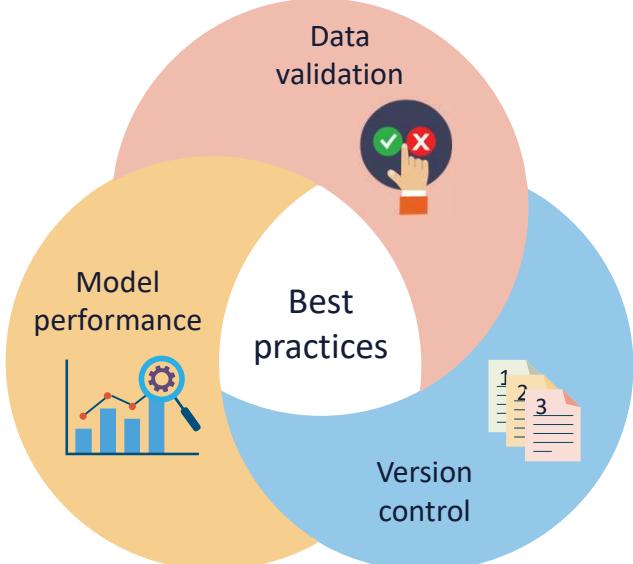


- **Validation tests:** define automated tests to validate model performance on unseen data
- **Regression tests:** set up tests to ensure new model versions perform as well as or better than previous versions
- **Monitoring:** continuously monitor validation results for anomalies

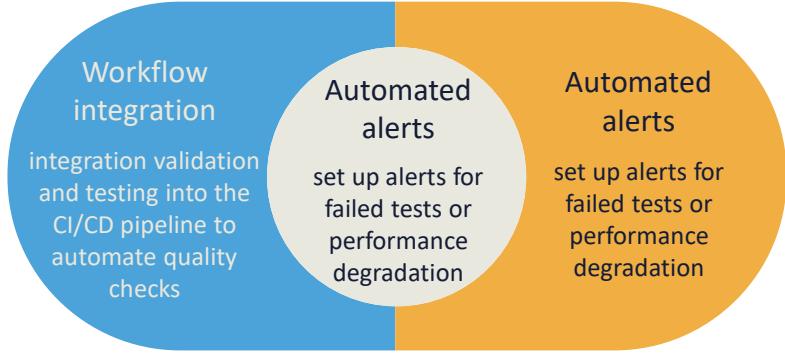
## Automating Model Testing in CI/CD Pipelines



## Best Practices for Automated Model Validation



## Integrating CI/CD with Model Validation and Testing



### Best Practices for CI/CD in ML

- **Documentation**: document CI/CD workflows, test cases, and validation criteria
- **Collaboration**: encourage collaboration between data scientists, engineers, and DevOps teams
- **Iterative improvement**: continuously improve CI/CD processes based on feedback and new requirements



## Knowledge Check

9. CI/CD in machine learning only improves code quality and does not affect deployment speed.
- A. True
  - B. False

## Knowledge Check

10. What is the primary purpose of CI/CD in machine learning?
- A. To reduce data storage costs
  - B. To automate and streamline model building, testing, and deployment
  - C. To increase the complexity of model deployment
  - D. To decentralize model training across multiple devices



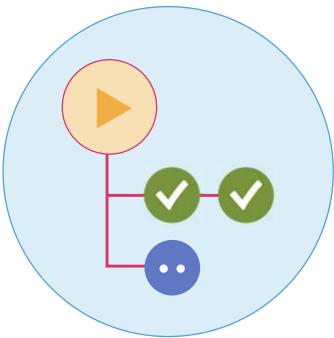
## Topic 7: Setting Up CI/CD for Pipeline for ML

- What is GitHub Actions for CI/CD?
- Key concepts in CI/CD for machine learning
- Benefits of automating ML pipelines using CI/CD
- Setting up a CI/CD pipeline in GitHub
- Writing a YAML configuration for GitHub Actions
- Automating model training and deployment



## Topic 7: Setting Up CI/CD for Pipeline for ML

- Integrating MLflow for experiment tracking
- Configuring MLflow in the CI/CD pipeline
- Running the CI/CD pipeline
- Validating pipeline performance

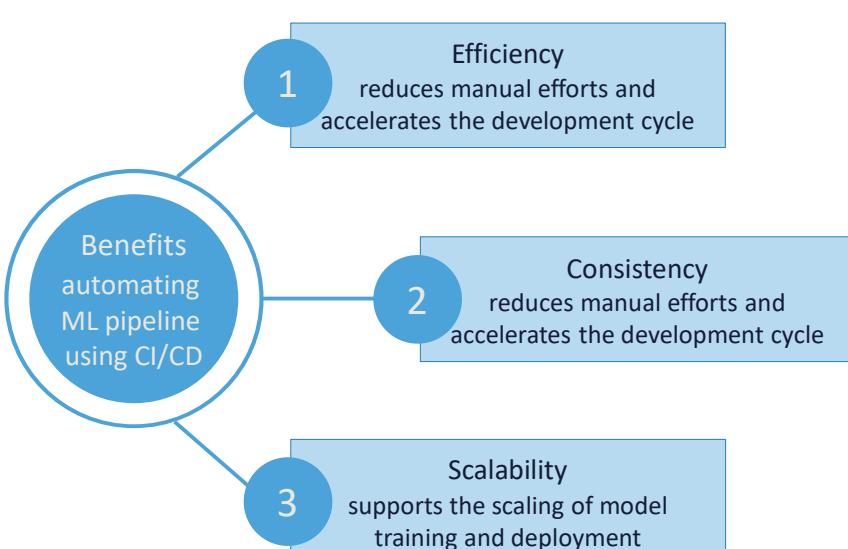


GitHub actions

## What Is GitHub Actions for CI/CD?

- **Definition:** GitHub Actions is a CI/CD platform that allows developers to automate workflows directly from their GitHub repositories
- **Benefits:** integrates seamlessly with GitHub repositories, supports custom workflows, and provides a robust environment for automating ML pipelines
- **Use cases:** automating code builds, tests, deployments, and model training

### Benefits of Automating ML Pipelines Using CI/CD

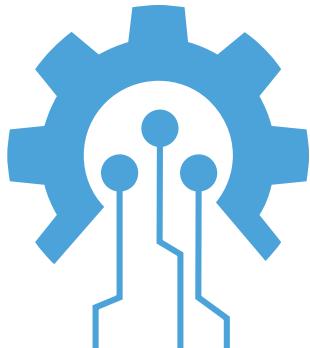


1 Efficiency  
reduces manual efforts and accelerates the development cycle

2 Consistency  
reduces manual efforts and accelerates the development cycle

3 Scalability  
supports the scaling of model training and deployment

## Automating Model Training and Deployment



- **Training job:** automate model training using CI/CD workflows
- **Deployment job:** automate model deployment to cloud platforms or web services
- **Artifact storage:** automate model deployment to cloud platforms or web services

## Integrating MLflow for Experiment Tracking

- **Logging models:** automatically log models and their parameters in MLflow
- **Tracking metrics:** track and compare model performance metrics
- **Model registry:** use MLflow's model registry to manage different model versions

```
1 import mlflow
2 import mlflow.sklearn
3 from sklearn.ensemble import RandomForestClassifier
4 from sklearn.metrics import accuracy_score
5 from sklearn.model_selection import train_test_split
6 from sklearn.datasets import load_iris
7
8 # Load data
9 data = load_iris()
10 X_train, X_test, y_train, y_test = train_test_split(data.data, data.target, test_size=0.2)
11
12 # Start an MLFlow experiment
13 mlflow.set_experiment("Iris Classification")
14
15 with mlflow.start_run():
16     # Model training
17     model = RandomForestClassifier(n_estimators=100, max_depth=3)
18     model.fit(X_train, y_train)
19
20     # Log model and parameters
21     mlflow.sklearn.log_model(model, "model")
22     mlflow.log_param("n_estimators", 100)
23     mlflow.log_param("max_depth", 3)
24
25     # Predict and log metrics
26     predictions = model.predict(X_test)
27     accuracy = accuracy_score(y_test, predictions)
28     mlflow.log_metric("accuracy", accuracy)
29
30     # Register model in the MLflow Model Registry
31     mlflow.register_model("runs:///{}/model".format(mlflow.active_run().info.run_id), "Iris_Model")
32
33 print("Experiment logged in MLflow.")
```

## Configuring MLflow in the CI/CD Pipeline

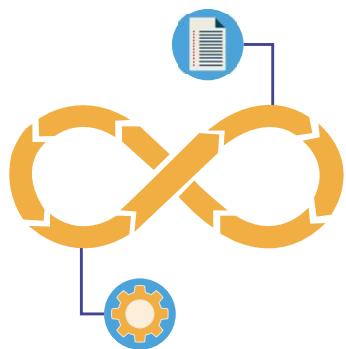
- **Install MLflow:** add MLflow as a CI/CD pipeline dependency
- **Log experiments:** use MLflow APIs to log experiments, metrics, and artifacts during training
- **Track model versions:** automatically track model versions in the MLflow registry

```

1 # Log Experiments: Use MLflow APIs to log experiments, metrics,
2 # and artifacts in the training script (train.py):
3 import mlflow
4 import mlflow.sklearn
5 from sklearn.ensemble import RandomForestClassifier
6 from sklearn.datasets import load_iris
7 from sklearn.model_selection import train_test_split
8
9 # Load data and split
10 data = load_iris()
11 X_train, X_test, y_train, y_test = train_test_split(data.data, data.target)
12
13 # Start MLflow logging
14 with mlflow.start_run():
15     # Train a model
16     model = RandomForestClassifier()
17     model.fit(X_train, y_train)
18
19     # Log model and parameters
20     mlflow.sklearn.log_model(model, "model")
21     mlflow.log_param("n_estimators", model.n_estimators)
22
23     # Log metrics
24     accuracy = model.score(X_test, y_test)
25     mlflow.log_metric("accuracy", accuracy)
26
27 #Track Model Versions: Register the model in the MLflow registry.
28 mlflow.register_model("runs://<RUN_ID>/model", "iris-classifier")
29

```

## Running the CI/CD Pipeline



- **Trigger events:** start the pipeline manually or based on trigger events like push or pull request
- **Monitoring:** monitor the pipeline execution through the GitHub Actions interface
- **Troubleshooting:** identify and fix errors in real-time using detailed logs



## Validating Pipeline Performance

- **Testing:** validate model performance using automated tests in the pipeline
- **Metrics:** review key performance metrics to ensure the model meets the desired criteria
- **Continuous improvement:** iterate on the pipeline based on performance feedback

## Knowledge Check

11. What is one key benefit of automating machine learning pipelines using CI/CD?
- A. Reduces deployment speed
  - B. Increase manual effort
  - C. Enhance model consistency across environments
  - D. Limits scalability



## Interactive Activity

### Designing a CI/CD pipeline

- Collaborate to design a CI/CD pipeline for ML
- Identify key stages and automation points
- Incorporate testing mechanisms for quality assurance

## Lab 9: Implement CI/CD Pipelines for ML with GitHub Actions

### Instructions

1. Set up a CI/CD pipeline for the Iris dataset with GitHub actions
2. Train a machine learning model using scikit-learn
3. Version and deploy the model using Continuous Machine Learning
4. Deploy to GitHub pages from the automated workflow

Estimated completion time: 25 minutes



## Topic 8: Monitoring and Maintaining ML Systems

- Why is monitoring important in ML?
- Key metrics to track for deployed models
- Overview of popular monitoring tools
- Comparing monitoring tools for ML models
- What is model drift and why does it matter?
- Setting up alerts for model drift



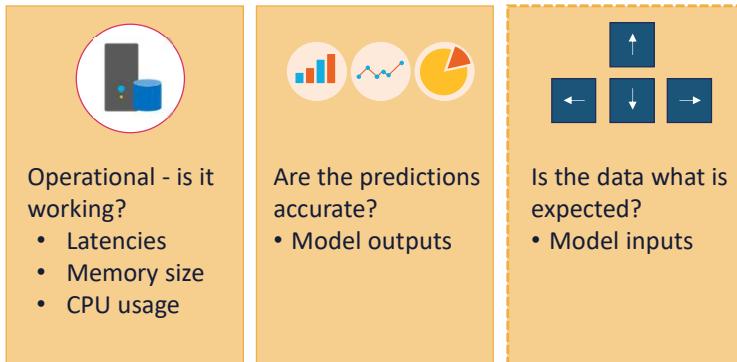
## Topic 8: Monitoring and Maintaining ML Systems

- Techniques for real-time performance monitoring
- Tools and frameworks for real-time monitoring
- Importance of continuous feedback loops
- Scaling monitoring for large-scale deployments

## Why Is Monitoring Important in ML?

- **Definition:** monitoring ML systems involves tracking the performance and health of machine learning models in production
- **Purpose:** ensures that models continue to perform as expected after deployment
- **Metrics:** helps detect issues like model drift, data quality problems, and system errors early

### ML monitoring



## Key Metrics to Track for Deployed Models

- 1 **Accuracy and precision**  
measure the correctness of model predictions
- 2 **Latency**  
time taken for the model to produce predictions
- 3 **Data distribution**  
monitor changes in input data characteristics
- 4 **Model drift**  
track deviations in model performance over time

## Overview of Popular Monitoring Tools



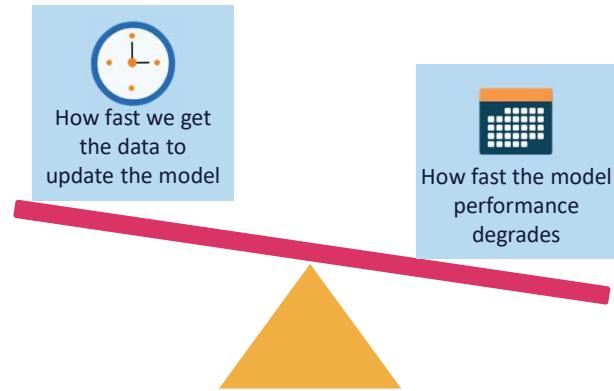
- **Prometheus:** an open-source monitoring system that collects metrics from configured targets
- **Grafana:** a powerful visualization tool used with Prometheus to create interactive dashboards
- **Seldon:** provides tools for monitoring, explaining, and deploying ML models in production

- **Prometheus vs. Grafana:** Prometheus is ideal for collecting and storing metrics, while Grafana excels at visualization
- **Seldon vs. others:** Seldon provides specialized tools for ML monitoring, including drift detection and explainability
- **Key considerations:** choose based on integration needs, scalability, and ease of use

## Comparing Monitoring Tools for ML Models

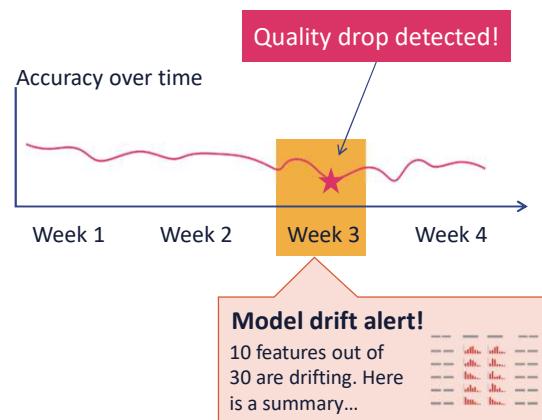
## What Is Model Drift and Why Does It Matter?

- **Definition:** model drift occurs when the statistical properties of the input data change over time, causing model performance to degrade
- **Impact:** leads to inaccurate predictions, reduced model performance, and potential business impact
- **Detection:** requires regular monitoring of model predictions and data characteristics



## Setting Up Alerts for Model Drift

- **Define threshold:** set thresholds for key metrics like accuracy and precision
- **Configure alerts:** set up alerts using monitoring tools like Prometheus and Grafana.
- **Automated actions:** trigger automated retraining or notifications when drift is detected



## Techniques for Real-Time Performance Monitoring

**Log-based monitoring:** capture and analyze logs for model predictions and performance

**Metric-based monitoring:** track key metrics like latency, throughput, and error rates

**Anomaly detection:** use statistical techniques or machine learning models to detect anomalies in real-time

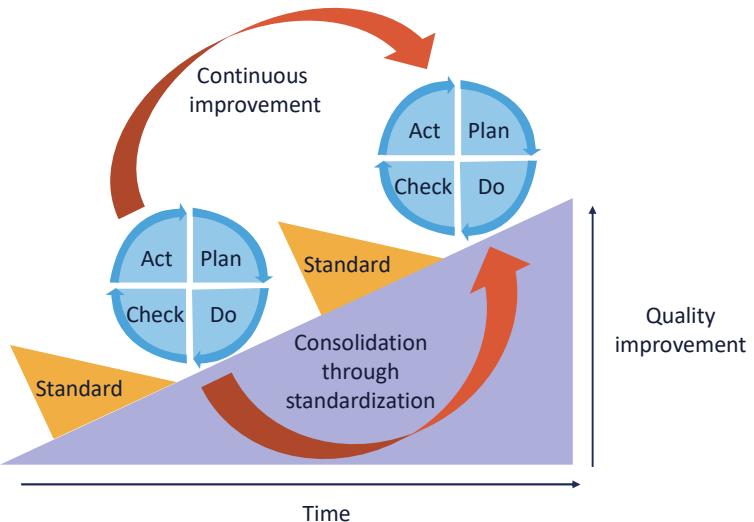


## Tools and Frameworks for Real-Time Monitoring

- **Prometheus & Grafana:** collects and visualizes real-time metrics
- **Seldon Core:** monitors ML models in production, providing insights into model performance
- **ELK Stack:** combines Elasticsearch, Logstash, and Kibana for log-based monitoring and visualization

## Importance of Continuous Feedback Loops

- **Definition:** continuous feedback loops involve collecting and using feedback on model performance to improve models
- **Benefits:** helps identify issues early and adapt models to changing conditions
- **Implementation:** integrate feedback from users and automated systems to update models



## Scaling Monitoring for Large-Scale Deployments

- **Distributed monitoring:** use distributed systems like Prometheus Federation for scalable monitoring
- **Data aggregation:** aggregate data from multiple sources for centralized monitoring
- **Resource management:** aggregate data from multiple sources for centralized monitoring



## Knowledge Check

12. Monitoring machine learning models in production is only necessary during the initial deployment phase.
- A. True
  - B. False

## Knowledge Check

13. Monitoring metrics like accuracy, precision, and latency helps detect issues in model performance early.
- A. True
  - B. False



## Topic 9: Implementing Monitoring Tools

- What are monitoring tools?
- Benefits of using Prometheus and Grafana
- Adding custom metrics to ML models
- Instrumenting the model for monitoring
- Setting up Prometheus for metrics collection
- Exposing metrics for Prometheus



## Topic 9: Implementing Monitoring Tools

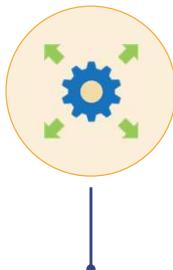
- Verifying metrics collection in Prometheus
- Setting up Grafana for visualization
- Creating a dashboard in Grafana
- Running test scenarios and verifying metrics

## What Are Monitoring Tools?



- **Definition:** monitoring tools are software systems that collect, process, and visualize metrics to track the performance and health of an application
- **Purpose:** ensure ML models perform as expected and help identify issues early
- **Common tools:** Prometheus for metrics collection and Grafana for visualization

## Benefits of Using Prometheus and Grafana



**Scalability**  
both tools can handle large-scale deployments with numerous metrics.

**Flexibility**  
custom metrics can be added to monitor specific aspects of ML models



**Visualization**  
Grafana provides interactive and customizable dashboards for real-time insights

## Setting Up Grafana for Visualization

1. Install Grafana on your machine or cloud environment
2. Add Prometheus as a data source in Grafana
3. Create a new dashboard for visualizing metrics

## Creating Dashboard in Grafana

- **Add panels:** use panels to display different metrics like inference latency, request counts, and error rates
- **Customize visuals:** customize graphs, tables, and alerts based on your monitoring needs
- **Save and share:** save the dashboard and share it with your team for collaborative monitoring





## Running Test Scenarios and Verifying Metrics

- **Simulation scenarios:** run test scenarios to simulate model drift or performance degradation
- **Verify metrics:** check that the metrics are accurately captured and displayed in Grafana
- **Adjust alerts:** adjust alert thresholds based on test results to refine monitoring

## Knowledge Check

14. Which of the following is **not** a purpose of using Prometheus and Grafana in monitoring ML models?
- A. Collect and visualize metrics
  - B. Detect data drift
  - C. Automatically fix model errors
  - D. Track system health and model performance

## Knowledge Check

15. Monitoring tools like Prometheus and Grafana help identify potential issues in ML models by collecting, processing, and visualizing various performance metrics.
- A. True
  - B. False

## Lab 10: Monitor ML Models with Prometheus and Grafana

### Instructions

1. Install monitoring tools Prometheus and Grafana
2. Train a machine learning model using scikit-learn
3. Set up Prometheus to scrape metrics from an ML model
4. Visualize those metrics in Grafana

Estimated completion time: 25 minutes

## Challenge Lab 3: Automate ML Workflows: CI/CD for Wheat Seeds Model with GitHub Actions

### Instructions

1. Set up a CI/CD pipeline for the wheat seeds dataset using GitHub Actions
2. Train a machine learning model using scikit-learn
3. Generate and deploy the model results as an HTML page to GitHub pages

Estimated completion time: 45 minutes



## Discussion

### Best Practices in MLOps

- Reflect on key MLOps practices from the training
- Discuss practical implementation strategies
- Identify improvements for MLOps adoption in your organization

## Agenda Review

Module 1: MLOps and Data

Module 2: Development, Experimentation, and Evaluation

### Module 3: Deployment, Monitoring, and Pipelines

- ✓ Model Serving and Deployment Strategies
- ✓ Legal and Compliance Issues in MLOps
- ✓ Containerizing ML Models with Docker
- ✓ Deploying Models to Cloud Platforms
- ✓ Federated Training and Edge Deployments
- ✓ CI/CD for ML
- ✓ Setting Up CI/CD Pipelines for ML
- ✓ Monitoring and Maintaining ML Systems
- ✓ Implementing Monitoring Tools

## Thank you

For additional courses, visit  
<https://www.globalknowledge.com/us-en/course/200509/mlops-fundamentals/>



## Appendix: Knowledge Check Answer Key



### Knowledge Check

#### Module 1: MLOps and Data

1. What is the primary goal of MLOps?
  - A. Managing the ML model lifecycle
  - B. Developing software applications
  - C. Creating data pipelines
  - D. Monitoring the application performance

## Knowledge Check

### Module 1: MLOps and Data

2. MLOps is solely concerned with the deployment of machine learning models.
  - A. True
  - B. False

## Knowledge Check

### Module 1: MLOps and Data

3. Tools like MLFlow and Kubeflow help automate various stages of the ML lifecycle.
  - A. True
  - B. False

## Knowledge Check

### Module 1: MLOps and Data

4. Which of the following tools is used for experiment tracking in MLOps?
- A. Dockers
  - B. MLFlow**
  - C. Kubernetes
  - D. Terraform

## Knowledge Check

### Module 1: MLOps and Data

5. Which of the following tools is used for managing a virtual environment in Python?
- A. Git
  - B. Docker
  - C. venv**
  - D. Prometheus

## Knowledge Check

### Module 1: MLOps and Data

6. Data versioning is primarily used to track changes in code.
- A. True
  - B. False

## Knowledge Check

### Module 1: MLOps and Data

7. Which of the following tools is commonly used for data versioning in ML projects?
- A. GitHub
  - B. Docker
  - C. DVC
  - D. Kubernetes

## Knowledge Check

### Module 1: MLOps and Data

8. Which of the following is a primary benefit of using feature stores in MLOps?
  - A. Reducing model training time by avoiding data preprocessing
  - B. Ensuring consistency of features across training and serving**
  - C. Eliminating the need for version control in ML projects
  - D. Automating the model deployment process

## Knowledge Check

### Module 1: MLOps and Data

9. What is the main purpose of feature versioning in the feature store?
  - A. Improve the speed of model prediction
  - B. Track changes in features and maintain reproducibility**
  - C. Automate data preprocessing and transformation
  - D. Eliminate the need for feature engineering

## Knowledge Check

### Module 1: MLOps and Data

10. Cross-validation is used to prevent overfitting by ensuring that the model performs well on different subsets of the training data.

- A. True
- B. False

## Knowledge Check

### Module 1: MLOps and Data

11. Which of the following metrics is most suitable for evaluating the performance of a regression model?

- A. Precision
- B. Mean Squared Error (MSE)
- C. Accuracy
- D. F1 Score

## Knowledge Check

### Module 1: MLOps and Data

12. A Sklearn pipeline can include both data preprocessing steps and model training steps in a single workflow.

- A. True
- B. False

## Knowledge Check

### Module 1: MLOps and Data

13. Which of the following tools is commonly used for automating the scheduling of training pipelines in machine learning?

- A. GitHub Actions
- B. Jenkins
- C. Apache Airflow
- D. Dockers

## Knowledge Check

### Module 2: Development, Experimentation and Evaluation

1. Modern approaches in MLOps emphasize an iterative process with automated pipelines and continuous integration.
  - A. True
  - B. False

## Knowledge Check

### Module 2: Development, Experimentation and Evaluation

2. What is the primary focus of the traditional approach to model development?
  - A. Continuous integration and deployment
  - B. Iterative experimentation
  - C. Static datasets and manual feature engineering
  - D. Automated pipelines

## Knowledge Check

### Module 2: Development, Experimentation and Evaluation

3. Local interpretability explains how a machine learning model makes decisions across the entire dataset.
- A. True
  - B. False

## Knowledge Check

### Module 2: Development, Experimentation and Evaluation

4. Which of the following techniques provides both global and local interpretability?
- A. LIME
  - B. SHAP
  - C. Gini Importance
  - D. Integrated Ingredients

## Knowledge Check

### Module 2: Development, Experimentation and Evaluation

5. Which of the following is a key criterion for selecting an algorithm in machine learning?
- A. Problem type
  - B. Data size
  - C. Algorithm popularity
  - D. All of the above

## Knowledge Check

### Module 2: Development, Experimentation and Evaluation

6. What is the primary benefit of experiment tracking in MLOps?
- A. Faster model training
  - B. Improved data preprocessing
  - C. Enables reproducibility and collaboration
  - D. Automates data collection

## Knowledge Check

### Module 2: Development, Experimentation and Evaluation

7. Which tool is commonly used for tracking experiments and visualizing metrics in machine learning projects?
- A. PyTorch
  - B. TensorFlow
  - C. MLflow
  - D. Jupyter

## Knowledge Check

### Module 2: Development, Experimentation and Evaluation

8. Which component of MLflow is responsible for tracking and logging experiments, parameters, and metrics?
- A. MLflow Registry
  - B. MLflow Projects
  - C. MLflow Tracking
  - D. MLflow Models

## Knowledge Check

### Module 2: Development, Experimentation and Evaluation

9. Bayesian optimization focuses the search on the least promising areas of the hyperparameter space.

- A. True
- B. False

## Knowledge Check

### Module 2: Development, Experimentation and Evaluation

10. Which of the following is a disadvantage of Grid Search for hyperparameter tuning?

- A. It randomly samples hyperparameters
- B. It is computationally expensive and time-consuming
- C. It uses probabilistic models to predict performance
- D. It is efficient for large-scale problems

## Knowledge Check

### Module 2: Development, Experimentation and Evaluation

11. Automated hyperparameter tuning always guarantees finding the optimal hyperparameter configuration.

- A. True
- B. False

## Knowledge Check

### Module 3: Deployment, Monitoring, and Pipelines

1. Match the model serving method with its best use case:

A. Batch serving	<u>D</u> 1. Minimizing downtime during updates
B. Real-time serving	<u>A</u> 2. Generating periodic reports
C. A/B testing	<u>C</u> 3. Comparing model performance
D. Blue-green deployment	<u>B</u> 4. Fraud detection
E. Canary deployment	<u>E</u> 5. Gradual rollout for new feature

## Knowledge Check

### Module 3: Deployment, Monitoring, and Pipelines

2. What is the primary advantage of deploying a model using containerization?
  - A. Simplifies hyperparameter tuning
  - B. Ensures consistent environment across deployments**
  - C. Reduces the need for feature engineering
  - D. Improves model accuracy automatically

## Knowledge Check

### Module 3: Deployment, Monitoring, and Pipelines

3. Machine learning models deployed in production can unintentionally lead to biased or unfair outcomes.
  - A. True**
  - B. False

## Knowledge Check

### Module 3: Deployment, Monitoring, and Pipelines

4. Which of the following best defines Docker?
  - A. A platform for managing large datasets
  - B. A tool for creating virtual machines
  - C. A platform for packaging applications and their dependencies into containers
  - D. A software for direct hardware access

## Knowledge Check

### Module 3: Deployment, Monitoring, and Pipelines

5. Which of the following is **not** a key Docker concept?
  - A. Docker Container
  - B. Docker Image
  - C. Dockerfile
  - D. Docker Host

## Knowledge Check

### Module 3: Deployment, Monitoring, and Pipelines

6. Cloud deployment enables organizations to scale machine learning models effectively by leveraging cloud platforms' flexibility and resources.

A. True

B. False

## Knowledge Check

### Module 3: Deployment, Monitoring, and Pipelines

7. Which of the following best describes Federated Learning?
- A. A centralized approach to collect and store data from multiple devices
- B. A method of training models across devices without centralizing data
- C. A technique to reduce model complexity
- D. A system that relies solely on cloud-based data processing

## Knowledge Check

### Module 3: Deployment, Monitoring, and Pipelines

8. What is the primary benefit of Edge Computing in relation to Federated Learning?
  - A. Centralizing data processing on the cloud
  - B. Reducing latency by processing data closer to its source**
  - C. Decreasing model accuracy
  - D. Increasing the need for high-bandwidth networks

## Knowledge Check

### Module 3: Deployment, Monitoring, and Pipelines

9. CI/CD in machine learning only improves code quality and does not affect deployment speed.
  - A. True
  - B. False**

## Knowledge Check

### Module 3: Deployment, Monitoring, and Pipelines

10. What is the primary purpose of CI/CD in machine learning?
- A. To reduce data storage costs
  - B. To automate and streamline model building, testing, and deployment**
  - C. To increase the complexity of model deployment
  - D. To decentralize model training across multiple devices

## Knowledge Check

### Module 3: Deployment, Monitoring, and Pipelines

11. What is one key benefit of automating machine learning pipelines using CI/CD?
- A. Reduces deployment speed
  - B. Increase manual effort
  - C. Enhance model consistency across environments**
  - D. Limits scalability

## Knowledge Check

### Module 3: Deployment, Monitoring, and Pipelines

12. Monitoring machine learning models in production is only necessary during the initial deployment phase.

- A. True
- B. False

## Knowledge Check

### Module 3: Deployment, Monitoring, and Pipelines

13. Monitoring metrics like accuracy, precision, and latency helps detect issues in model performance early.

- A. True
- B. False

## Knowledge Check

### Module 3: Deployment, Monitoring, and Pipelines

14. Which of the following is **not** a purpose of using Prometheus and Grafana in monitoring ML models?
- A. Collect and visualize metrics
  - B. Detect data drift
  - C. Automatically fix model errors
  - D. Track system health and model performance

## Knowledge Check

### Module 3: Deployment, Monitoring, and Pipelines

15. Monitoring tools like Prometheus and Grafana help identify potential issues in ML models by collecting, processing, and visualizing various performance metrics.
- A. True
  - B. False